



ceph


**BLUESTORE: A NEW, FASTER STORAGE BACKEND
FOR CEPH**

SAGE WEIL
VAULT - 2016.04.21

OUTLINE



- Ceph background and context
 - FileStore, and why POSIX failed us
 - NewStore – a hybrid approach
- BlueStore – a new Ceph OSD backend
 - Metadata
 - Data
- Performance
- Upcoming changes
- Summary



MOTIVATION



- Object, block, and file storage in a single cluster
 - All components scale horizontally
 - No single point of failure
 - Hardware agnostic, commodity hardware
 - Self-manage whenever possible
 - Open source (LGPL)
-
- “A Scalable, High-Performance Distributed File System”
 - “performance, reliability, and scalability”



CEPH COMPONENTS



OBJECT



RGW

A web services gateway for object storage, compatible with S3 and Swift

BLOCK



RBD

A reliable, fully-distributed block device with cloud platform integration

FILE



CEPHFS

A distributed file system with POSIX semantics and scale-out metadata management

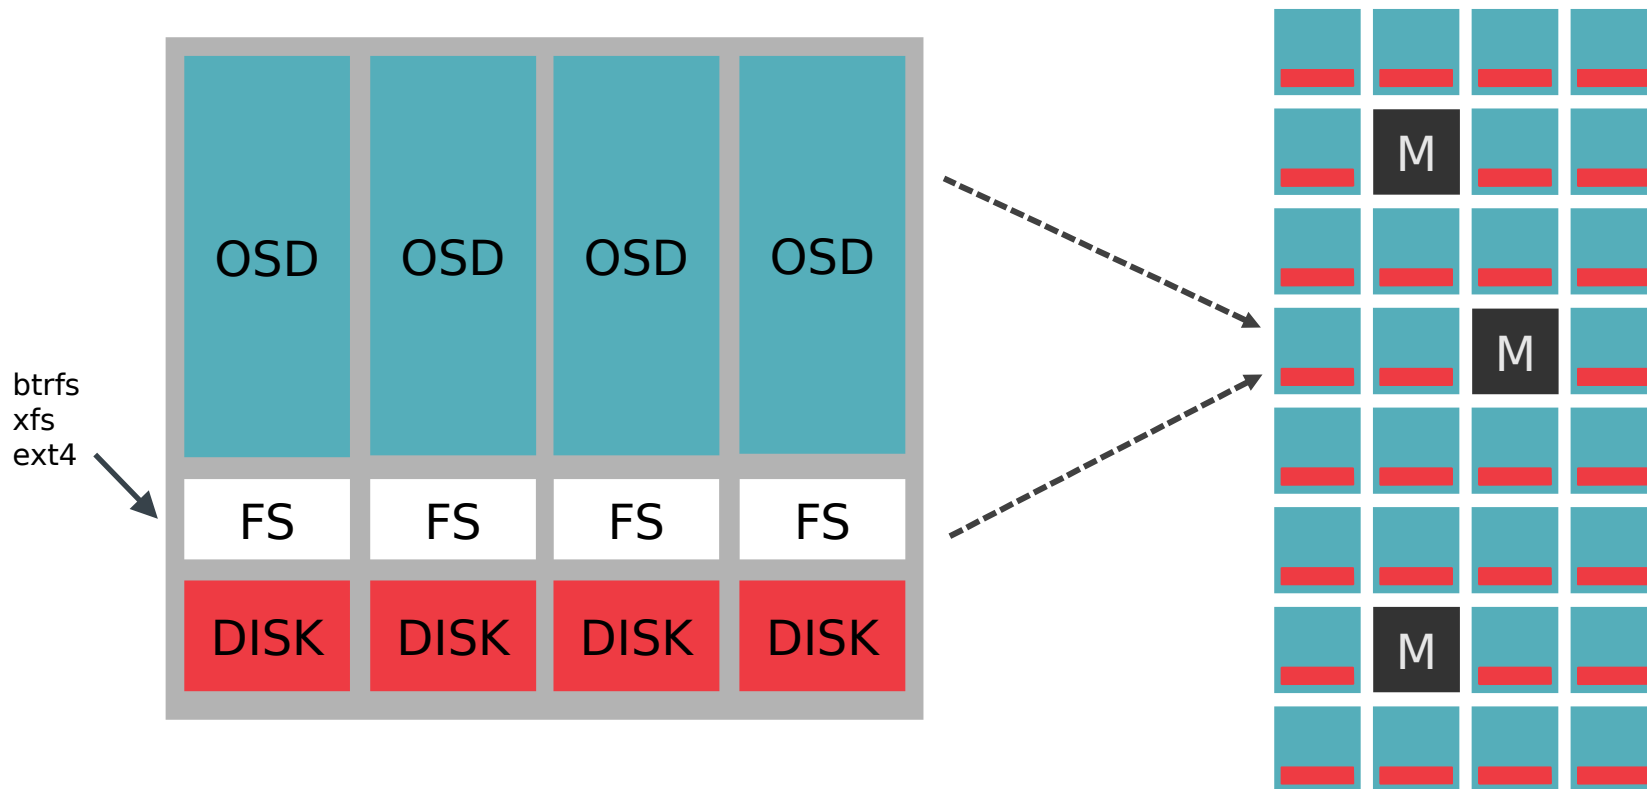
LIBRADOS

A library allowing apps to directly access RADOS (C, C++, Java, Python, Ruby, PHP)

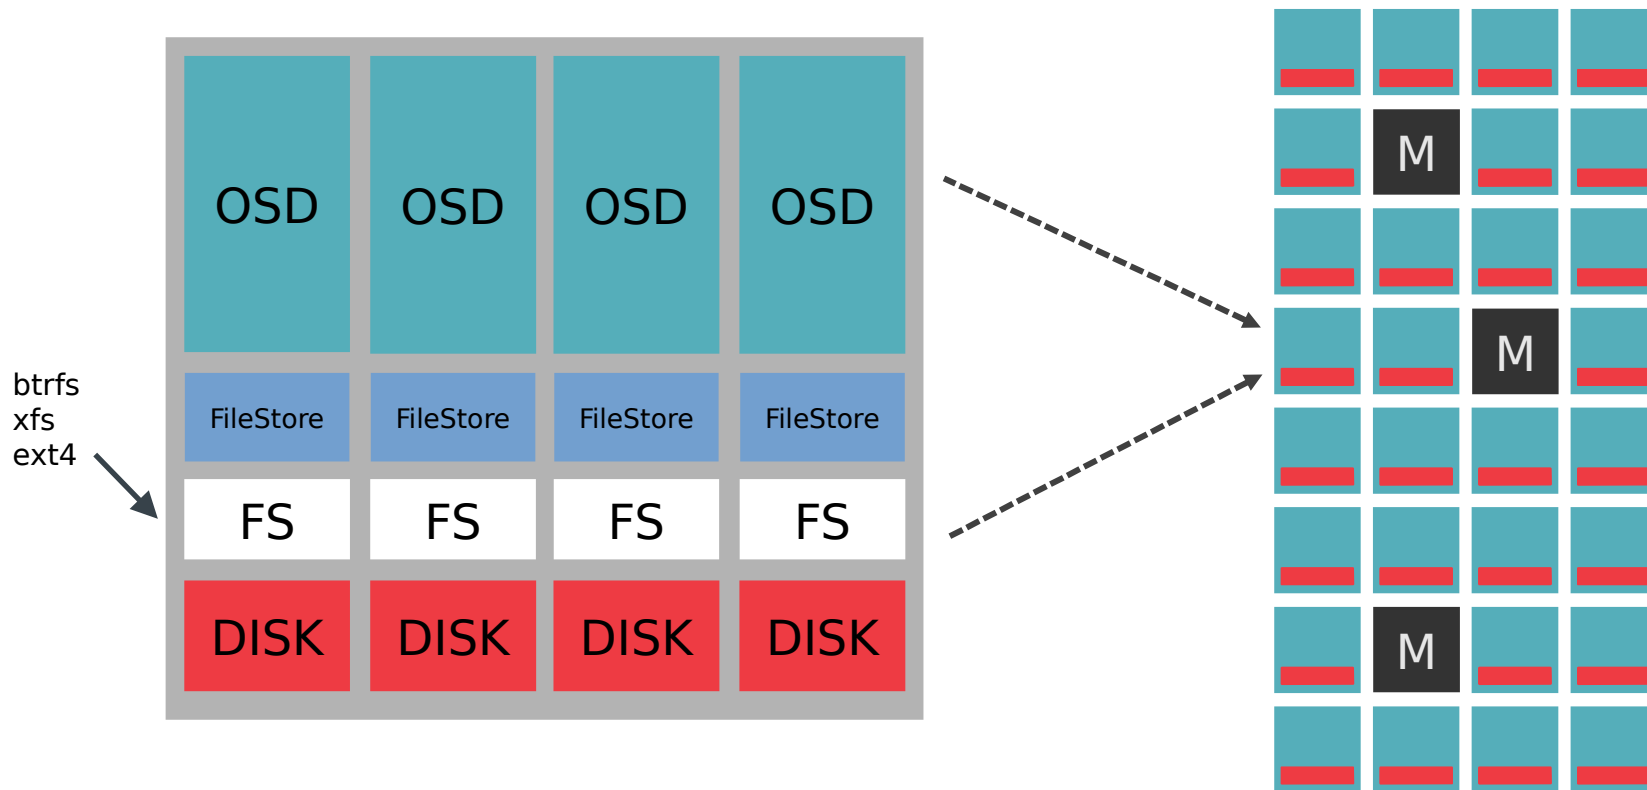
RADOS

A software-based, reliable, autonomous, distributed object store comprised of self-healing, self-managing, intelligent storage nodes and lightweight monitors

OBJECT STORAGE DAEMONS (OSDs)



OBJECT STORAGE DAEMONS (OSDS)



OBJECTSTORE AND DATA MODEL



- ObjectStore
 - abstract interface for storing local data
 - EBOFS, FileStore
- EBOFS
 - a user-space **e**xtent-**b**ased **o**bject **f**ile **s**ystem
 - deprecated in favor of FileStore on btrfs in 2009
- Object
 - data (file-like byte stream)
 - attributes (small key/value)
 - omap (unbounded key/value)
- Collection
 - placement group shard (slice of the RADOS pool)
 - sharded by 32-bit hash value
- All writes are transactions
 - **A**tomic + **C**onsistent + **D**urable
 - **I**solation provided by OSD

FILESTORE



- FileStore
 - PG = collection = directory
 - object = file
- Leveldb
 - large xattr spillover
 - object omap (key/value) data
- Originally just for development...
 - later, only supported backend (on XFS)
- `/var/lib/ceph/osd/ceph-123/`
 - `current/`
 - `meta/`
 - `osdmap123`
 - `osdmap124`
 - `0.1_head/`
 - `object1`
 - `object12`
 - `0.7_head/`
 - `object3`
 - `object5`
 - `0.a_head/`
 - `object4`
 - `object6`
 - `db/`
 - `<leveldb files>`

- ...but others are arbitrarily large/complex

1

POSIX FAILS: TRANSACTIONS



- Btrfs transaction hooks

```
/* trans start and trans end are dangerous, and only for
 * use by applications that know how to avoid the
 * resulting deadlocks
 */
#define BTRFS_IOC_TRANS_START  _IO(BTRFS_IOCTL_MAGIC, 6)
#define BTRFS_IOC_TRANS_END    _IO(BTRFS_IOCTL_MAGIC, 7)
```

- Writeback ordering

```
#define BTRFS_MOUNT_FLUSHONCOMMIT      (1 << 7)
```

- What if we hit an error? ceph-osd process dies?

```
#define BTRFS_MOUNT_WEDGEONTRANSABORT  (1 << ...)
```

- There is no rollback...

POSIX FAILS: TRANSACTIONS




- Write-ahead journal
 - serialize and journal every `ObjectStore::Transaction`
 - then write it to the file system
- Btrfs parallel journaling
 - periodic sync takes a snapshot
 - on restart, rollback, and replay journal against appropriate snapshot
- XFS/ext4 write-ahead journaling
 - periodic sync, then trim old journal entries
 - on restart, replay entire journal
 - lots of ugly hackery to deal with events that aren't idempotent
 - e.g., renames, collection delete + create, ...
- **full data journal → we double write everything → ~halve disk throughput**

POSIX FAILS: ENUMERATION



- Ceph objects are distributed by a 32-bit hash
- Enumeration is in hash order
 - scrubbing
 - “backfill” (data rebalancing, recovery)
 - enumeration via librados client API
- POSIX readdir is not well-ordered
- Need $O(1)$ “split” for a given shard/range
- Build directory tree by hash-value prefix
 - split any directory when size $> \sim 100$ files
 - merge when size $< \sim 50$ files
 - read entire directory, sort in-memory

```
...
DIR_A/
DIR_A/A03224D3_qwer
DIR_A/A247233E_zxcv
...
DIR_B/
DIR_B/DIR_8/
DIR_B/DIR_8/B823032D_foo
DIR_B/DIR_8/B8474342_bar
DIR_B/DIR_9/
DIR_B/DIR_9/B924273B_baz
DIR_B/DIR_A/
DIR_B/DIR_A/BA4328D2_asdf
...
```



NEWSTORE

NEW STORE GOALS

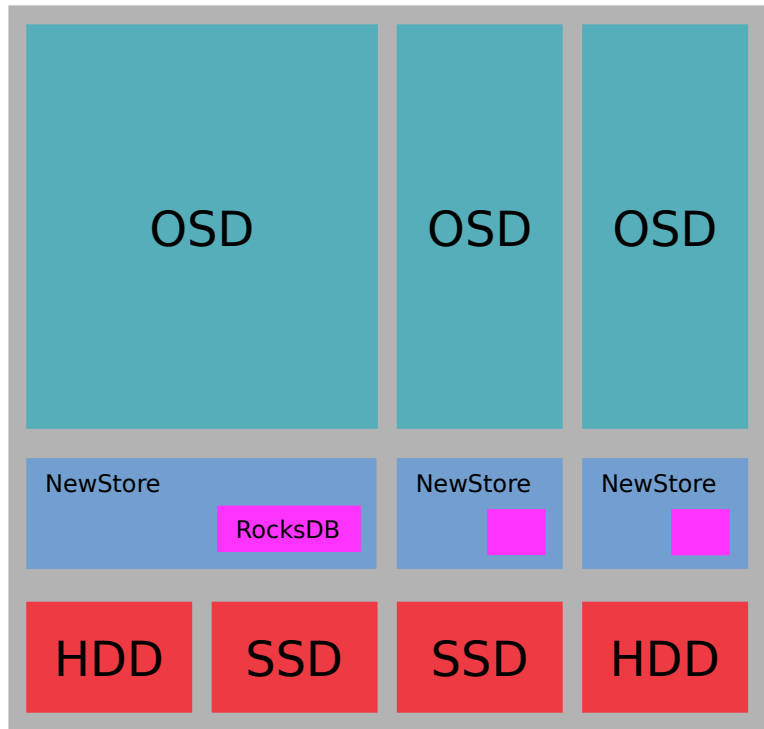


- More natural transaction atomicity
- Avoid double writes
- Efficient object enumeration
- Efficient clone operation
- Efficient splice (“move these bytes from object X to object Y”)
- Efficient IO pattern for HDDs, SSDs, NVMe
- Minimal locking, maximum parallelism (between PGs)
- Advanced features
 - full data and metadata checksums
 - compression

NEWSTORE – WE MANAGE NAMESPACE



- POSIX has the wrong metadata model for us
- Ordered key/value is perfect match
 - well-defined object name sort order
 - efficient enumeration and random lookup
- NewStore = rocksdb + object files
 - `/var/lib/ceph/osd/ceph-123/`
 - `db/`
 - `<rocksdb, leveldb, whatever>`
 - `blobs.1/`
 - `0`
 - `1`
 - `...`
 - `blobs.2/`
 - `100000`
 - `100001`
 - `...`



NEWSTORE FAIL: CONSISTENCY OVERHEAD




- RocksDB has a write-ahead log “journal”
 - XFS/ext4(/btrfs) have their own journal (tree-log)
 - Journal-on-journal has high overhead
 - each journal manages half of overall consistency, but incurs same overhead
- `write(2) + fsync(2)` to new blobs.2/10302
 - 1 write + flush to block device
 - 1 write + flush to XFS/ext4 journal
 - `write(2) + fsync(2)` on RocksDB log
 - 1 write + flush to block device
 - 1 write + flush to XFS/ext4 journal

NEWSTORE FAIL: ATOMICITY NEEDS WAL



- We can't overwrite a POSIX file as part of a atomic transaction
- Writing overwrite data to a new file means many files for each object
- Write-ahead logging?
 - put overwrite data in a “WAL” records in RocksDB
 - commit atomically with transaction
 - then overwrite original file data
 - but we're back to a double-write of overwrites...
- Performance sucks again
- Overwrites dominate RBD block workloads

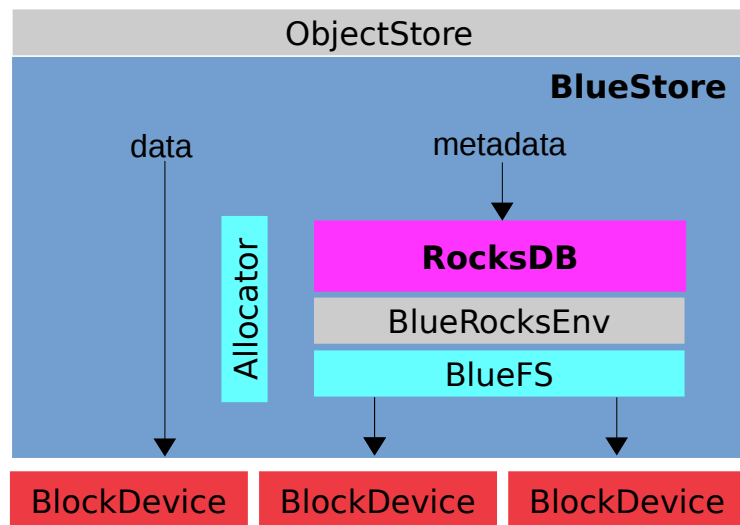


BLUESTORE

BLUESTORE




- BlueStore = **Block** + **NewStore**
 - consume raw block device(s)
 - key/value database (RocksDB) for metadata
 - data written directly to block device
 - pluggable block Allocator
- We must share the block device with RocksDB
 - implement our own rocksdb::Env
 - implement tiny “file system” BlueFS
 - make BlueStore and BlueFS share



ROCKSDB: BLUEROCKSENV + BLUEFS



- `class BlueRocksEnv : public rocksdb::EnvWrapper`
 - passes file IO operations to BlueFS
- BlueFS is a super-simple “file system”
 - all metadata loaded in RAM on start/mount
 - no need to store block free list
 - coarse allocation unit (1 MB blocks)
 - all metadata lives in written to a journal
 - journal rewritten/compacted when it gets large
- Map “directories” to different block devices
 - `db.wal/` - on NVRAM, NVMe, SSD
 - `db/` - level0 and hot SSTs on SSD
 - `db.slow/` - cold SSTs on HDD
- BlueStore periodically balances free space



ROCKSDB: JOURNAL RECYCLING




- Problem: 1 small (4 KB) Ceph write → 3-4 disk IOs!
 - BlueStore: write **4 KB of user data**
 - rocksdb: append record to WAL
 - **write update block** at end of log file
 - fsync: XFS/ext4/BlueFS journals **inode size/alloc update** to its journal
- fallocate(2) doesn't help
 - data blocks are not pre-zeroed; fsync still has to update alloc metadata
- rocksdb LogReader only understands two modes
 - read until end of file (need accurate file size)
 - read all valid records, then ignore zeros at end (need zeroed tail)

ROCKSDB: JOURNAL RECYCLING (2)



- Put old log files on recycle list (instead of deleting them)
- LogWriter
 - overwrite old log data with new log data
 - include log number in each record
- LogReader
 - stop replaying when we get garbage (bad CRC)
 - or when we get a valid CRC but record is from a previous log incarnation
- Now we get one log append → one IO!
- Upstream in RocksDB!
 - but missing a bug fix (PR #881)
- Works with normal file-based storage, or BlueFS



METADATA

BLUESTORE METADATA



- Partition namespace for different metadata
 - S* – “superblock” metadata for the entire store
 - B* – block allocation metadata
 - C* – collection name → cnode_t
 - O* – object name → onode_t or enode_t
 - L* – write-ahead log entries, promises of future IO
 - M* – omap (user key/value data, stored in objects)
 - V* – overlay object data (obsolete?)



- Per object metadata
 - Lives directly in key/value pair
 - Serializes to ~200 bytes
- Unique object id (like `ino_t`)
- Size in bytes
- Inline attributes (user attr data)
- Block pointers (user byte data)
 - Overlay metadata
- Omap prefix/ID (user k/v data)

```
struct bluestore_onode_t {  
    uint64_t nid;  
    uint64_t size;  
    map<string,bufferptr> attrs;  
    map<uint64_t,bluestore_extent_t> block_map;  
    map<uint64_t,bluestore_overlay_t> overlay_map;  
    uint64_t omap_head;  
};
```

```
struct bluestore_extent_t {  
    uint64_t offset;  
    uint32_t length;  
    uint32_t flags;  
};
```



- Collection metadata
 - Interval of object namespace

```
shard pool hash name bits
C<NOSHARD,12,3d321e00> "12.e123d3" = <25>
```

```
shard pool hash name snap gen
0<NOSHARD,12,3d321d88,foo,NOSNAP,NOGEN> = ...
0<NOSHARD,12,3d321d92,bar,NOSNAP,NOGEN> = ...
```

```
0<NOSHARD,12,3d321e02,baz,NOSNAP,NOGEN> = ...
0<NOSHARD,12,3d321e12,zip,NOSNAP,NOGEN> = ...
0<NOSHARD,12,3d321e12,dee,NOSNAP,NOGEN> = ...
```

```
0<NOSHARD,12,3d321e38,dah,NOSNAP,NOGEN> = ...
```

```
struct spg_t {
    uint64_t pool;
    uint32_t hash;
    shard_id_t shard;
};
```

```
struct bluestore_cnode_t {
    uint32_t bits;
};
```

- Nice properties
 - Ordered enumeration of objects
 - We can “split” collections by adjusting metadata



- Extent metadata
 - *Sometimes* we share blocks between objects (usually clones/snaps)
 - We need to reference count those extents
 - We still want to split collections and repartition extent metadata by hash

```
shard pool hash name snap gen
0<NOSHARD,12,3d321d92,bar,NOSNAP,NOGEN> = onode
0<NOSHARD,12,3d321e02> = enode
0<NOSHARD,12,3d321e02,baz,NOSNAP,NOGEN> = onode
0<NOSHARD,12,3d321e12> = enode
0<NOSHARD,12,3d321e12,zip,NOSNAP,NOGEN> = onode
0<NOSHARD,12,3d321e12,dee,NOSNAP,NOGEN> = onode
0<NOSHARD,12,3d321e38,dah,NOSNAP,NOGEN> = onode
```

```
struct bluestore_enode_t {
    struct record_t {
        uint32_t length;
        uint32_t refs;
    };
    map<uint64_t,record_t> ref_map;

    void add(uint64_t offset, uint32_t len,
             unsigned ref=2);
    void get(uint64_t offset, uint32_t len);
    void put(uint64_t offset, uint32_t len,
             vector<bluestore_extent_t> *release);
};
```



DATA PATH

DATA PATH BASICS



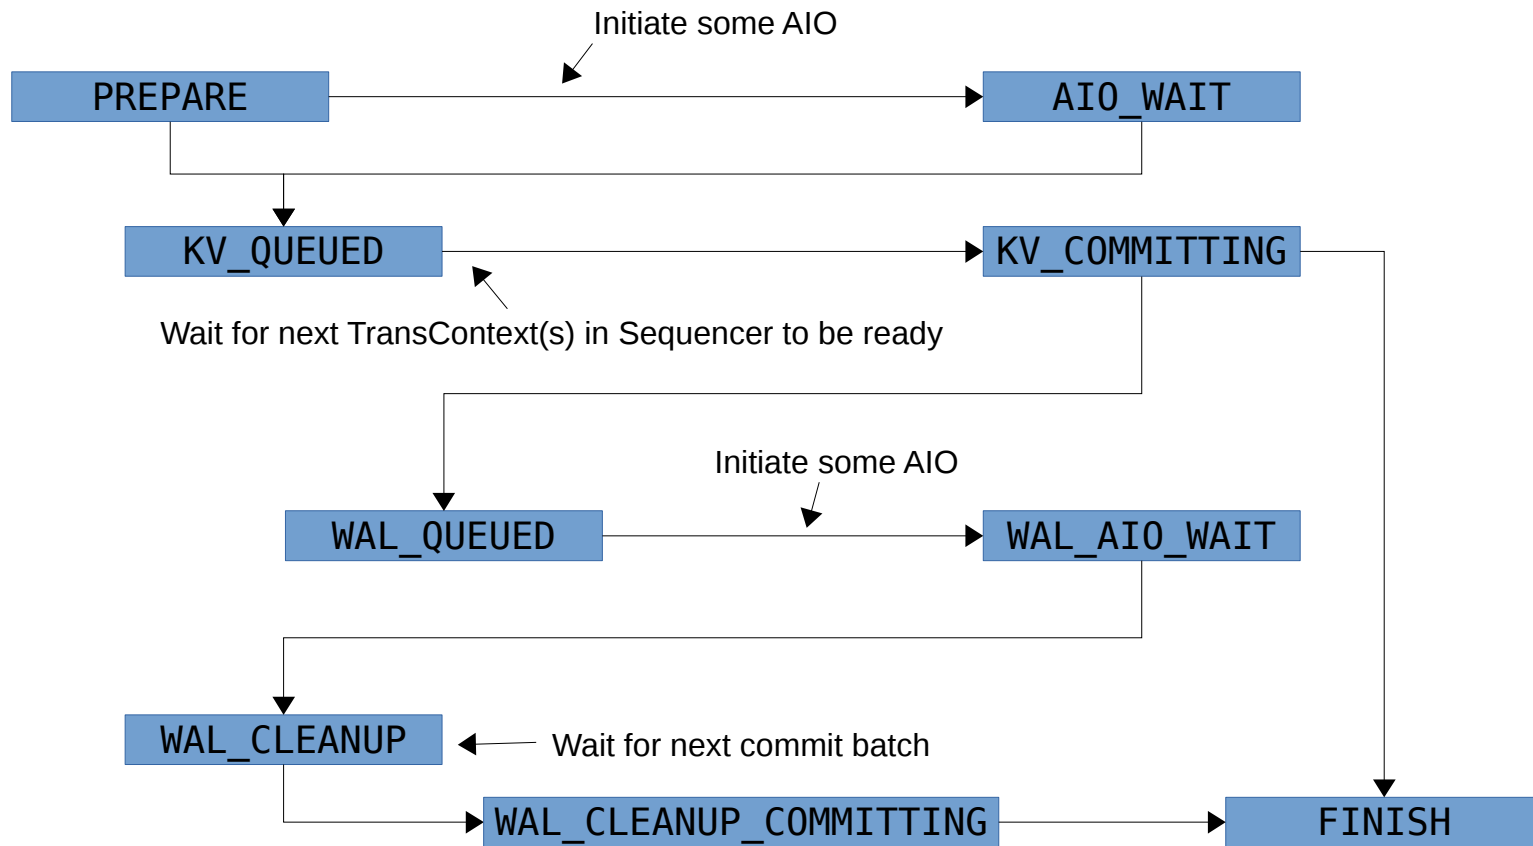
Terms

- Sequencer
 - An independent, totally ordered queue of transactions
 - One per PG
- TransContext
 - State describing an executing transaction

Two ways to write

- New allocation
 - Any write larger than **min_alloc_size** goes to a new, unused extent on disk
 - Once that IO completes, we commit the transaction
- WAL (write-ahead-logged)
 - Commit temporary promise to (over)write data with transaction
 - includes data!
 - Do async overwrite
 - Clean up temporary k/v pair

TRANSCONTEXT STATE MACHINE



Sequencer queue

PREPARE
AIO_WAIT
KV_QUEUED
AIO_WAIT
KV_COMMITTING
KV_COMMITTING
KV_COMMITTING
FINISH
WAL_QUEUED
FINISH
WAL_AIO_WAIT
WAL_CLEANUP
WAL_CLEANUP
FINISH
FINISH

AIO, O_DIRECT, AND CACHING



- From open(2) man page

Applications should avoid mixing O_DIRECT and normal I/O to the same file, and especially to overlapping byte regions in the same file.

- By default, all IO is AIO + O_DIRECT
 - but sometimes we want to cache (e.g., POSIX_FADV_WILLNEED)
- BlueStore mixes direct and buffered IO
 - O_DIRECT read(2) and write(2) invalidate and/or flush pages... racily
 - we avoid mixing them on the same pages
 - disable readahead: posix_fadvise(fd, 0, 0, POSIX_FADV_RANDOM)
- But it's not quite enough...
 - moving to fully user-space cache

BLOCK FREE LIST



- FreelistManager
 - persist list of free extents to key/value store
 - prepare incremental updates for allocate or release
- Initial implementation
 - `<offset> = <length>`
 - keep in-memory copy
 - enforce an ordering on commits
 - `del 1600=100000`
 - `put 1700=99990`
 - small initial memory footprint, very expensive when fragmented
- New approach
 - `<offset> = <region bitmap>`
 - where region is N blocks (1024?)
 - no in-memory state
 - use k/v **merge** operator to XOR allocation or release
 - `merge 10=0000000011`
 - `merge 20=1110000000`
 - RocksDB log-structured-merge tree coalesces keys during compaction

BLOCK ALLOCATOR



- Allocator
 - abstract interface to allocate new space
- StupidAllocator
 - bin free extents by size (powers of 2)
 - choose sufficiently large extent closest to hint
 - highly variable memory usage
 - btree of free extents
 - implemented, works
 - based on ancient ebfs policy
- BitmapAllocator
 - hierarchy of indexes
 - L1: 2 bits = 2^6 blocks
 - L2: 2 bits = 2^{12} blocks
 - ...
 - 00 = all free, 11 = all used,
 - 01 = mix
 - fixed memory consumption
 - ~35 MB RAM per TB



- Let's support them natively!
- 256 MB zones / bands
 - must be written sequentially, but not all at once
 - libzbc supports ZAC and ZBC HDDs
 - host-managed or host-aware
- SMRAAllocator
 - write pointer per zone
 - used + free counters per zone
 - Bonus: almost no memory!
- IO ordering
 - must ensure allocated writes reach disk in order
- Cleaning
 - store k/v hints
 - zone → object hash
 - pick emptiest closed zone, scan hints, move objects that are still there
 - opportunistically rewrite objects we read if the zone is flagged for cleaning soon



FANCY STUFF

WE WANT FANCY STUFF



Full data checksums

- We scrub... periodically
- We want to validate checksum on **every** read

Compression

- 3x replication is expensive
- Any scale-out cluster is expensive

WE WANT FANCY STUFF



Full data checksums

- We scrub... periodically
- We want to validate checksum on **every** read
- More data with extent pointer
 - 4KB for 32-bit csum per 4KB block
 - bigger onode: 300 → 4396 bytes
 - larger csum blocks?

Compression

- 3x replication is expensive
- Any scale-out cluster is expensive
- Need largish extents to get compression benefit (64 KB, 128 KB)
 - overwrites need to do read/modify/write

INDIRECT EXTENT STRUCTURES



onode_t

```
map<uint64_t,bluestore_lextent_t> extent_map;
```

```
struct bluestore_lextent_t {  
    uint64_t blob_id;  
    uint64_t b_off, b_len;  
};
```

onode may reference a piece of a blob,
or multiple pieces

ref counted (when in enode)

csum block size can vary

onode_t or enode_t

```
map<uint64_t,bluestore_blob_t> blob_map;
```

```
struct bluestore_blob_t {  
    vector<bluestore_pextent_t> extents;  
    uint32_t logical_length; ///< uncompressed length  
    uint32_t flags;          ///< FLAGS_*  
    uint16_t num_refs;  
    uint8_t csum_type;        ///< CSUM_*  
    uint8_t csum_block_order;  
    vector<char> csum_data;  
};
```

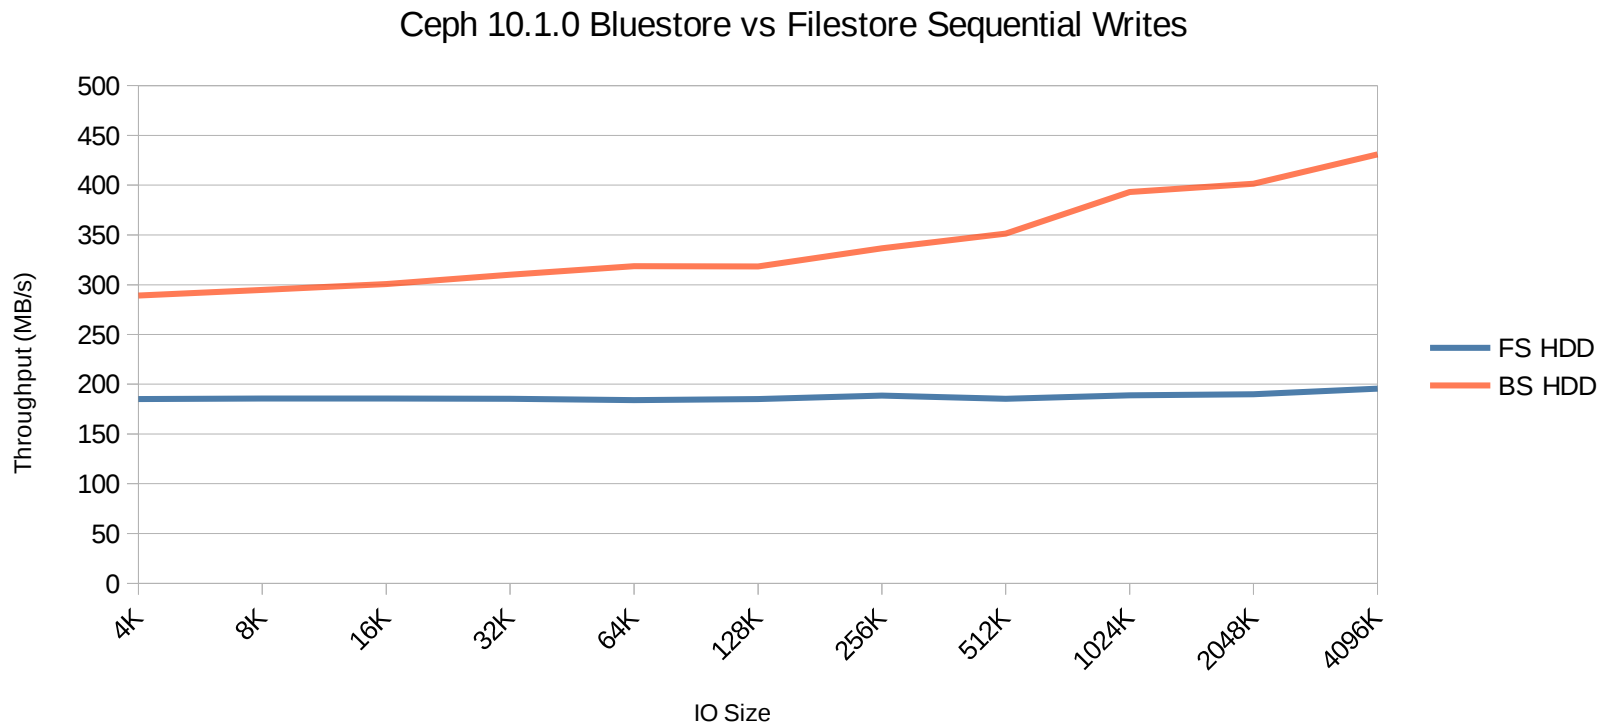
```
struct bluestore_pextent_t {  
    uint64_t offset, length;  
};
```

data on block device



PERFORMANCE


HDD: SEQUENTIAL WRITE




HDD: RANDOM WRITE



Ceph 10.1.0 Bluestore vs Filestore Random Writes




Ceph 10.1.0 Bluestore vs Filestore Random Writes



HDD: SEQUENTIAL READ



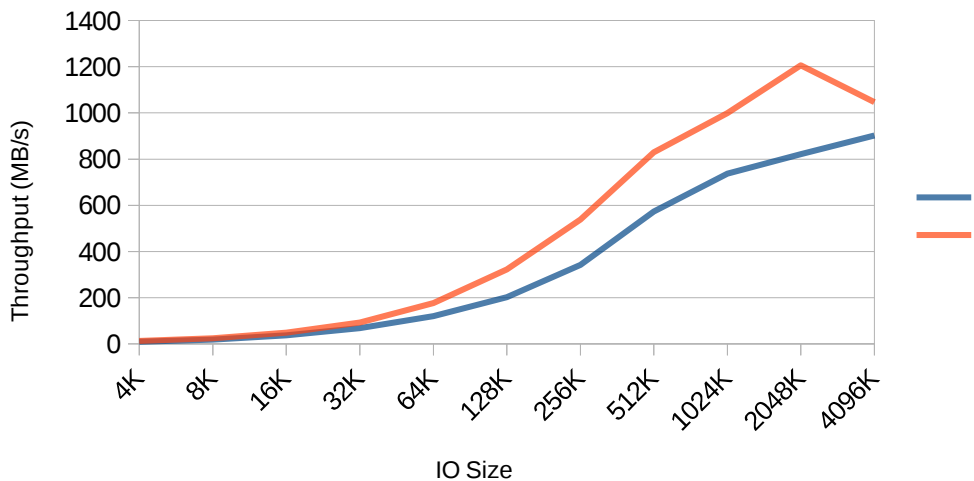
Ceph 10.1.0 Bluestore vs Filestore Sequential Reads



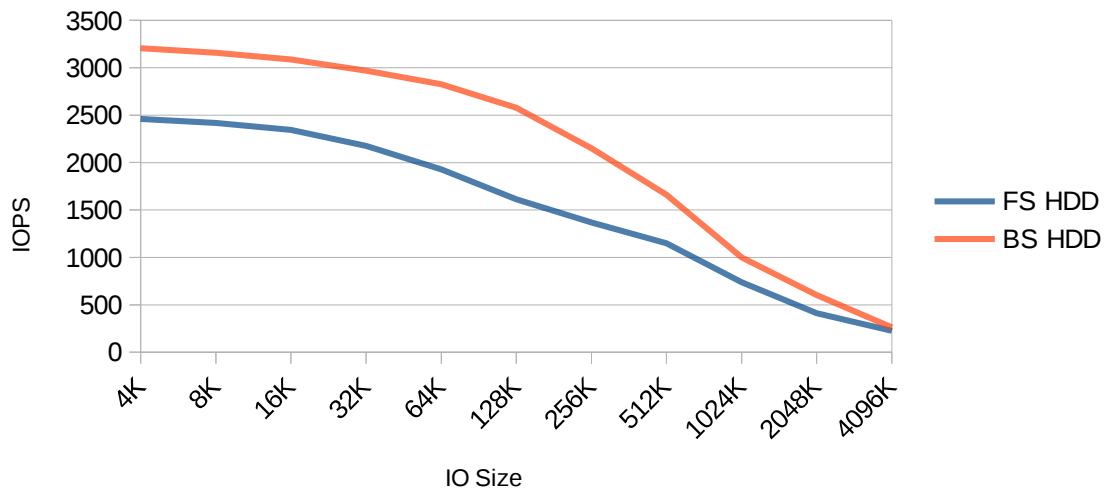
HDD: RANDOM READ



Ceph 10.1.0 Bluestore vs Filestore Random Reads



Ceph 10.1.0 Bluestore vs Filestore Random Reads



SSD AND NVME?



- NVMe journal
 - random writes ~2x faster
 - some testing anomalies (problem with test rig kernel?)
- SSD only
 - similar to HDD result
 - small write benefit is more pronounced
- NVMe only
 - more testing anomalies on test rig.. WIP



- Experimental backend in Jewel v10.2.z (just released)
 - enable experimental unrecoverable data corrupting features = bluestore rocksdb
 - ceph-disk --bluestore DEV
 - no multi-device magic provisioning just yet
- The goal...
 - stable in Kraken (Fall '16)?
 - default in Luminous (Spring '17)?

SUMMARY



- Ceph is great
- POSIX was poor choice for storing objects
- Our new BlueStore backend is awesome
- RocksDB rocks and was easy to embed
- Log recycling speeds up commits (now upstream)
- Delayed merge will help too (coming soon)

THANK YOU!

Sage Weil
CEPH PRINCIPAL ARCHITECT



sage@redhat.com



[@liewegas](https://twitter.com/liewegas)



ceph