

Red Hat Ceph Performance & Sizing Guide

Jose De la Rosa

September 7, 2016



Agenda

1. Ceph Overview
 - a. Architecture
 - b. Approach to storing data
2. Test methodology
 - a. What was tested
 - b. How it was tested & measured
3. Results
 - a. Key findings
 - b. Recommendations



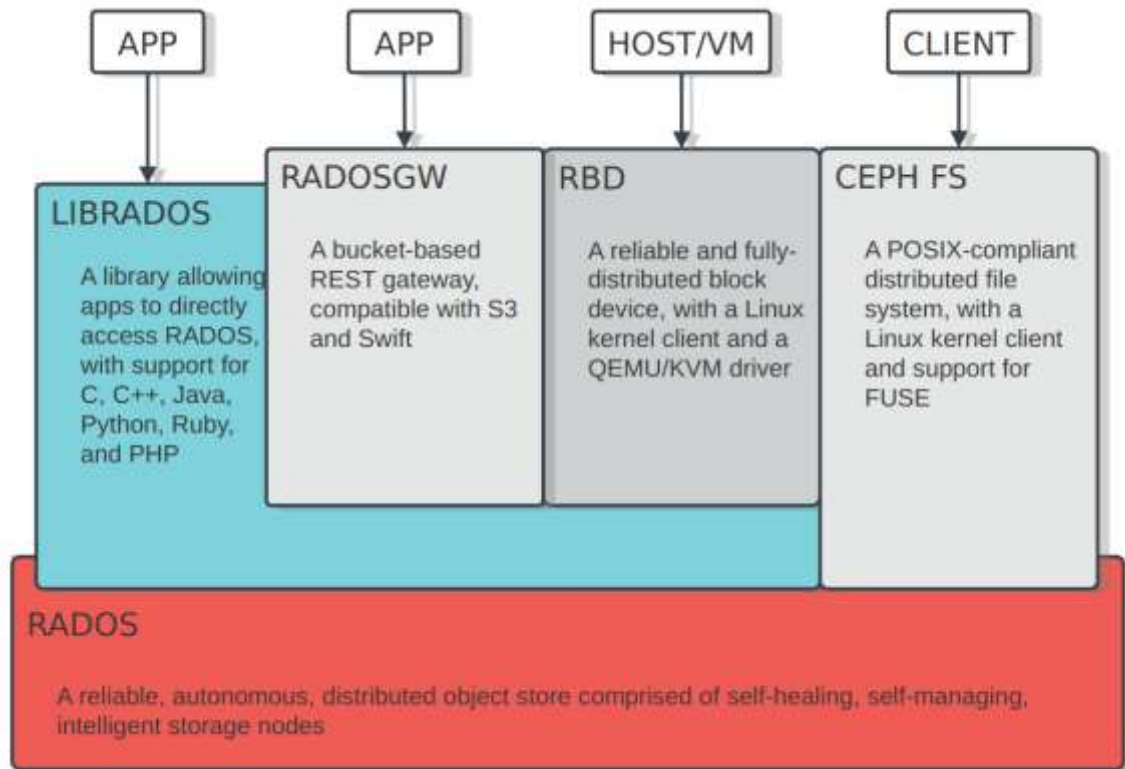
Ceph Overview

Ceph Overview

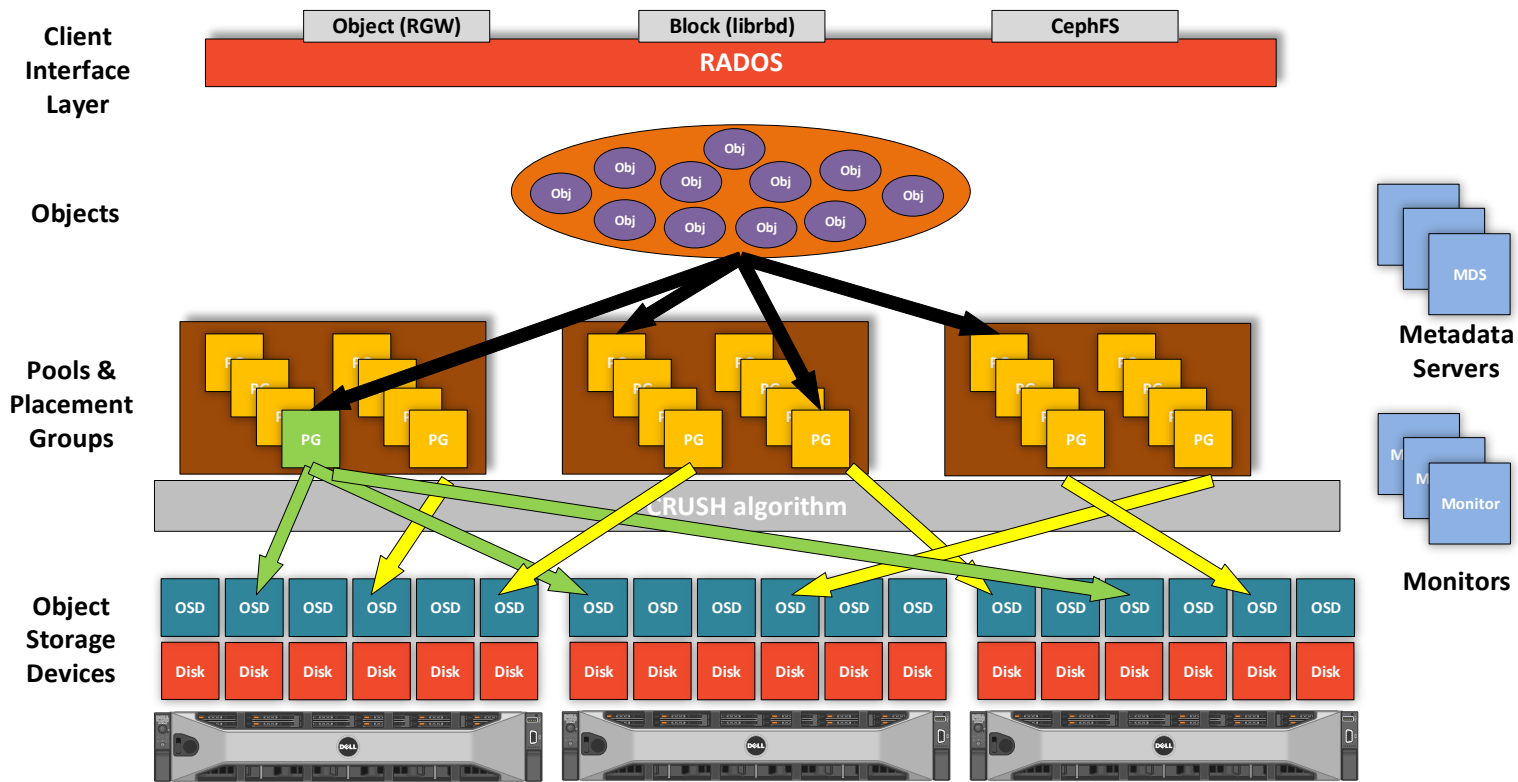
1. Open source software defined storage: logical storage services and capabilities are abstracted from the underlying physical storage systems.
2. Provides object, block and file system storage. All data is stored as objects.
3. Massively scalable to thousands of storage nodes.
4. Self-healing with no single point of failure: If a node fails, it is automatically detected and data rebalances to ensure availability.



Client access



Architecture




Storage Protection Method

3x Replication

Data copied 3 times


Actual disk usage = 33%




Erasure Coded (3+2)

Data split into 3 + 2 disks used for parity

Actual disk usage = 60%



Reading and Writing Data



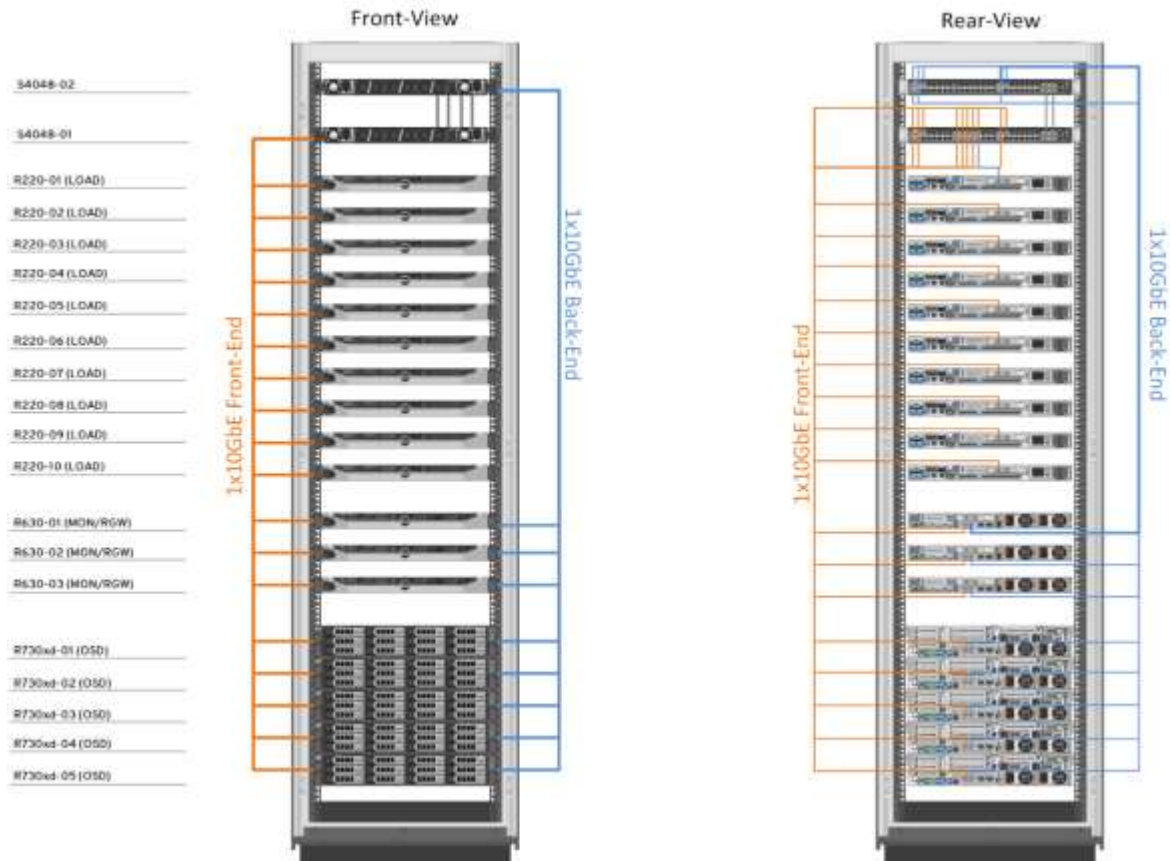
Test Setup & Methodology

Server configuration

| Testbed Details | | | |
|-----------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-----------------------------------------------------------------------|----------------------------------|
| Ceph tier | Storage Nodes (5) | Monitors (3) | Clients (10) |
| Platform | Dell PowerEdge R730xd | Dell PowerEdge R630 | Dell PowerEdge R220 |
| CPU | 2x Intel Xeon E5-2630 v3 2.4GHz | 2x Intel Xeon E5-2650 v3 2.3 GHz | 1x Intel Celeron G1820 2.7 GHz |
| Memory | 4x 16 GB 1866 MHz DDR4 | 8x 16 GB 2133MHz DDR4 | 4x 4 GB 1600 MHz DDR3 |
| Network | 1x Intel X520/2P I350 LOM | 1x Intel X520/2P I350 LOM | 1x Intel X520/2P I350 |
| Storage | PERC H730 Mini / 1 GB Cache Up to 16x: SEAGATE 4 TB SAS (ST4000NM0005) Up to 3x: Intel DC S3700 SSD 200 GB SATA (SSDSC2BA20) 1x Intel DC P3700 SSD 800 GB NVMe | PERC H730 Mini / 1 GB Cache 6x SEAGATE 500 GB SAS (ST9500620S) | 1x Toshiba 50 GB SATA (DT01ACA0) |



Network Topology



Configuration Guidance

1. General rules of thumb
 - 1 Core-GHz per OSD
 - SATA/SAS SSD to HDD Ratio: 1:4 - 1:5
 - NVME SSD to HDD Ratio: 1:17-1:18
 - 16GB RAM Baseline + 2-3GB per OSD
2. More details at <https://www.redhat.com/en/resources/red-hat-ceph-storage-hardware-configuration-guide>



Storage node configurations tested

| OSD to Journal Ratio [drives] | 12+3 | 16+0 | 16+1 |
|-------------------------------|-------------------------------------------|-------------------------------------------|-------------------------------------------|
| OSD node configuration | 12+3 | 16+0 | 16+1 |
| HDDs | 12 | 16 | 16 |
| HDD RAID mode | Single-disk RAID0 | Single-disk RAID0 | Single-disk RAID0 / HBA mode |
| SATA SSDs | 3 | 0 | 0 |
| SSD RAID mode | JBOD | JBOD | JBOD |
| NVMe SSDs | 0 | 0 | 1 |
| Network | 1x 10 GbE Front-End 1x 10 GbE Back-End | 1x 10 GbE Front-End 1x 10 GbE Back-End | 1x 10 GbE Front-End 1x 10 GbE Back-End |

Benchmarking with CBT

1. For benchmark automation, the open source utility Ceph Benchmarking Tool (CBT) was used.
2. It supports different drivers for examining different layers of the storage stack:
 - **radosbench - uses librados API (used in this study)**
 - librbdfio – test block storage without KVM/QEMU instances
 - kvmrbdfio – test block volumes attached to KVM/QEMU instances
3. Available at <https://github.com/ceph/cbt>



Factors that influence performance

1. Device used for journaling (SSD vs. HDD)
2. RAID0 vs. pass-through (HBA) mode
3. Number of clients (single stream vs. parallel access)
4. Data-protection mode (3-way replication vs. erasure coding)



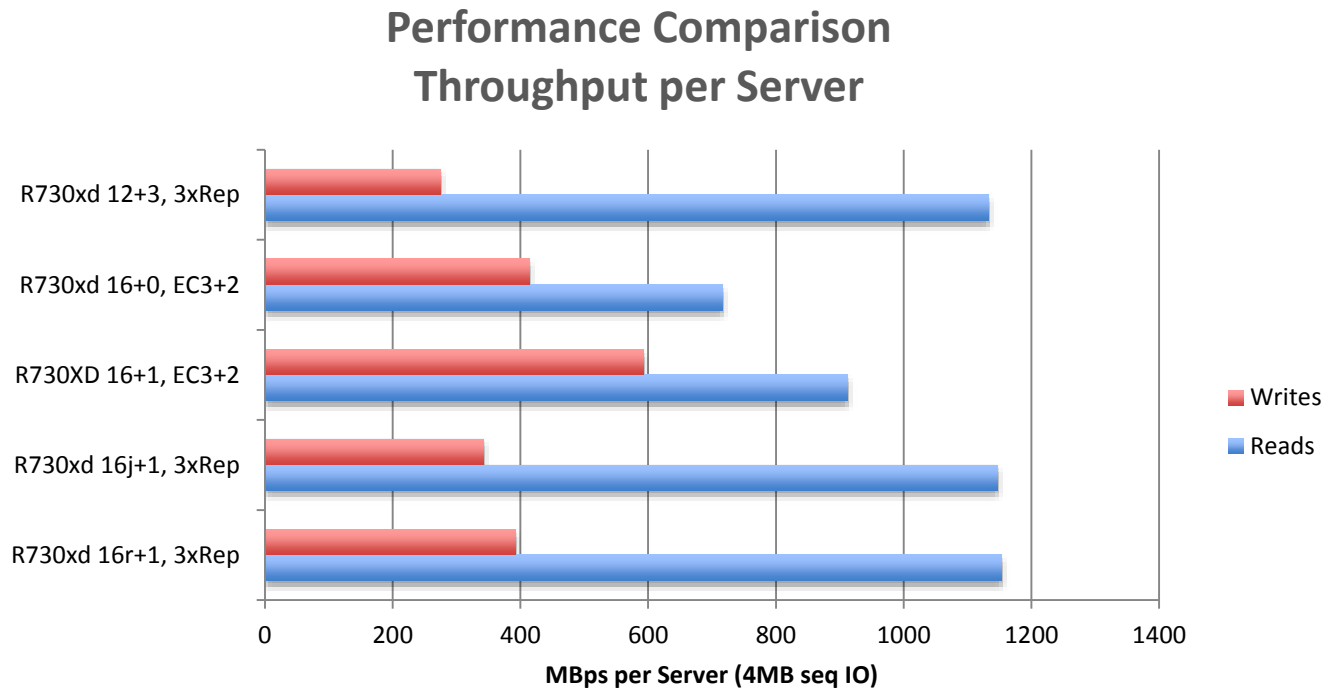
Test Matrix

| Server configuration | PowerEdge R730xd 12+3, 3xRep | PowerEdge R730xd 16+0, EC3+2 | PowerEdge R730xd 16r+1, 3xRep | PowerEdge R730xd 16+1, EC 3+2 | PowerEdge R730xd 16j+1, 3xRep |
|-------------------------------------------|---------------------------------|---------------------------------|-----------------------------------|-----------------------------------|-------------------------------------|
| OS disk | 2x 500 GB 2.5" | 2x 500 GB 2.5" | 2x 500 GB 2.5" | 2x 500 GB 2.5" | 2x 500 GB 2.5" |
| Data disk type | HDD 7.2K SAS 12Gbps, 4TB | HDD 7.2K SAS 12Gbps, 4TB | HDD 7.2K SAS 12Gbps, 4TB | HDD 7.2K SAS 12Gbps, 4TB | HDD 7.2K SAS 12Gbps, 4TB |
| HDD quantity | 12 | 16 | 16 | 16 | 16 |
| Number of Ceph write journal devices | 3 | 0 | 1 | 1 | 1 |
| Ceph write journal device type | Intel SATA SSD S3710 (6Gb/s) | n/a | Intel P3700 PCIe NVMe HHHL AIC | Intel P3700 PCIe NVMe HHHL AIC | Intel P3700 PCIe NVMe HHHL AIC |
| Ceph write journal device size (GB) | 200 | 0 | 800 | 800 | 800 |
| Controller model | PERC H730, 1 GB Cache | PERC H730, 1 GB Cache | PERC H730, 1 GB Cache | PERC H730, 1 GB Cache | PERC H730, 1 GB Cache |
| PERC Controller configuration for HDDs | RAID | RAID | RAID | RAID | JBOD (PERC pass-through mode) |
| Raw capacity for Ceph OSDs (TB) | 48 | 64 | 64 | 64 | 64 |




Benchmark Test Results

Throughput / server



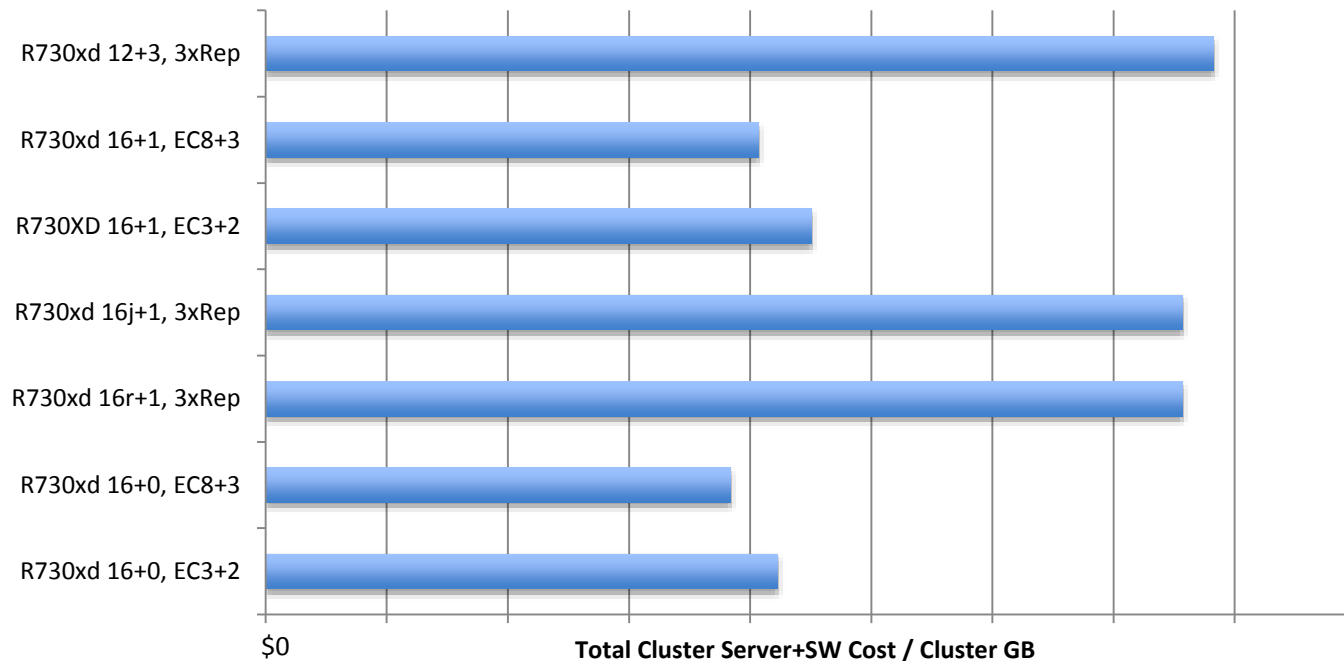
Overall Solution Price/Performance

Solution Price/Performance Comparison 500TB Usable Cluster (less \$ per MBps is better)

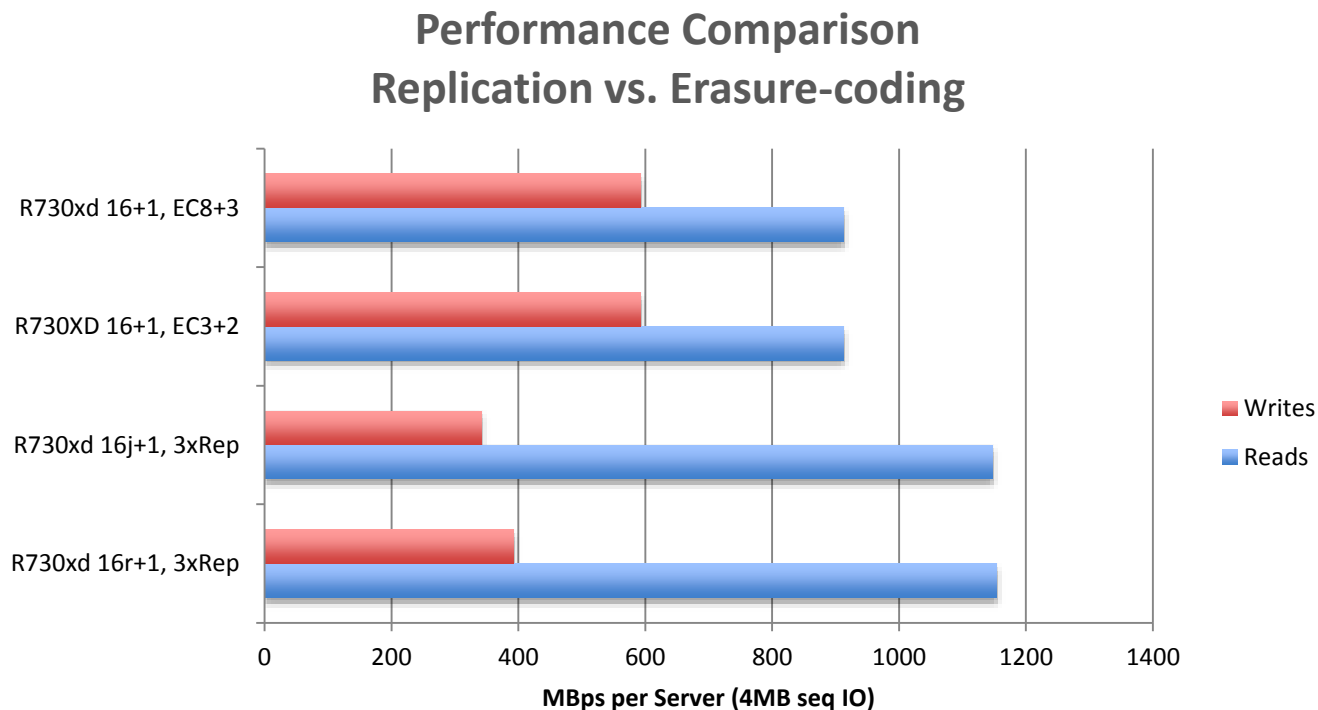


Overall Solution Price Capacity

Solution Price/Capacity Comparison
(less \$ per GB is better)




Replication vs. Erasure Coding



JBOD vs. RAID0

Performance Comparison JBOD vs. RAID0 Config



Performance conclusions

1. **Replication mode** yielded better performance for **read operations** and the **erasure-coded** mode proved better for **write operations**.
2. The PowerEdge **R730xd 16+1 3x replication** configuration yielded optimal price for **read-write throughput-oriented workloads**.
3. The PowerEdge **R730xd 12+3 3x replication** configuration yielded optimal price for **read-only throughput-oriented workloads**.
4. The PowerEdge **R730xd 16+1 erasure-coded** configuration proved to be the choice for **write-heavy operations**.
5. When used with Ceph Storage, Dell & Red Hat recommend the usage of single-drive RAID0 mode on PowerEdge R730xd with PERC H730.



Sizing Recommendations

| Storage Capacity | Extra Small | Small | Medium |
|-------------------------|----------------------|----------------------|----------------------|
| Cluster Capacity | 100 TB+ | 500 TB+ | 1 PB+ |
| Throughput-Optimized | >4x R730xd (8U) | >8x R730xd (16U) | NA |
| | 1x server/2U chassis | 1x server/2U chassis | |
| | 16x 6 TB HDD | 16x 6 TB HDD | |
| | 1x 800 GB NVMe SSD | 1x 800 GB NVMe SSD | |
| | 2x 10 GbE | 2x 10 GbE | |
| | 3x Replication | 3x Replication | |
| Cost/Capacity-Optimized | NA | NA | >15x R730xd (30U) |
| | | | 1x server/2U chassis |
| | | | 16x 8 TB HDD |
| | | | 1x HHHL AIC SSD |
| | | | 2x 10 GbE |
| | | | 8:3 Erasure-coding |

Observations

1. Obey SSD to HDD ratio
2. Hardware matters, look at RAID controllers if you use HDDs
3. Don't use RAID controllers on SSDs
4. SSD sequential write bandwidth becomes a bottleneck
5. Random workloads should go on Flash-only
6. 10GbE Bonding not necessary with ≤ 16 drives



Recommended Reading

Dell PowerEdge R730xd Red Hat Ceph Storage Performance and Sizing Guide

http://en.community.dell.com/techcenter/cloud/m/dell_cloud_resources/20442913

Enterprise Ceph: Every Way, Your Way

https://www.redhat.com/files/summit/session-assets/2016/SS88828-enterprise-ceph_every-way-your-way.pdf

