



managing a distributed  
storage system at scale

sage weil – inktank  
LISA – 12.12.12

# outline

- the scale-out world
- what is it, what it's for, how it works
- how you can use it
  - librados
  - radosgw
  - RBD, the ceph block device
  - distributed file system
- how do you deploy and manage it
- roadmap
- why we do this, who we are

why should you care about another  
storage system?

# the old reality

- direct-attached storage
  - raw disks
  - RAID controller
- network-attached storage
  - a few large arrays
  - NFS, CIFS, maybe iSCSI
- SAN
  - expensive enterprise array
- scalable capacity and performance, to a point
- poor management scalability

# new user demands

- “cloud”
  - many disk volumes, with automated provisioning
  - larger data sets
    - RESTful object storage (e.g., S3, Swift)
    - NoSQL distributed key/value stores (Cassandra, Riak)
- “big data”
  - distributed processing
  - “unstructured data”
- ...and the more traditional workloads

# requirements

- diverse use-cases
  - object storage
  - block devices (for VMs) with snapshots, cloning
  - shared file system with POSIX, coherent caches
  - structured data... files, block devices, or objects?
- **scale**
  - terabytes, petabytes, exabytes
  - heterogeneous hardware
  - reliability and fault tolerance

# cost

- near-linear function of size or performance
- incremental expansion
  - no fork-lift upgrades
- no vendor lock-in
  - choice of hardware
  - choice of software
- open

# time

- ease of administration
- no manual data migration, load balancing
- painless scaling
  - expansion **and** contraction
  - seamless migration
- devops friendly



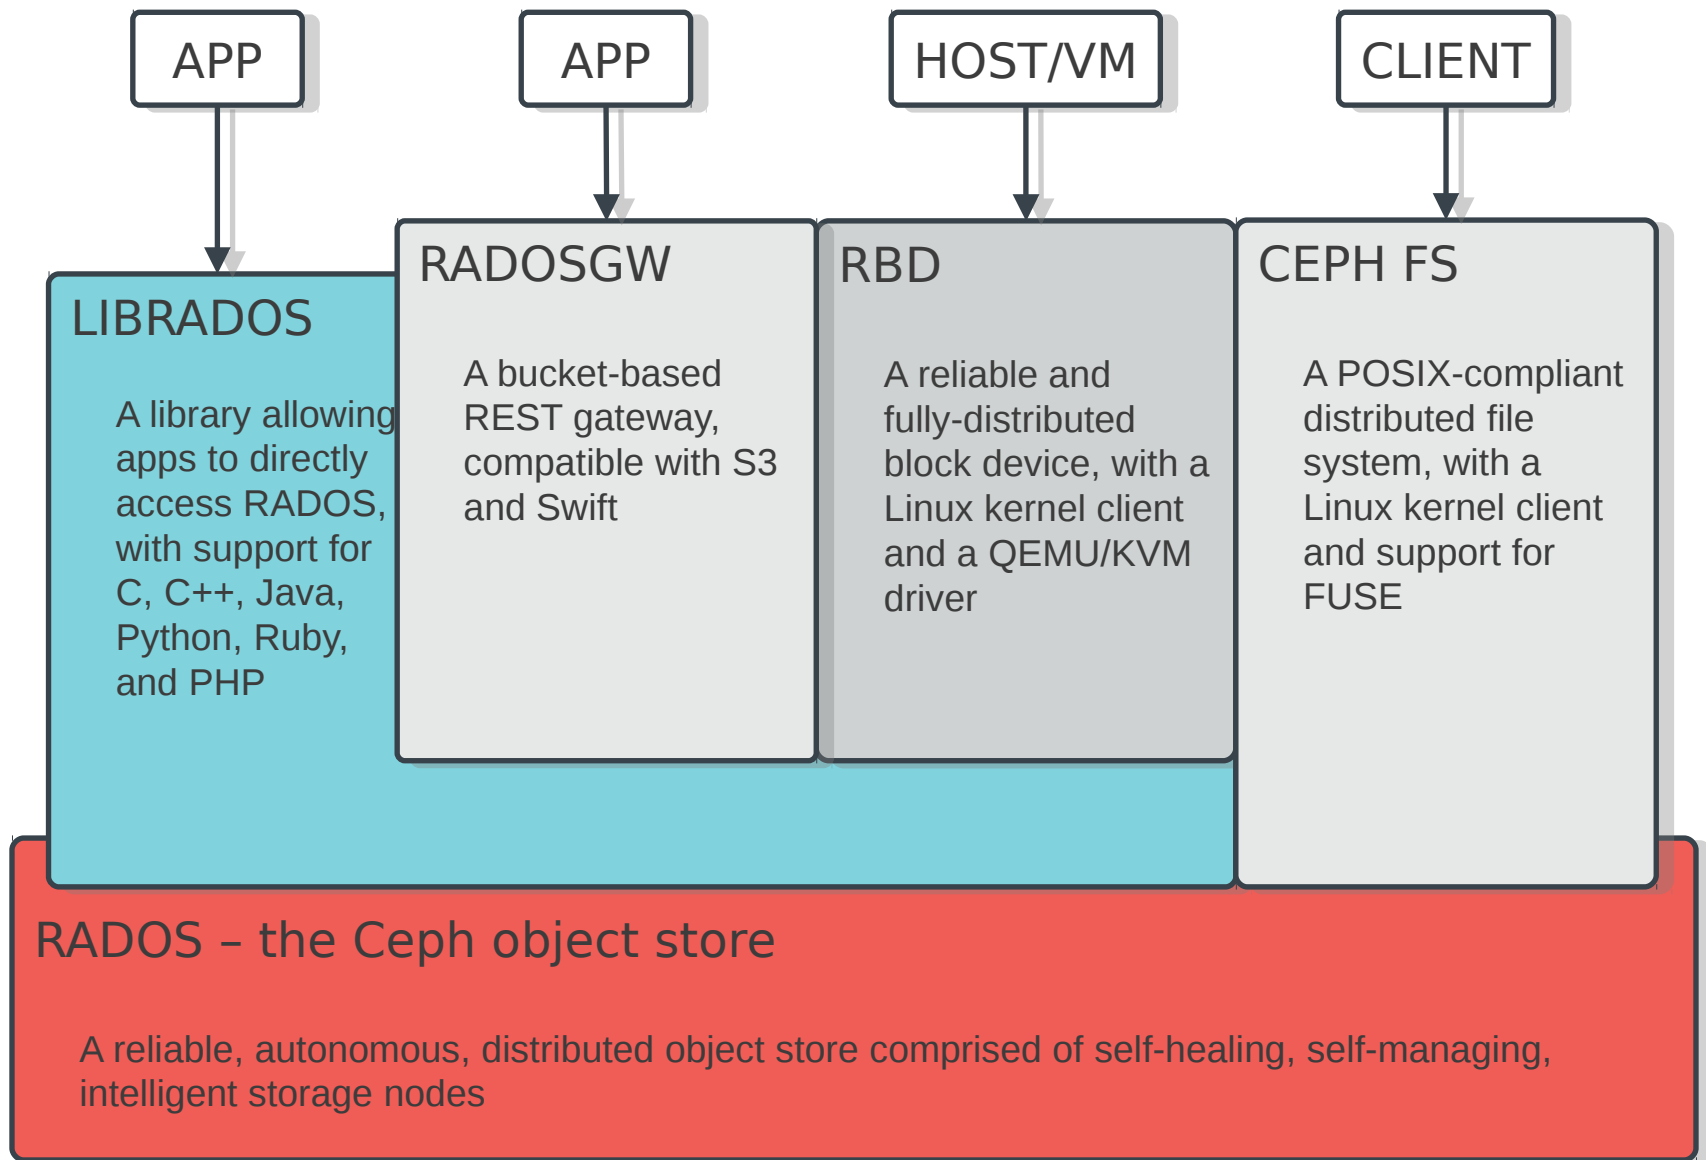
# being a good devops citizen

- seamless integration with infrastructure tools
  - Chef, Juju, Puppet, Crowbar, or home-grown
  - collectd, statsd, graphite, nagios
  - OpenStack, CloudStack
- simple processes for
  - expanding or contracting the storage cluster
  - responding to failure scenarios: initial healing, or mop-up
- evolution of “Unix philosophy”
  - everything is just daemon
  - CLI, REST APIs, JSON
  - solve one problem, and solve it well

what is **ceph**?

# unified storage system

- objects
  - native API
  - RESTful
- block
  - thin provisioning, snapshots, cloning
- file
  - strong consistency, snapshots



# distributed storage system

- data center scale
  - 10s to 10,000s of machines
  - terabytes to exabytes
- fault tolerant
  - no single point of failure
  - commodity hardware
- self-managing, self-healing

# ceph object model

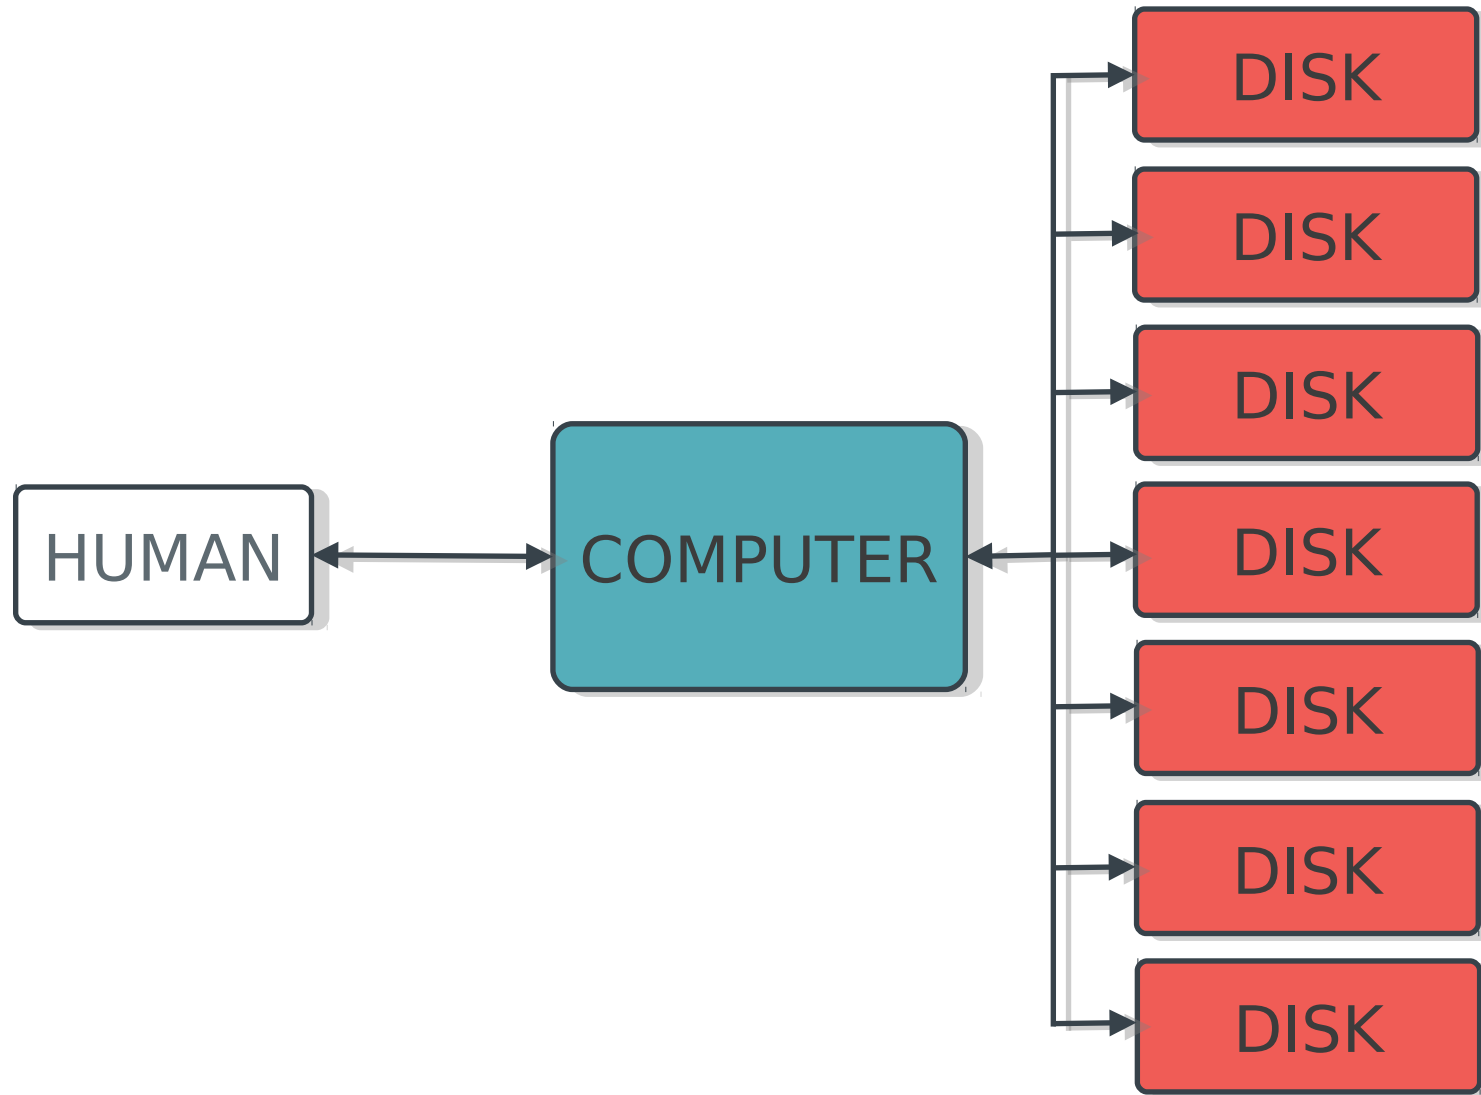
- pools
  - 1s to 100s
  - independent object namespaces or collections
  - replication level, placement policy
- objects
  - bazillions
  - blob of data (bytes to gigabytes)
  - attributes (e.g., “version=12”; bytes to kilobytes)
  - key/value bundle (bytes to gigabytes)

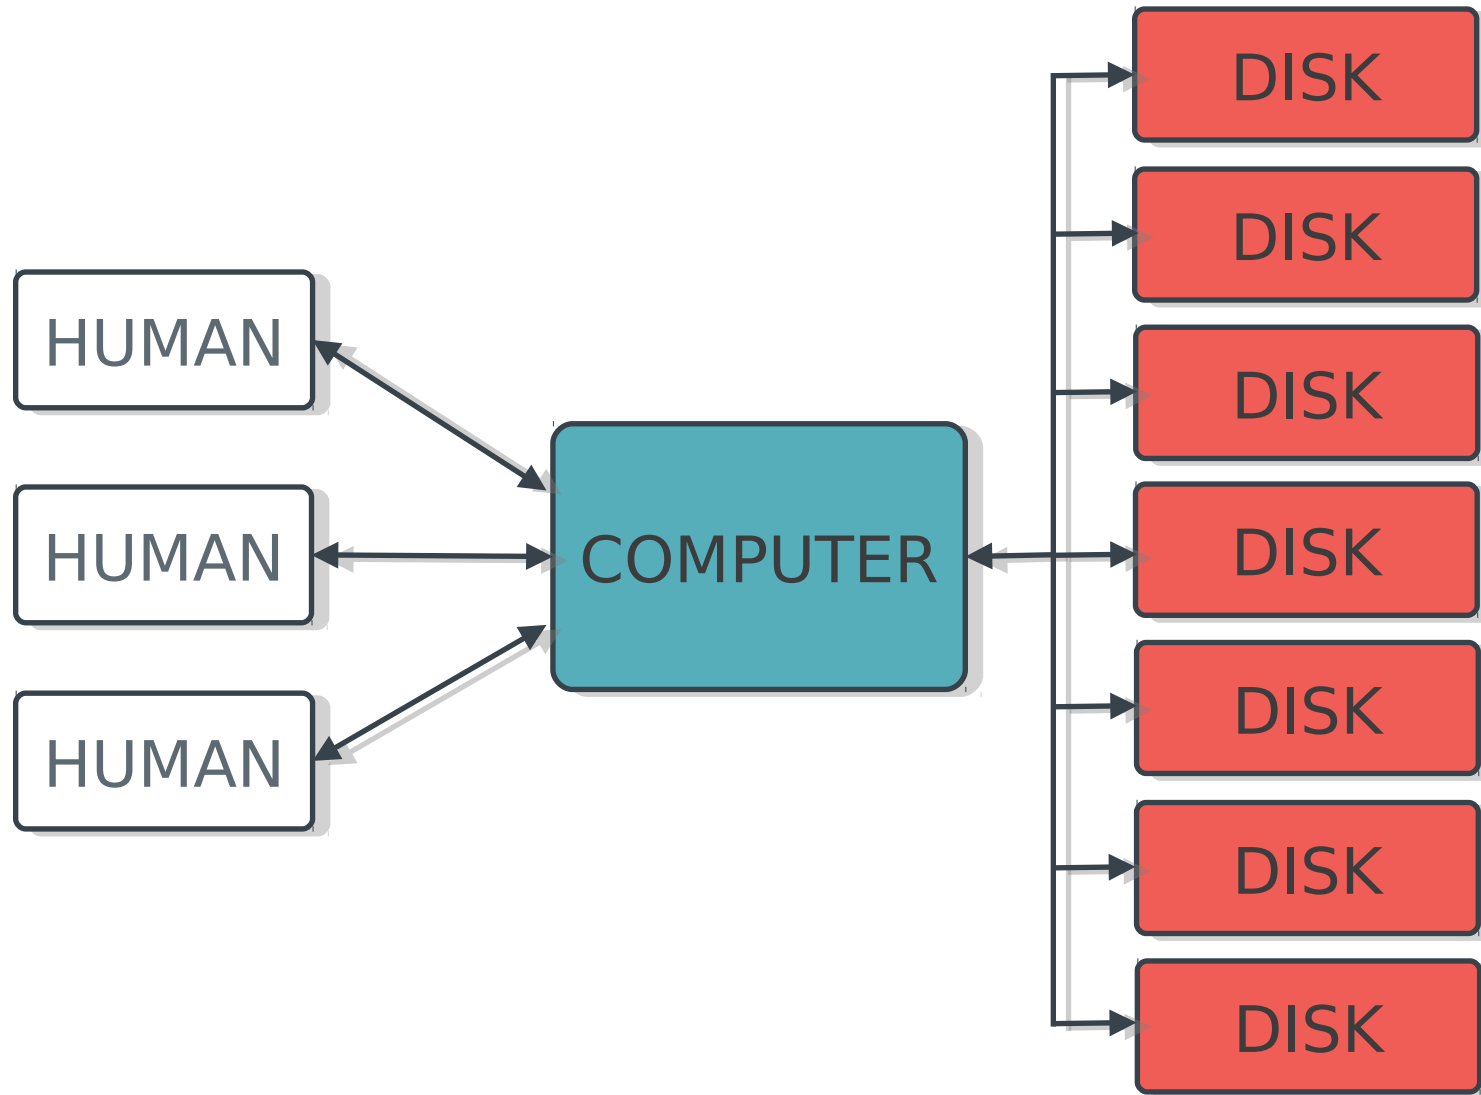
# why start with **objects**?

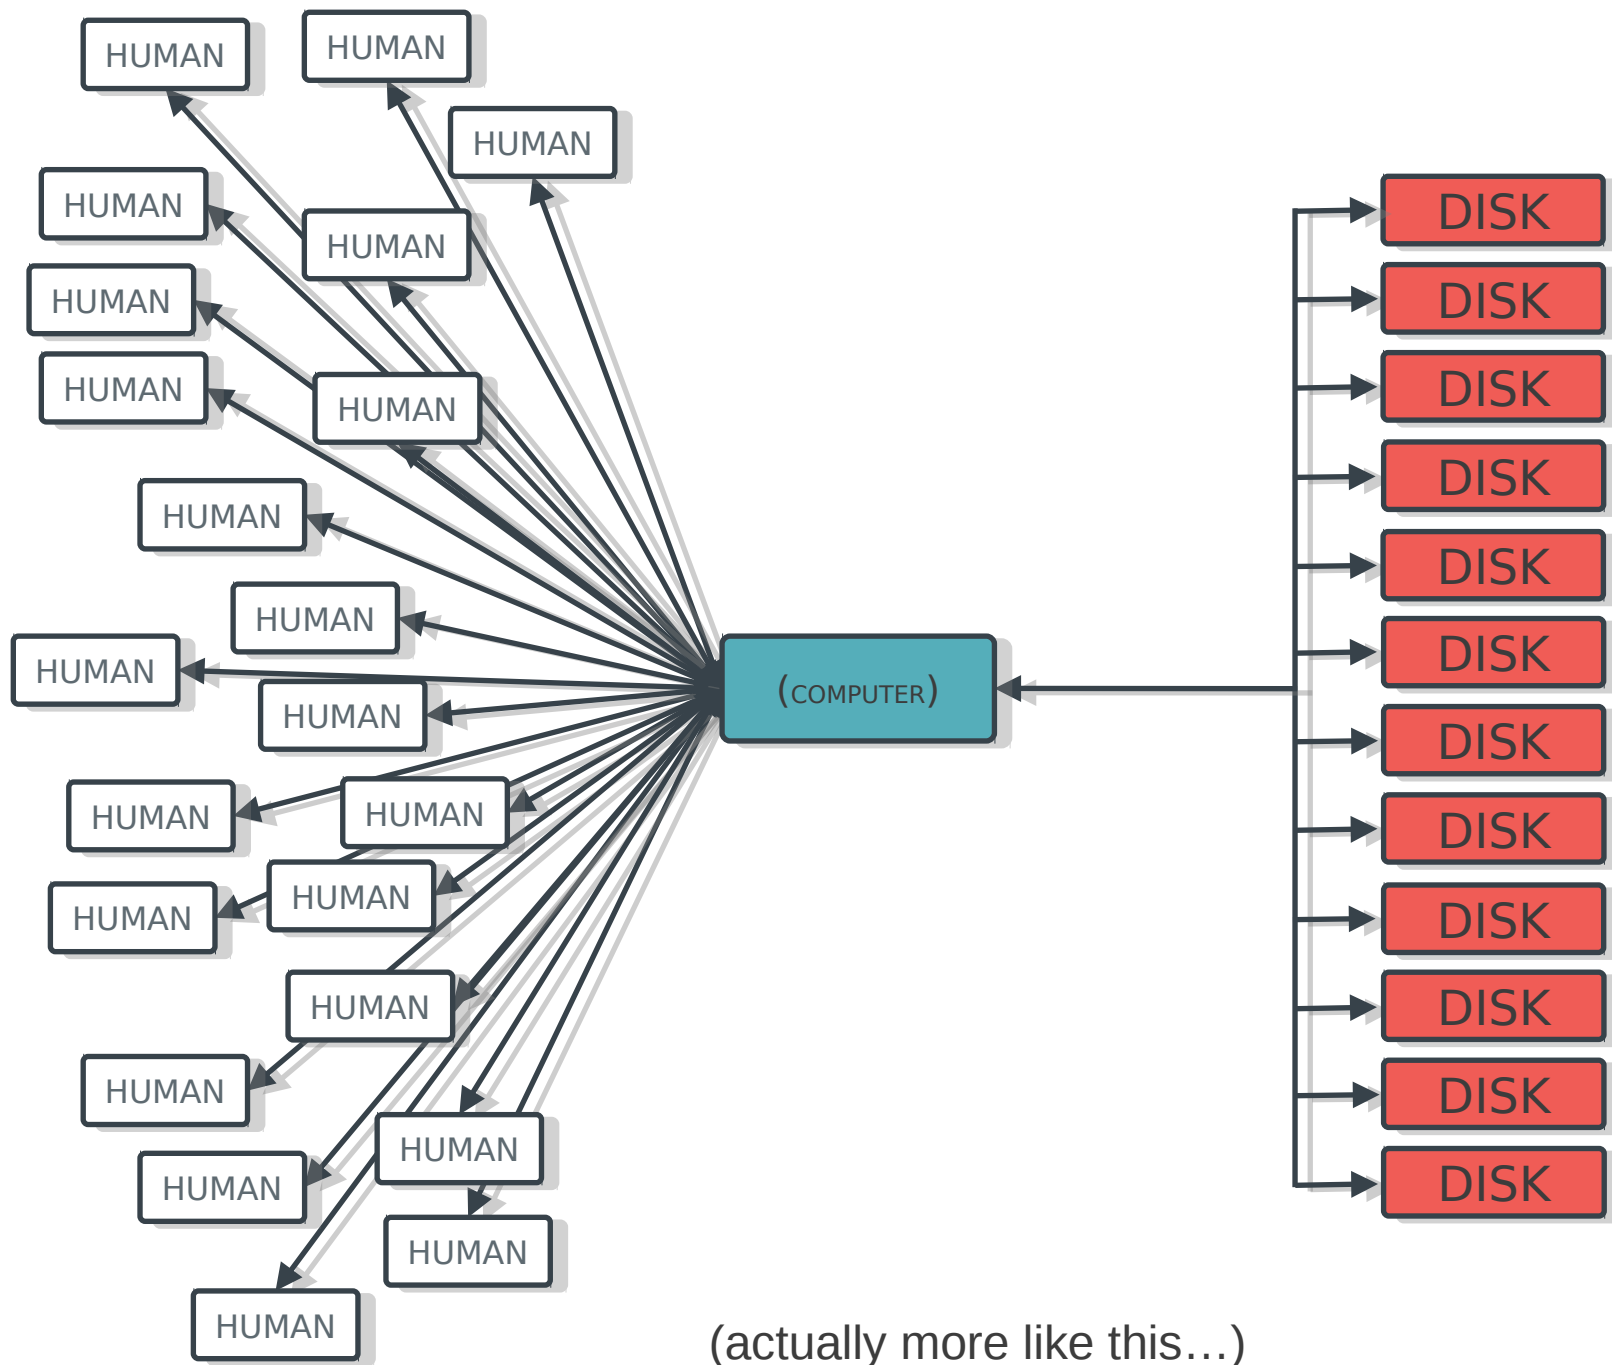
- more useful than (disk) blocks
  - names in a simple flat namespace
  - variable size
  - simple API with rich semantics
- more scalable than files
  - no hard-to-distribute hierarchy
  - update semantics do not span objects
  - workload is trivially parallel

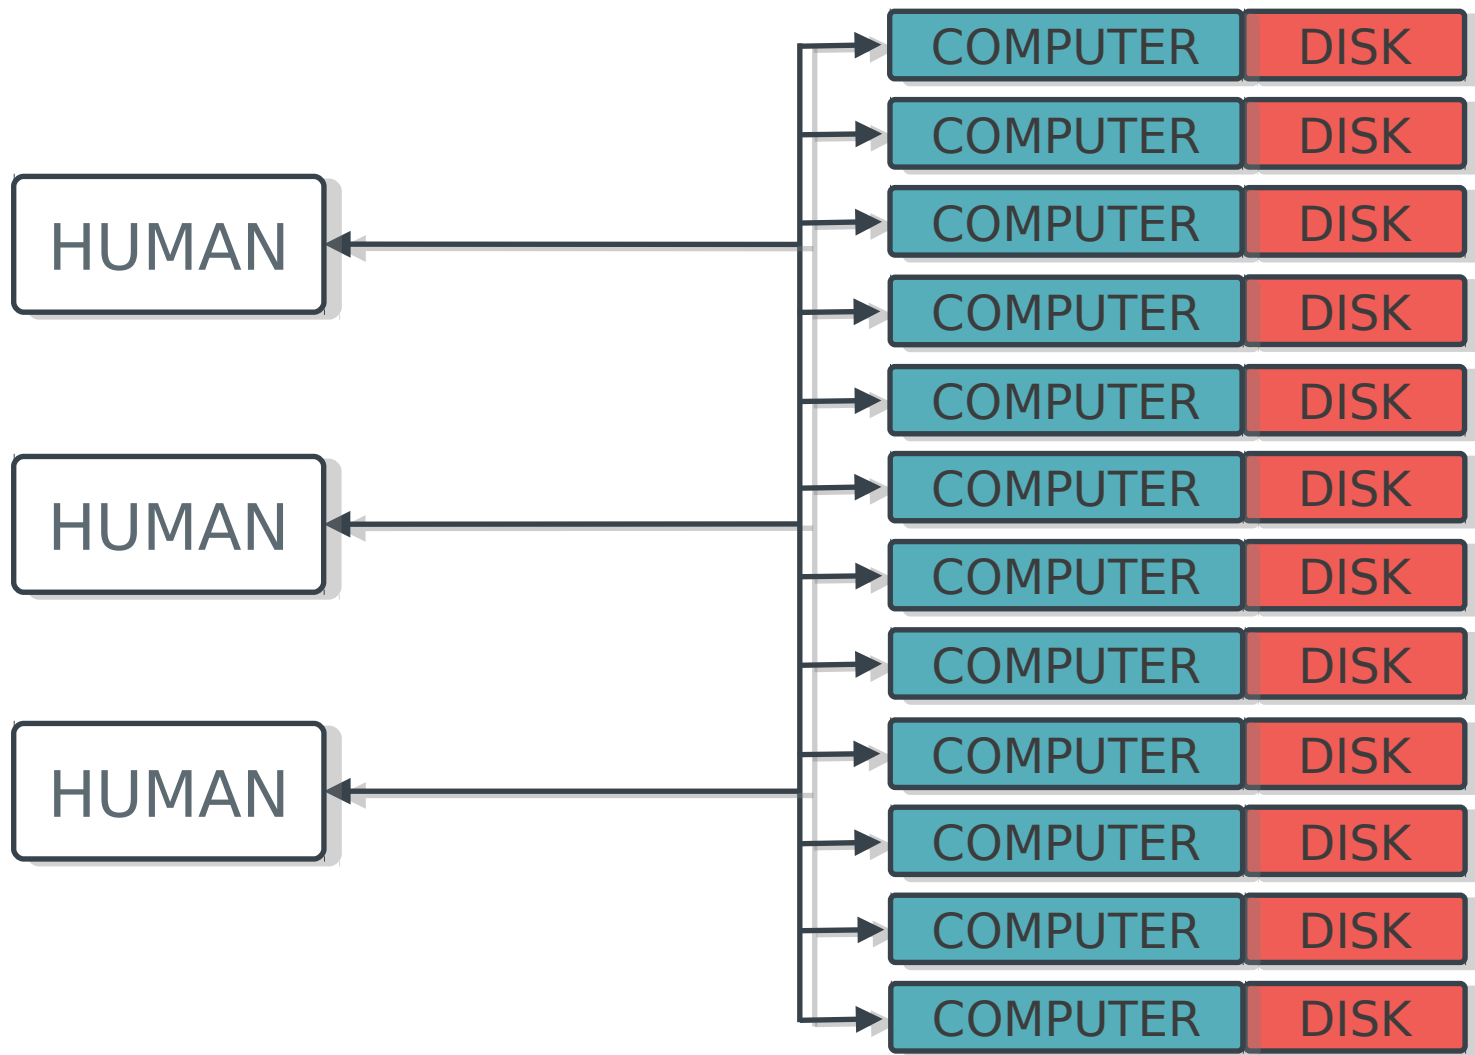
How do you design  
a storage system that scales?

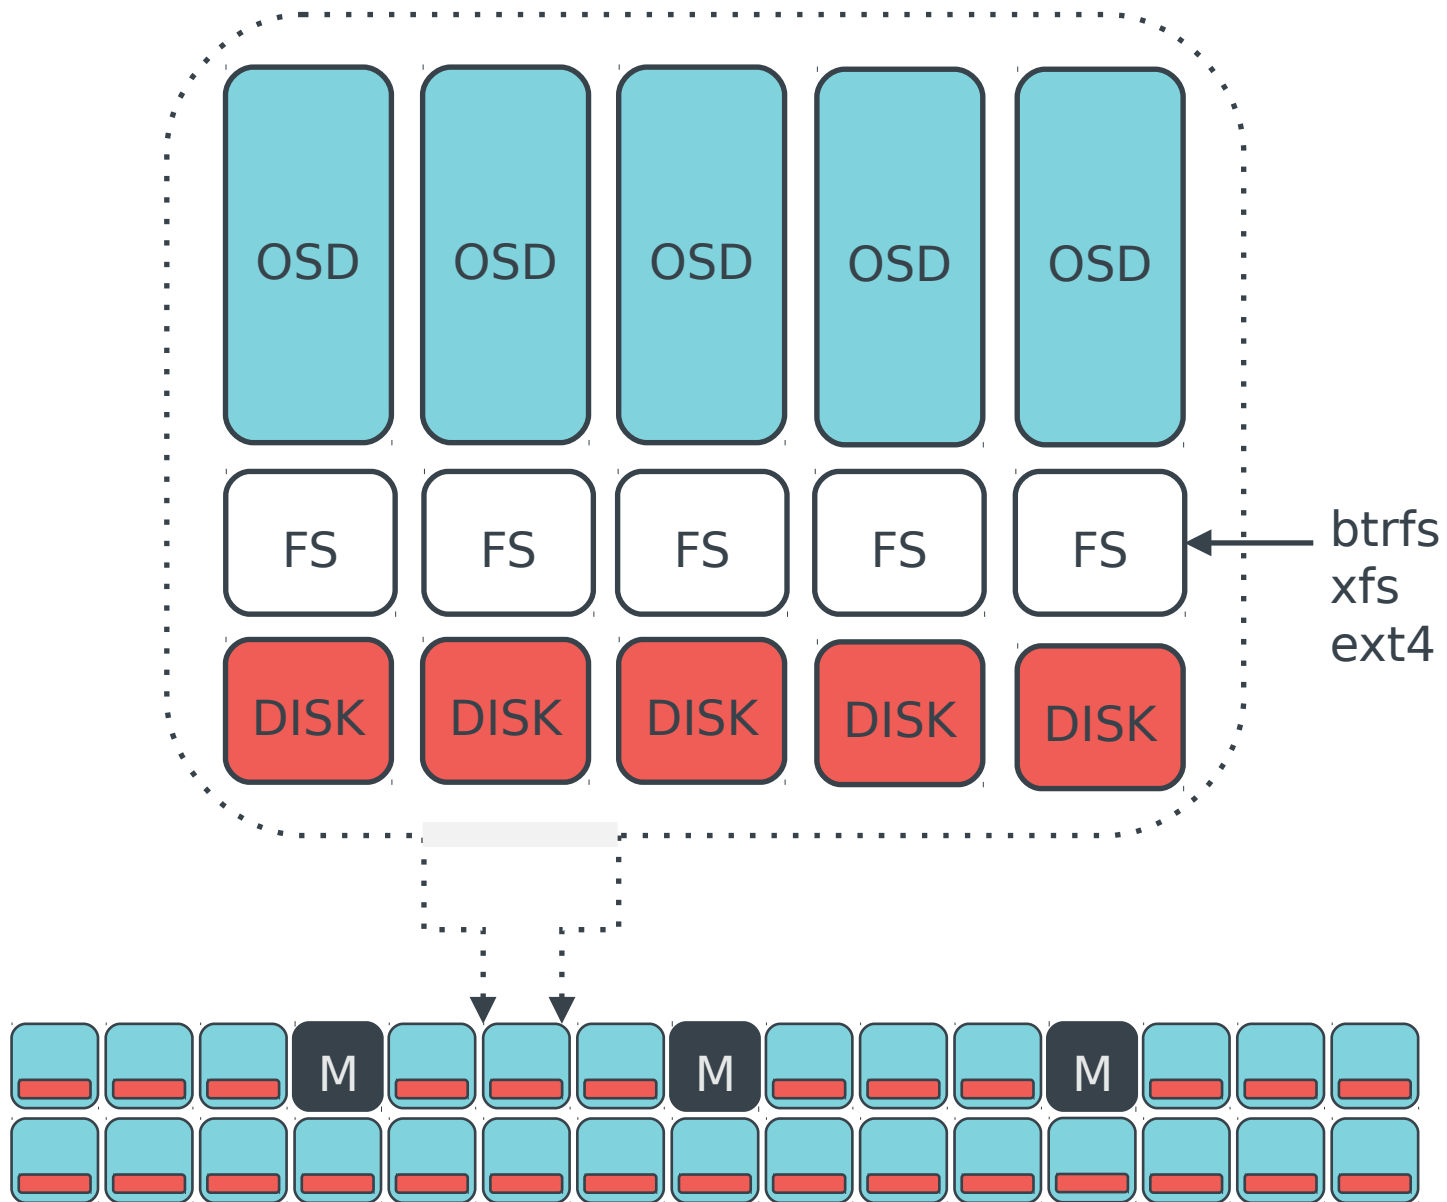








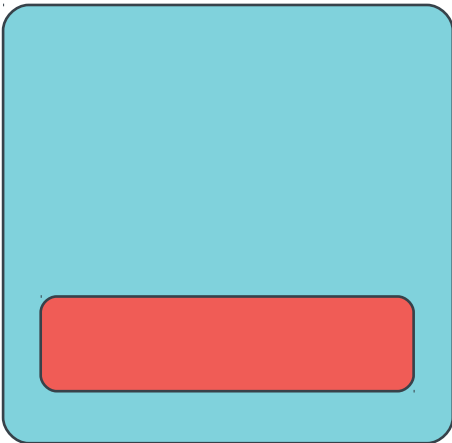






## Monitors

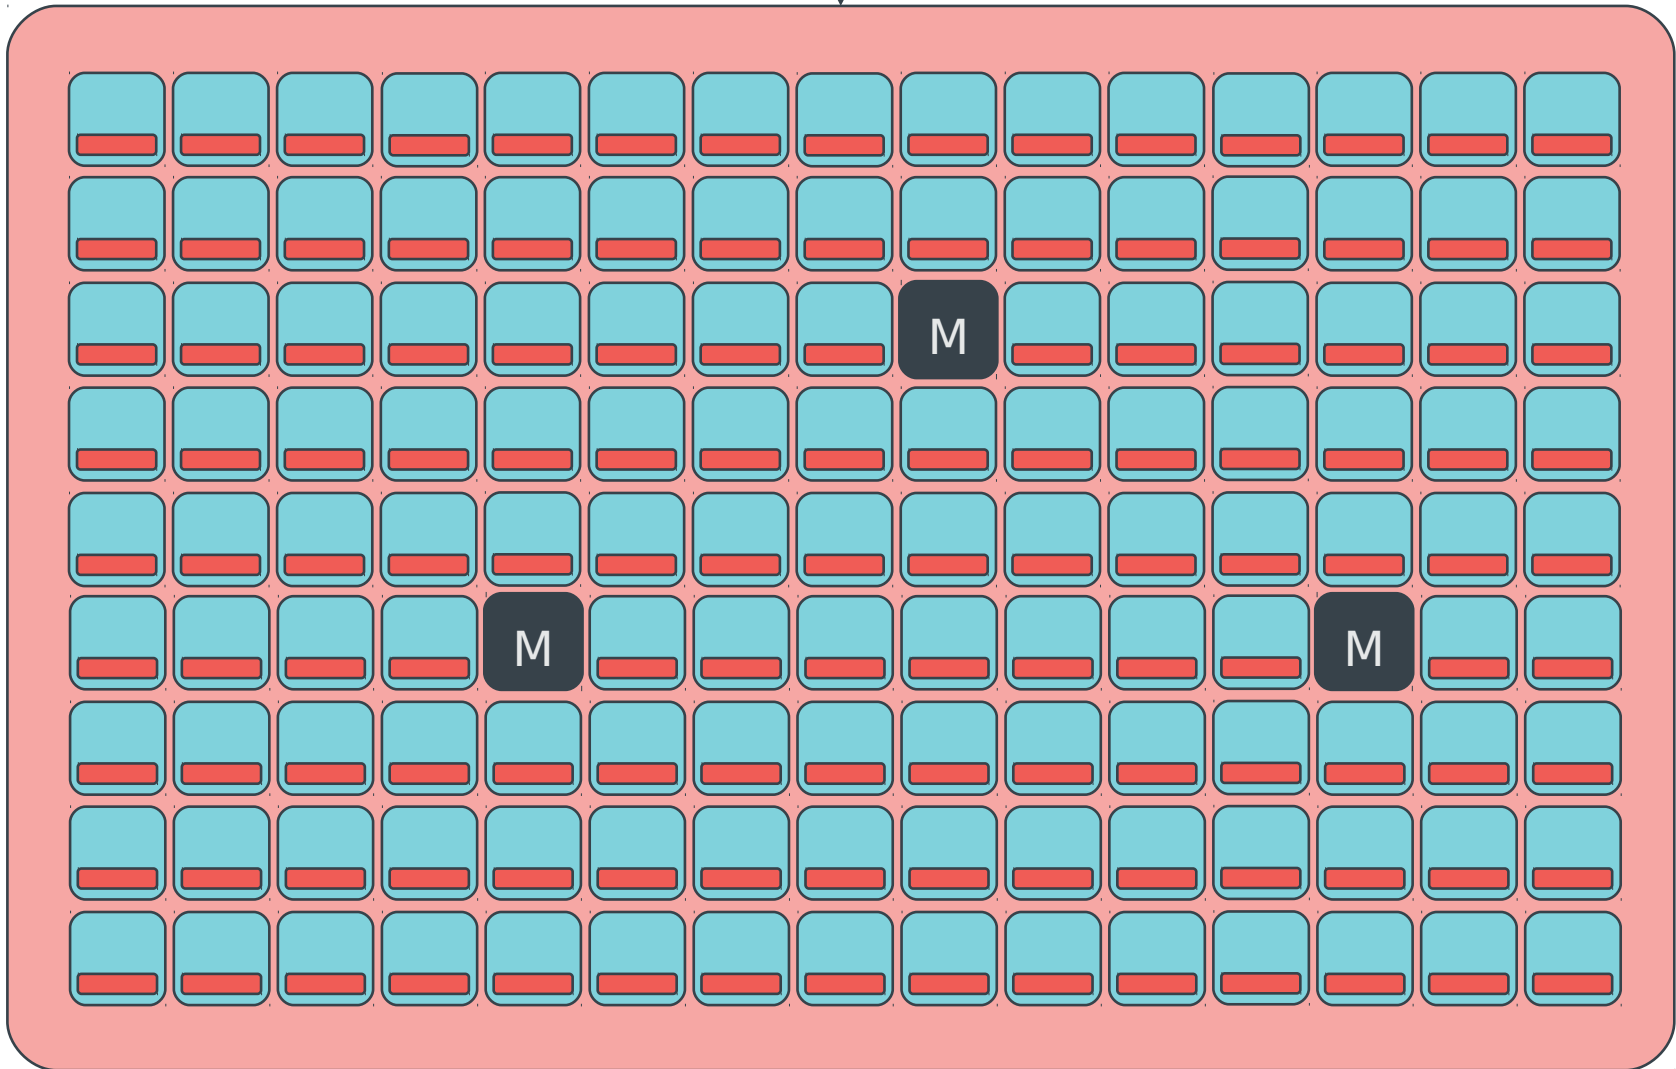
- maintain cluster membership and state
- provide consensus for distributed decision-making
- small, odd number
- these do not served stored objects to clients



## Object Storage Daemons (OSDs)

- 10s to 10000s in a cluster
- one per disk, SSD, or RAID group
  - hardware agnostic
- serve stored objects to clients
- intelligently peer to perform replication and recovery tasks

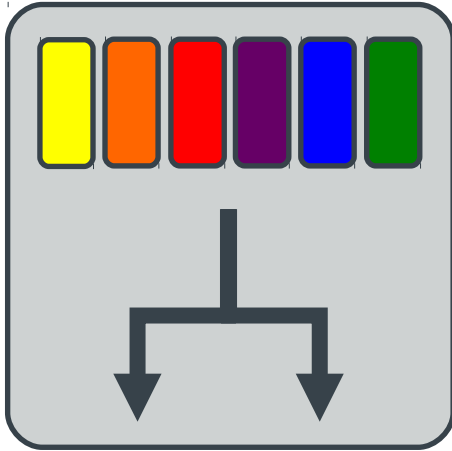
HUMAN



# data distribution

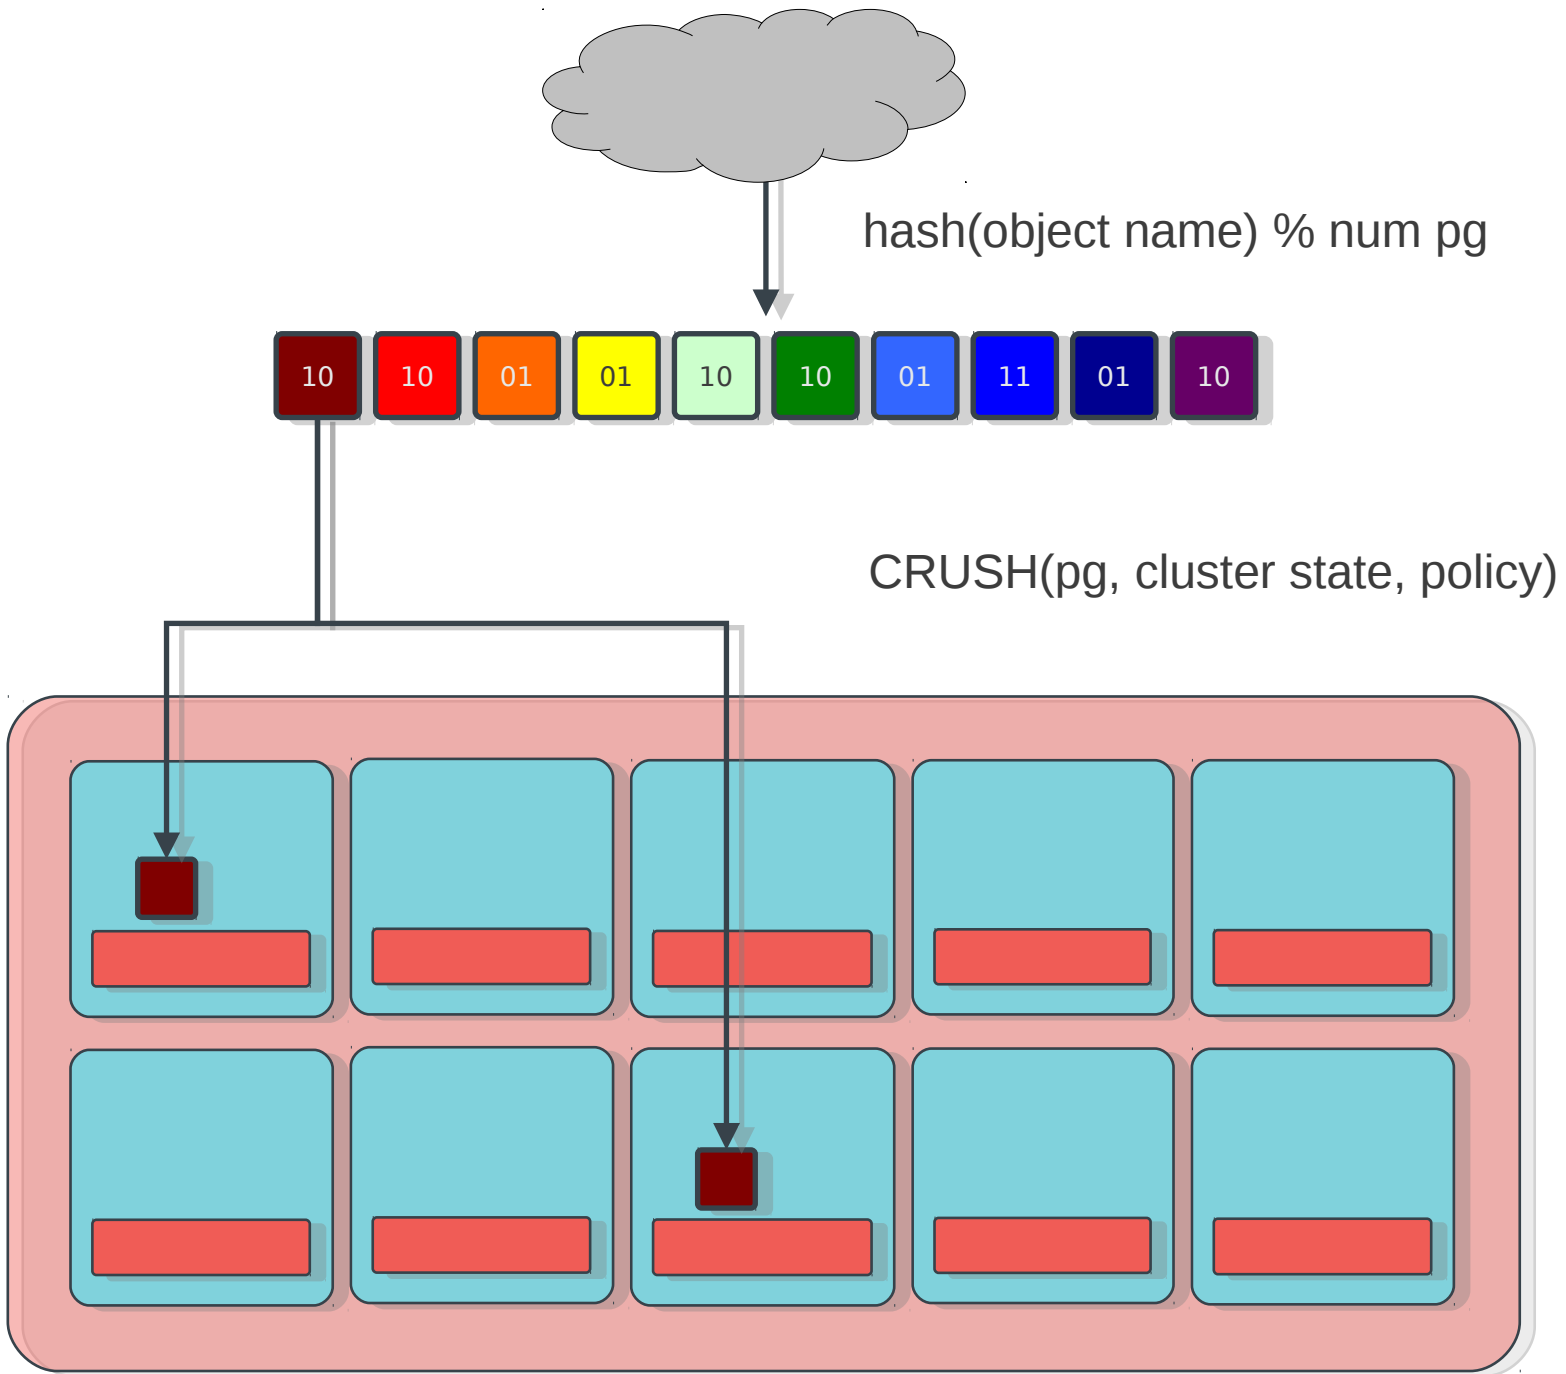
- all objects are replicated N times
- objects are automatically placed, balanced, migrated in a **dynamic** cluster
- must consider physical infrastructure
  - ceph-osds on hosts in racks in rows in data centers
- three approaches
  - pick a spot; remember where you put it
  - pick a spot; write down where you put it
  - calculate where to put it, where to find it

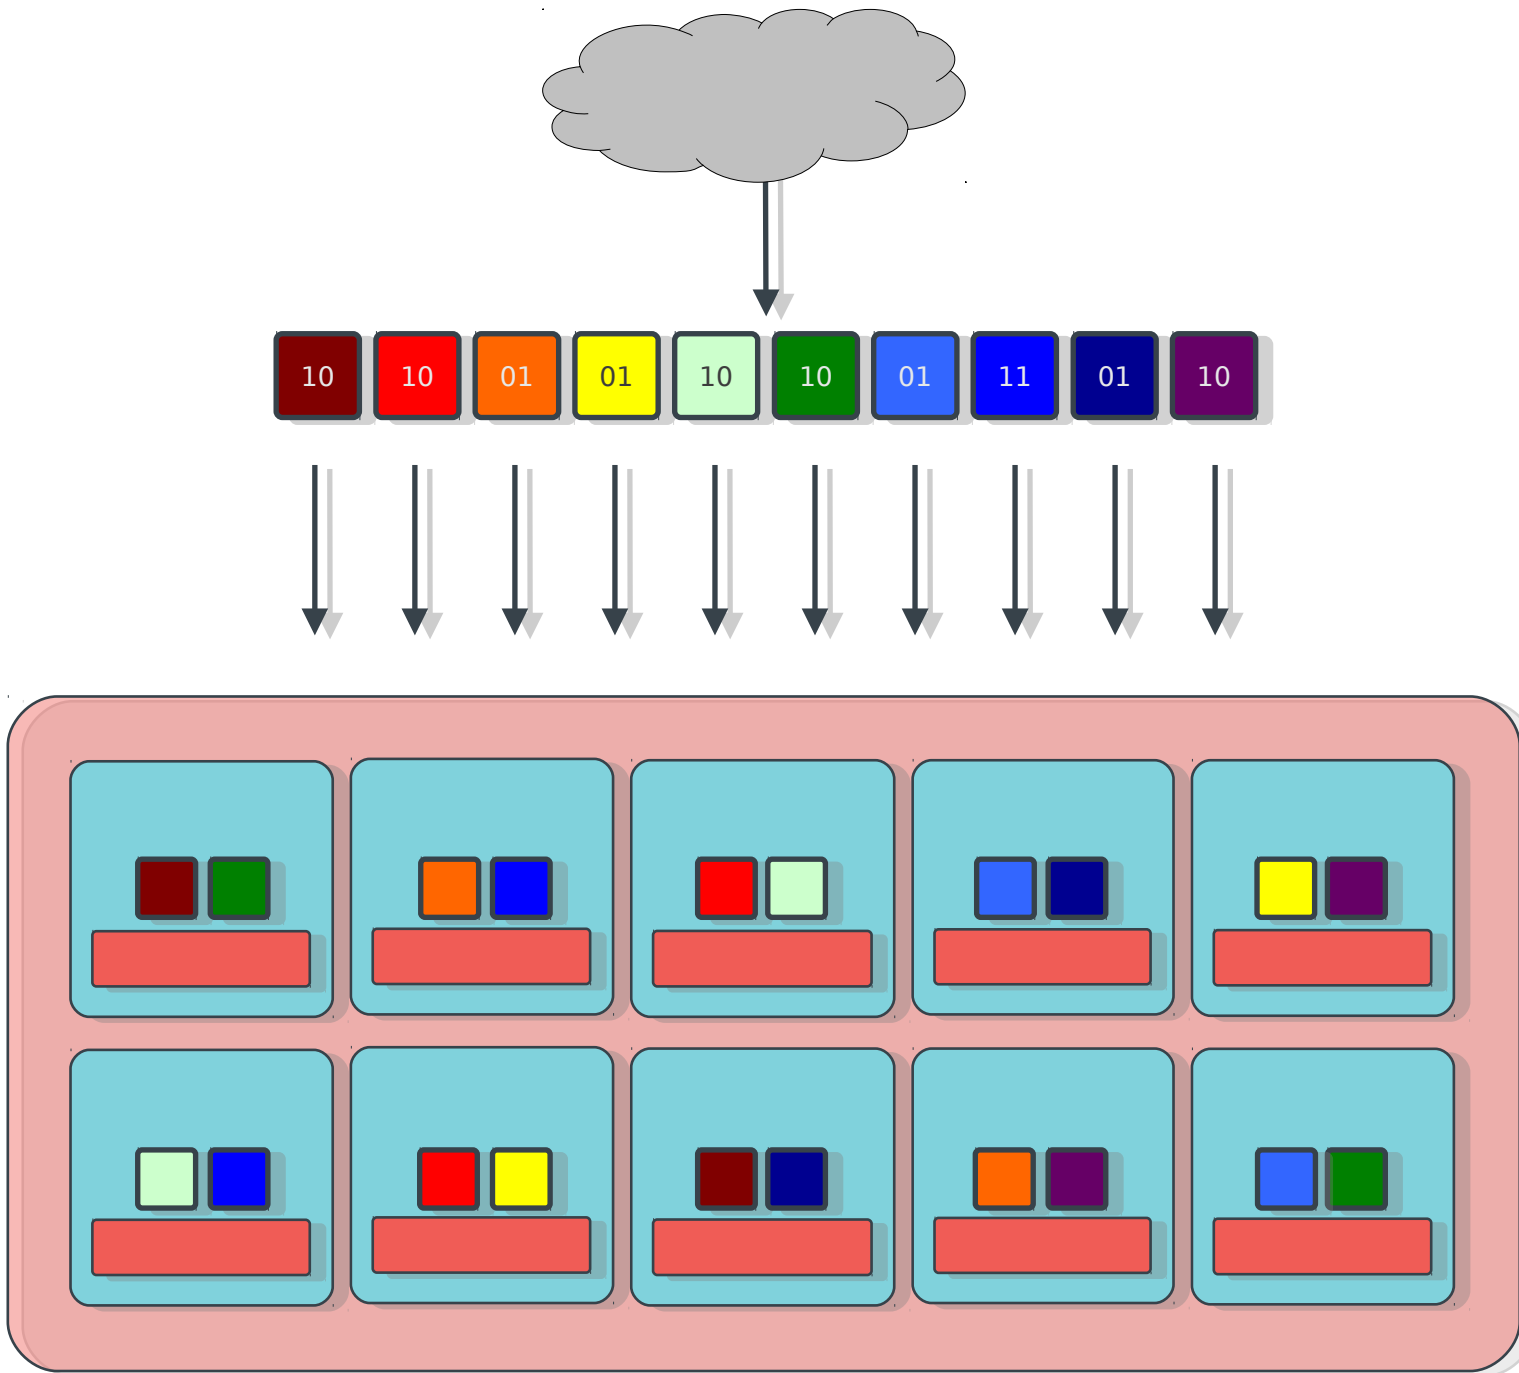




## CRUSH

- pseudo-random placement algorithm
  - fast calculation, **no lookup**
  - repeatable, deterministic
- statistically uniform distribution
- stable mapping
  - limited data migration on change
- rule-based configuration
  - infrastructure topology aware
  - adjustable replication
  - allows weighting

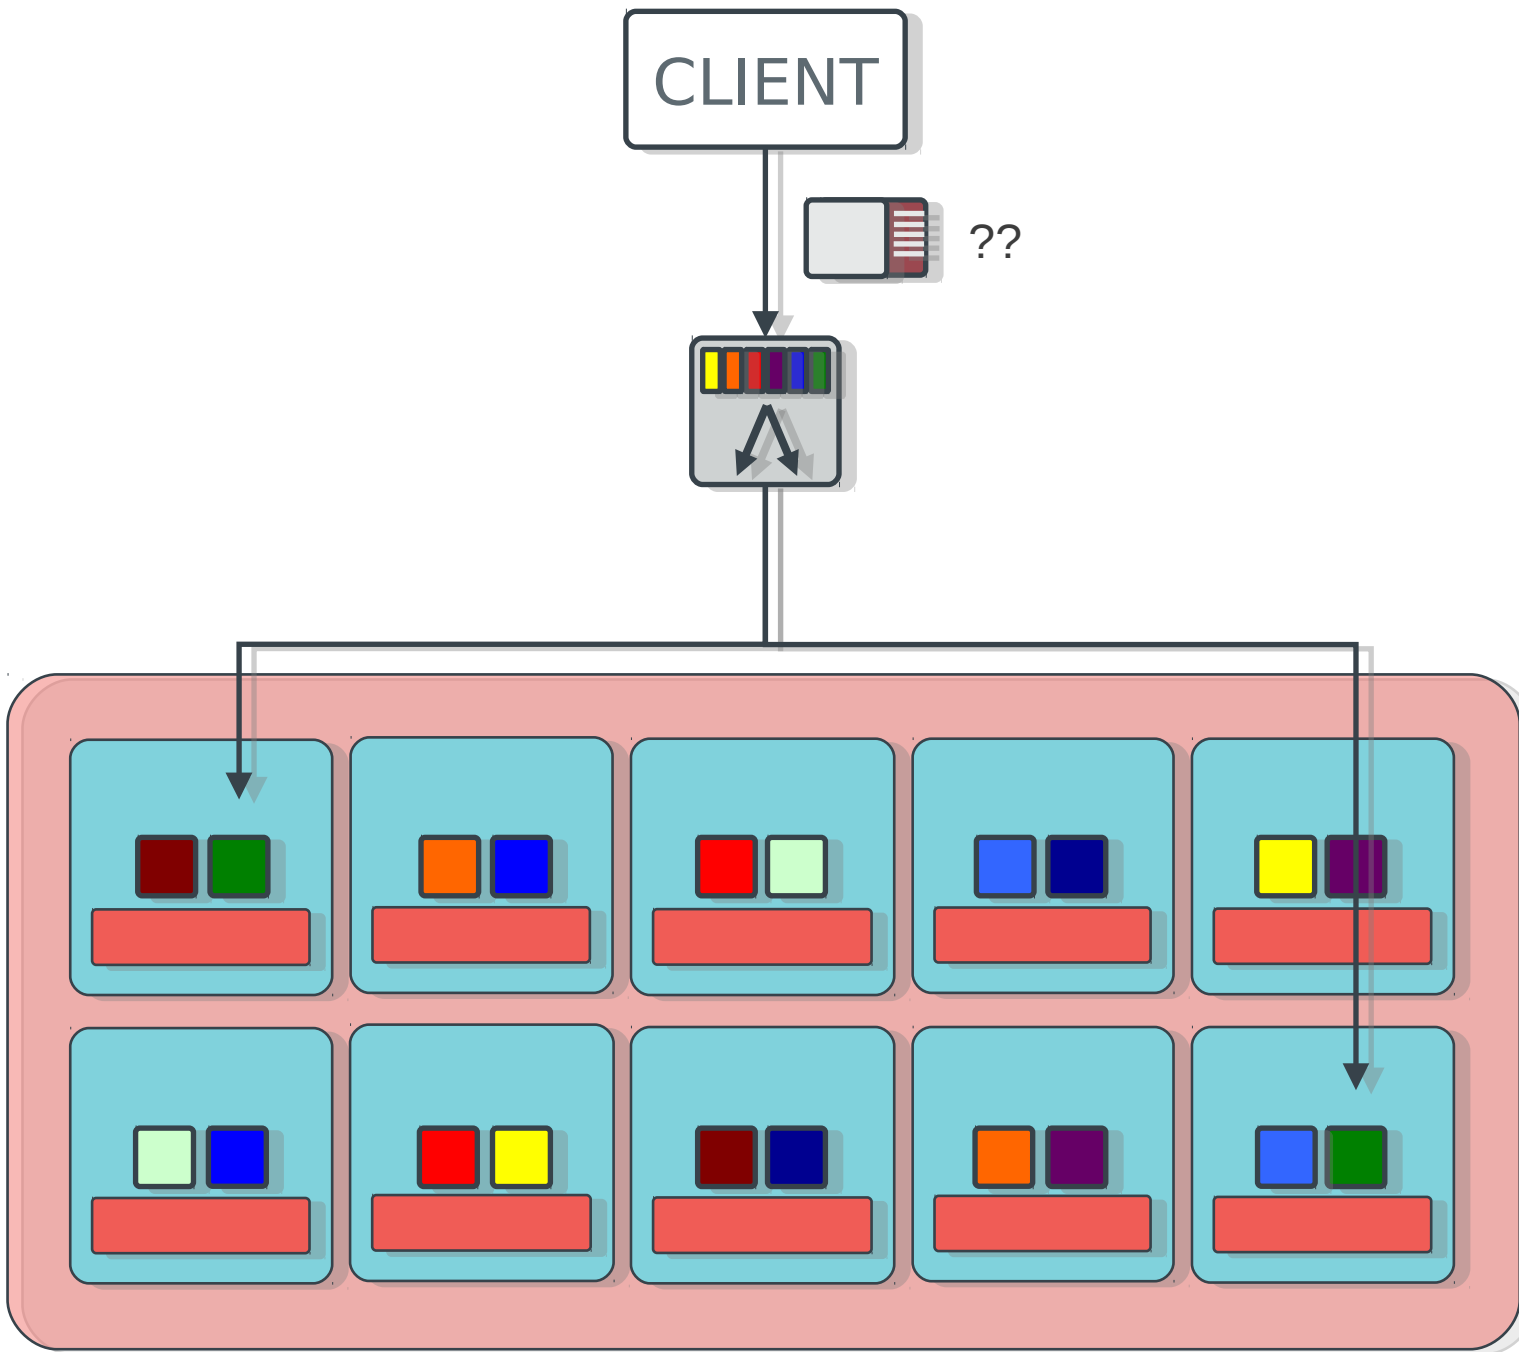


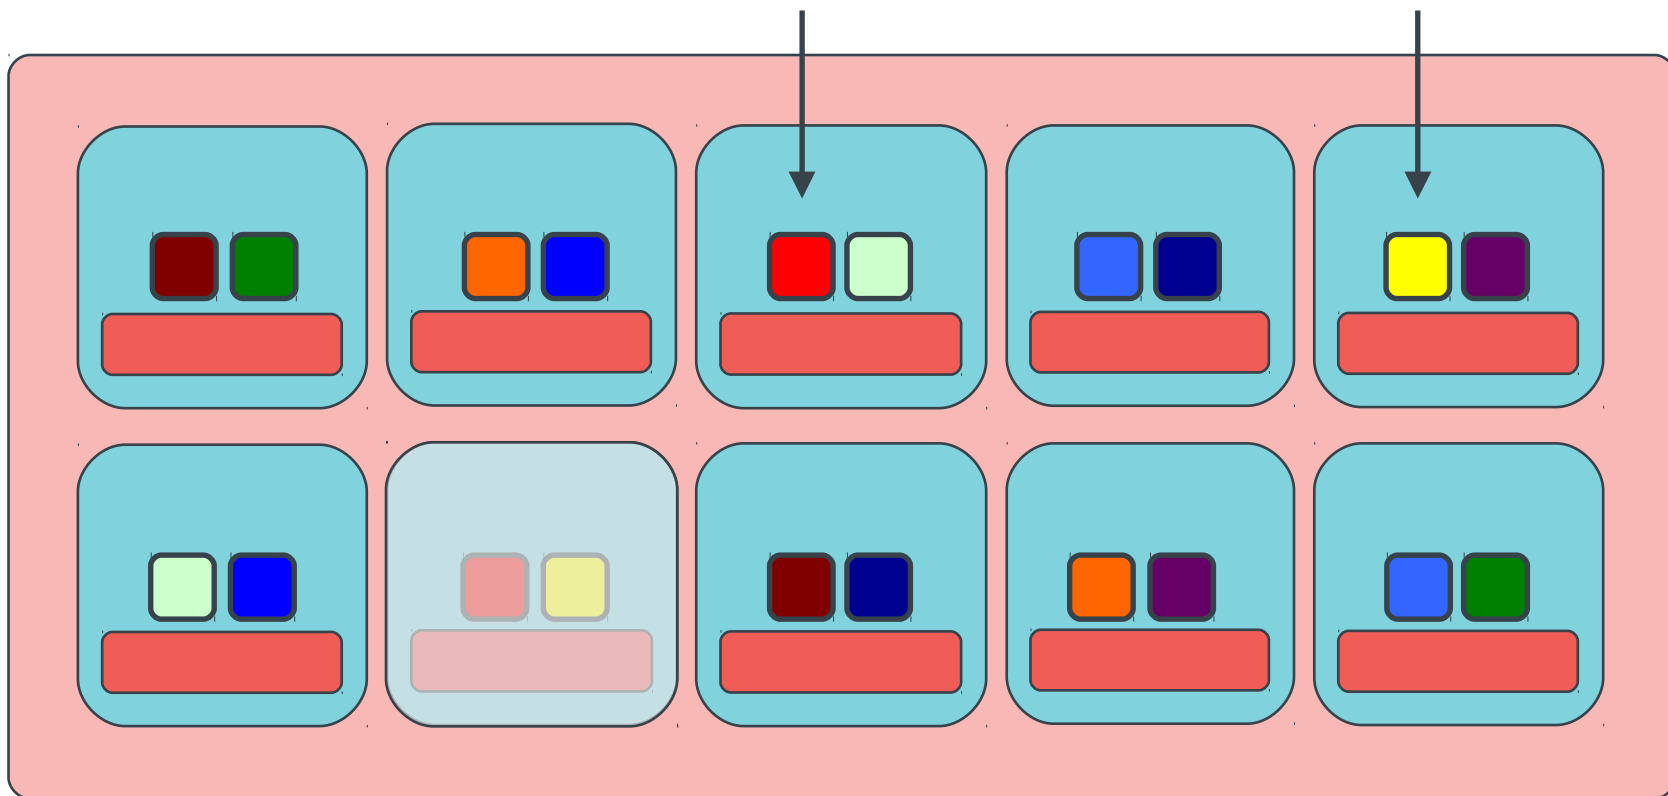


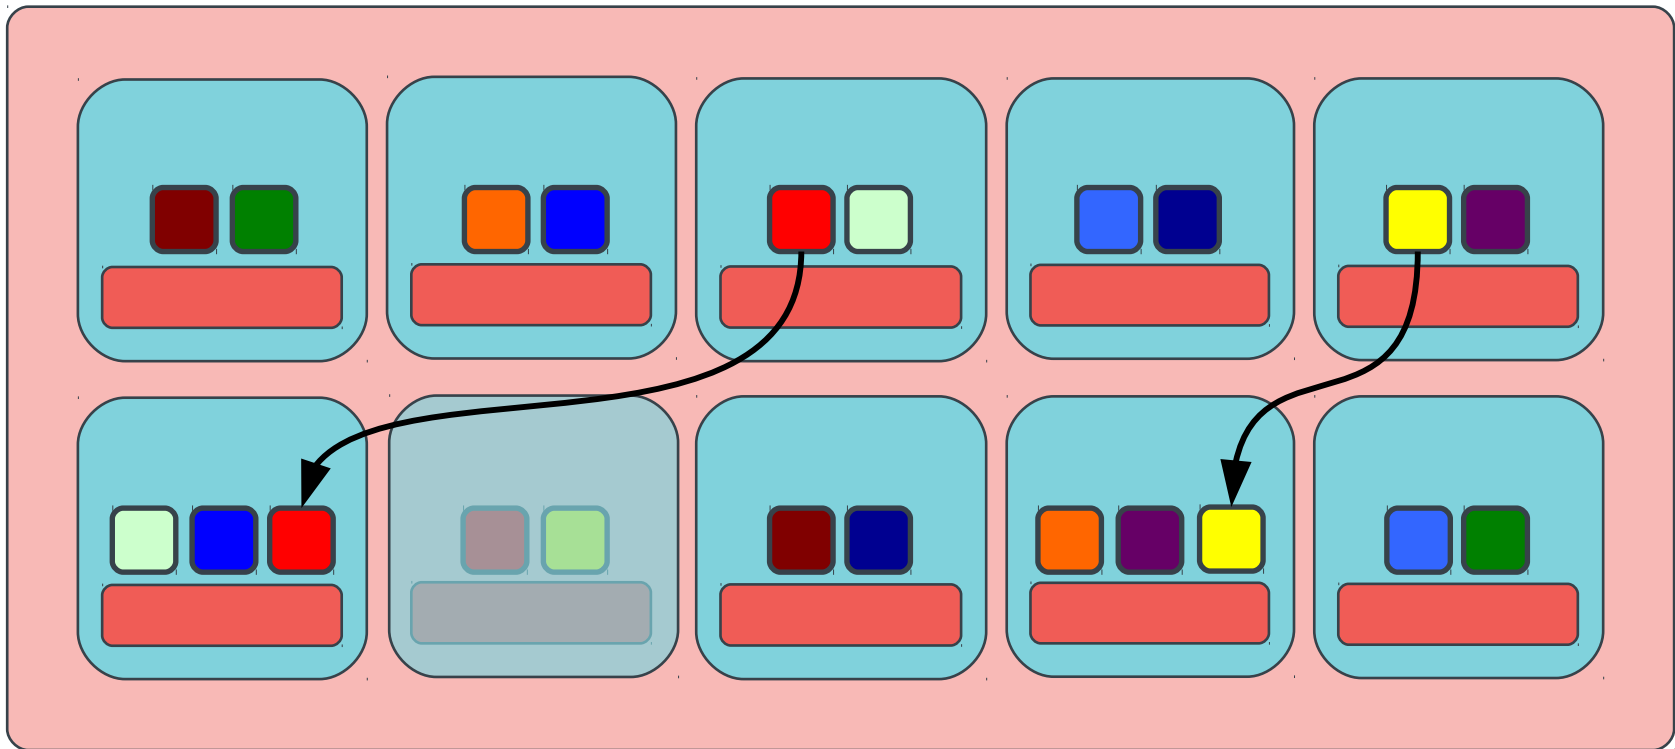
# RADOS

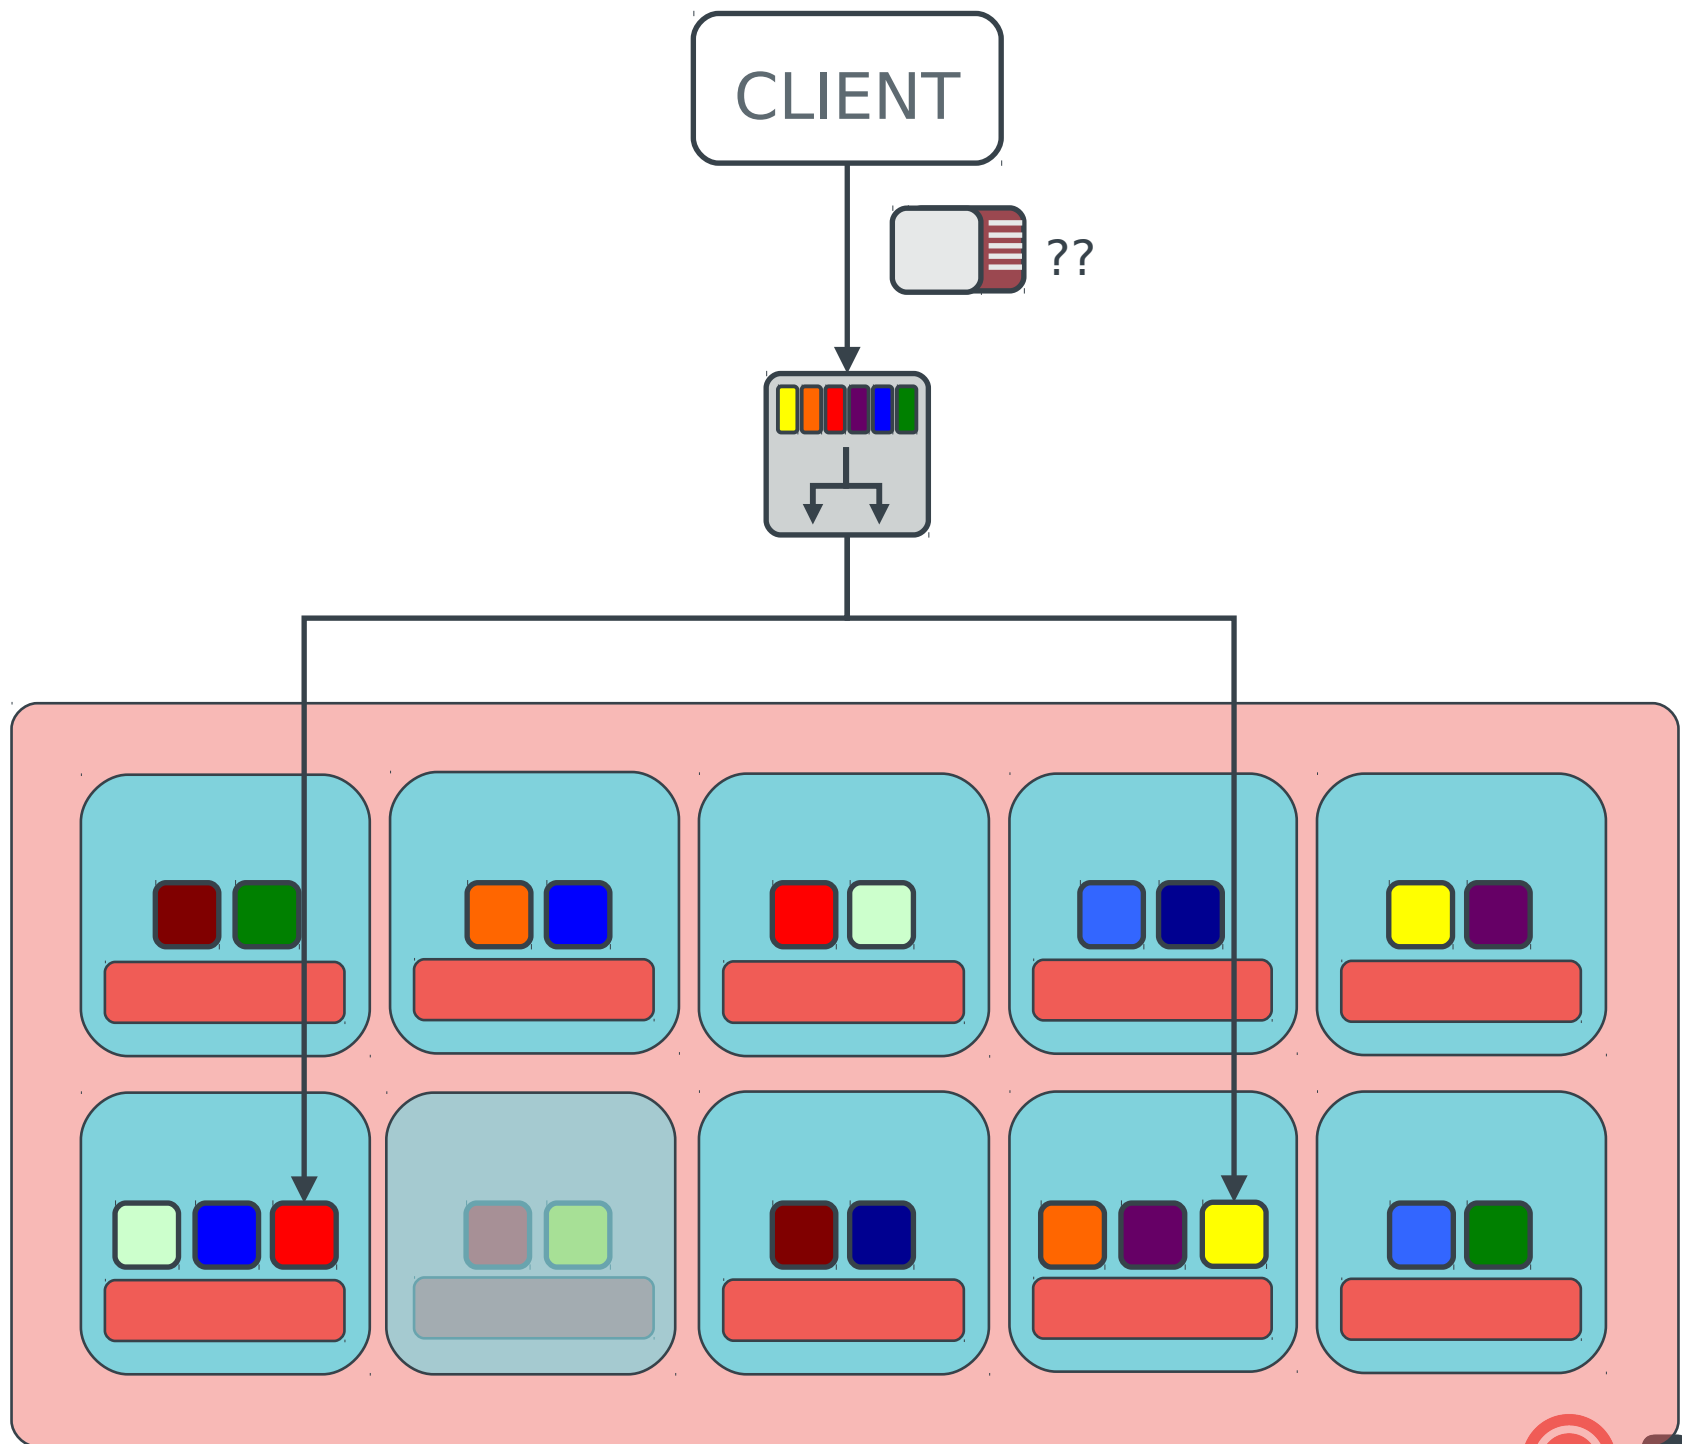
- monitors publish **osd map** that describes cluster state
  - ceph-osd node status (up/down, weight, IP)
  - CRUSH function specifying desired data distribution
- object storage daemons (OSDs)
  - safely replicate and store object
  - migrate data as the cluster changes over time
  - coordinate based on shared view of reality
- decentralized, distributed approach allows
  - massive scales (10,000s of servers or more)
  - the illusion of a single copy with consistent behavior



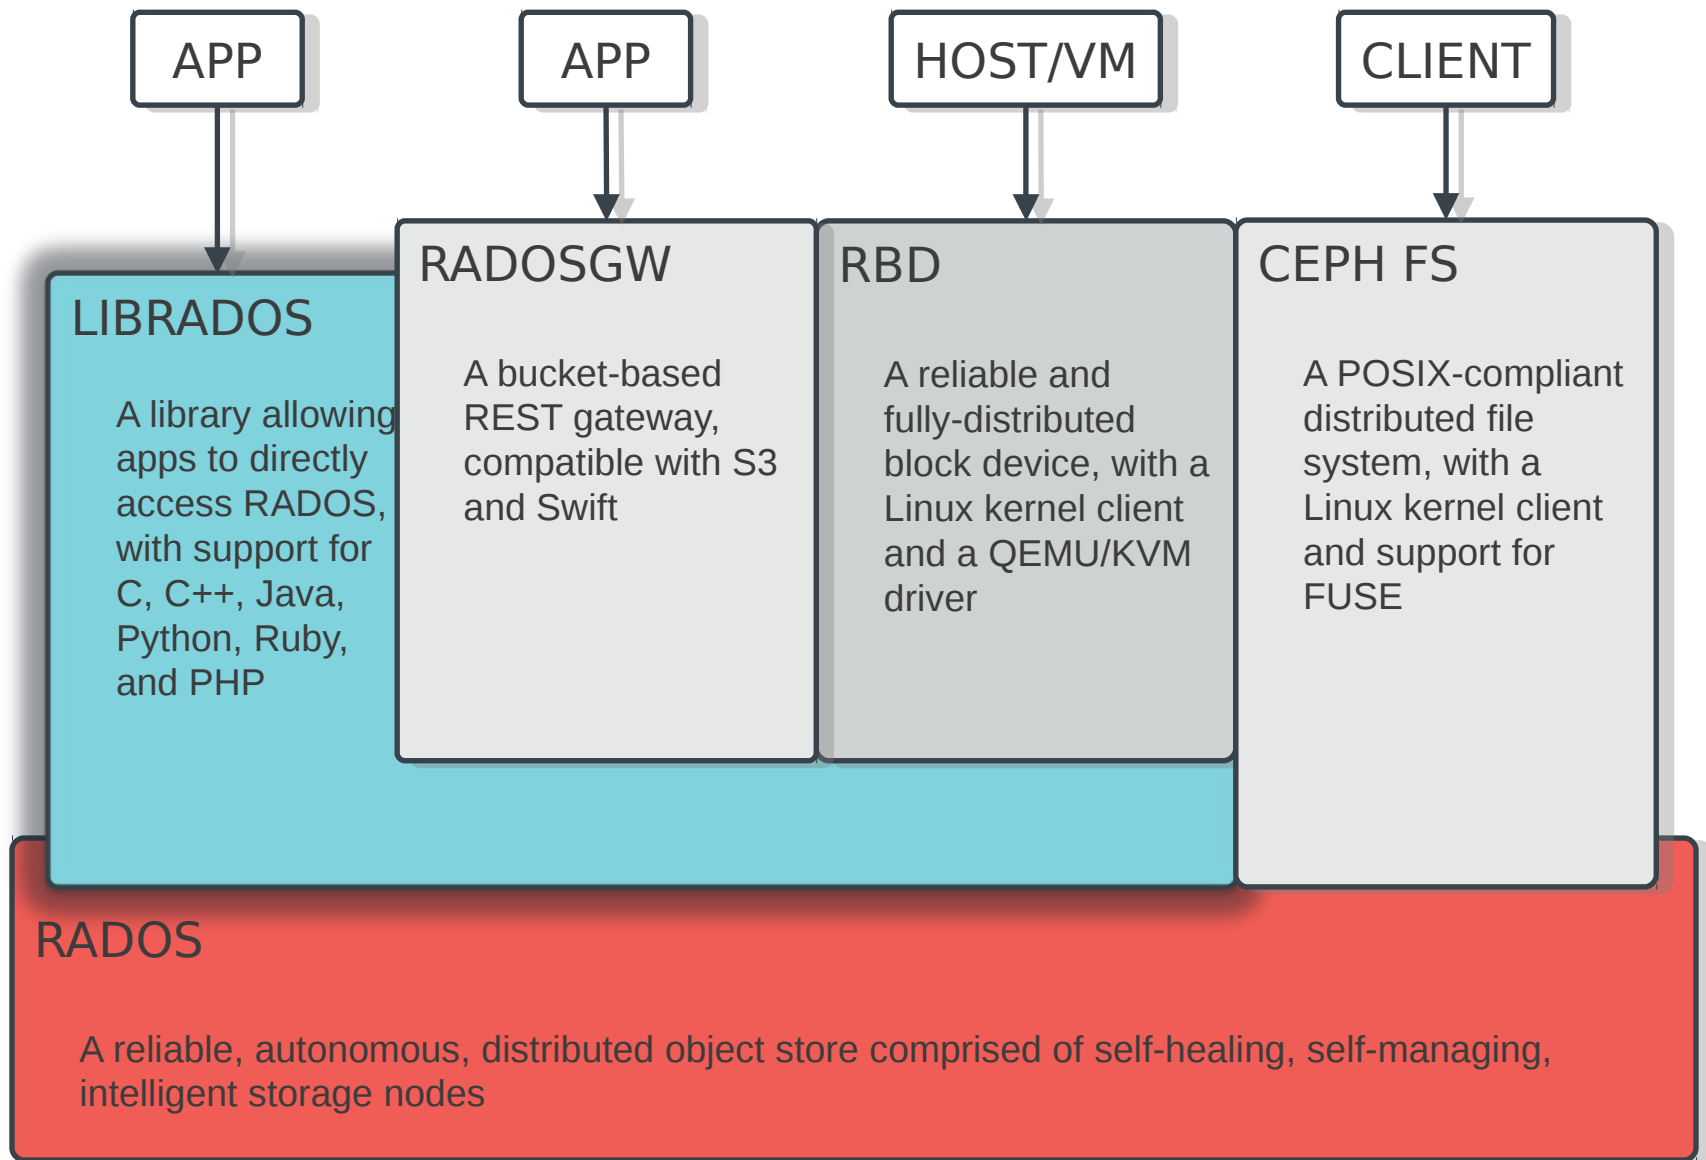


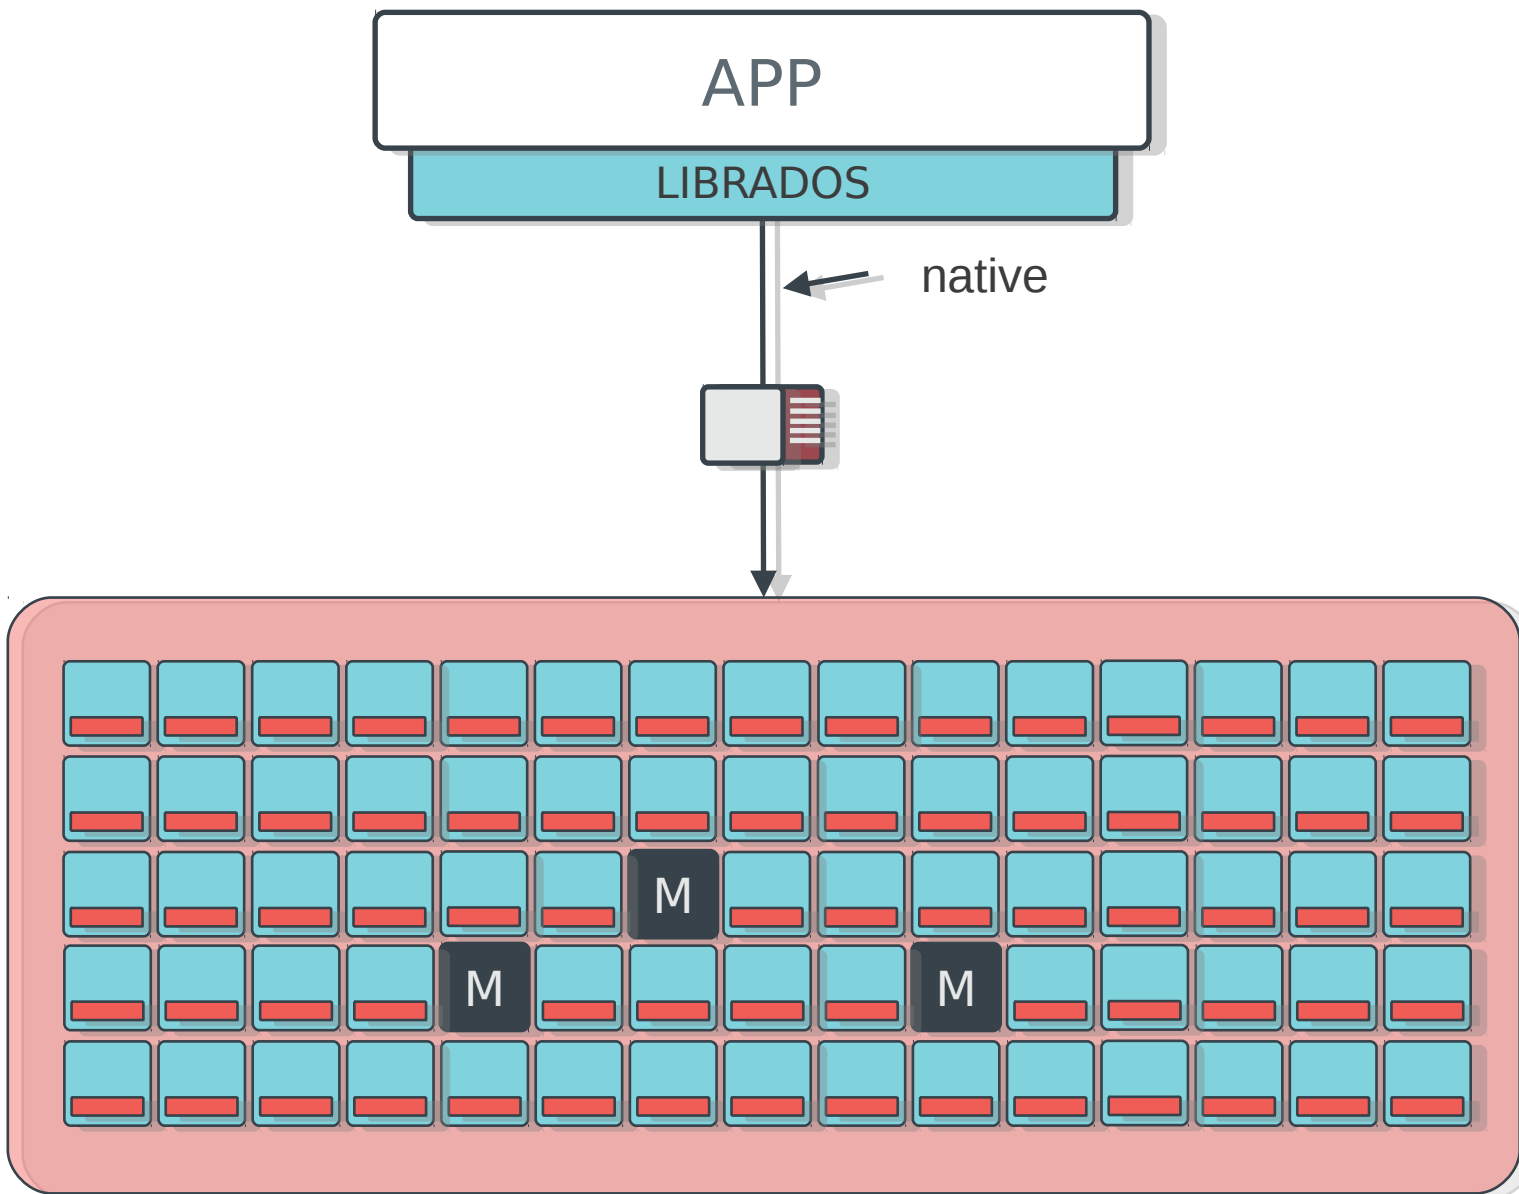












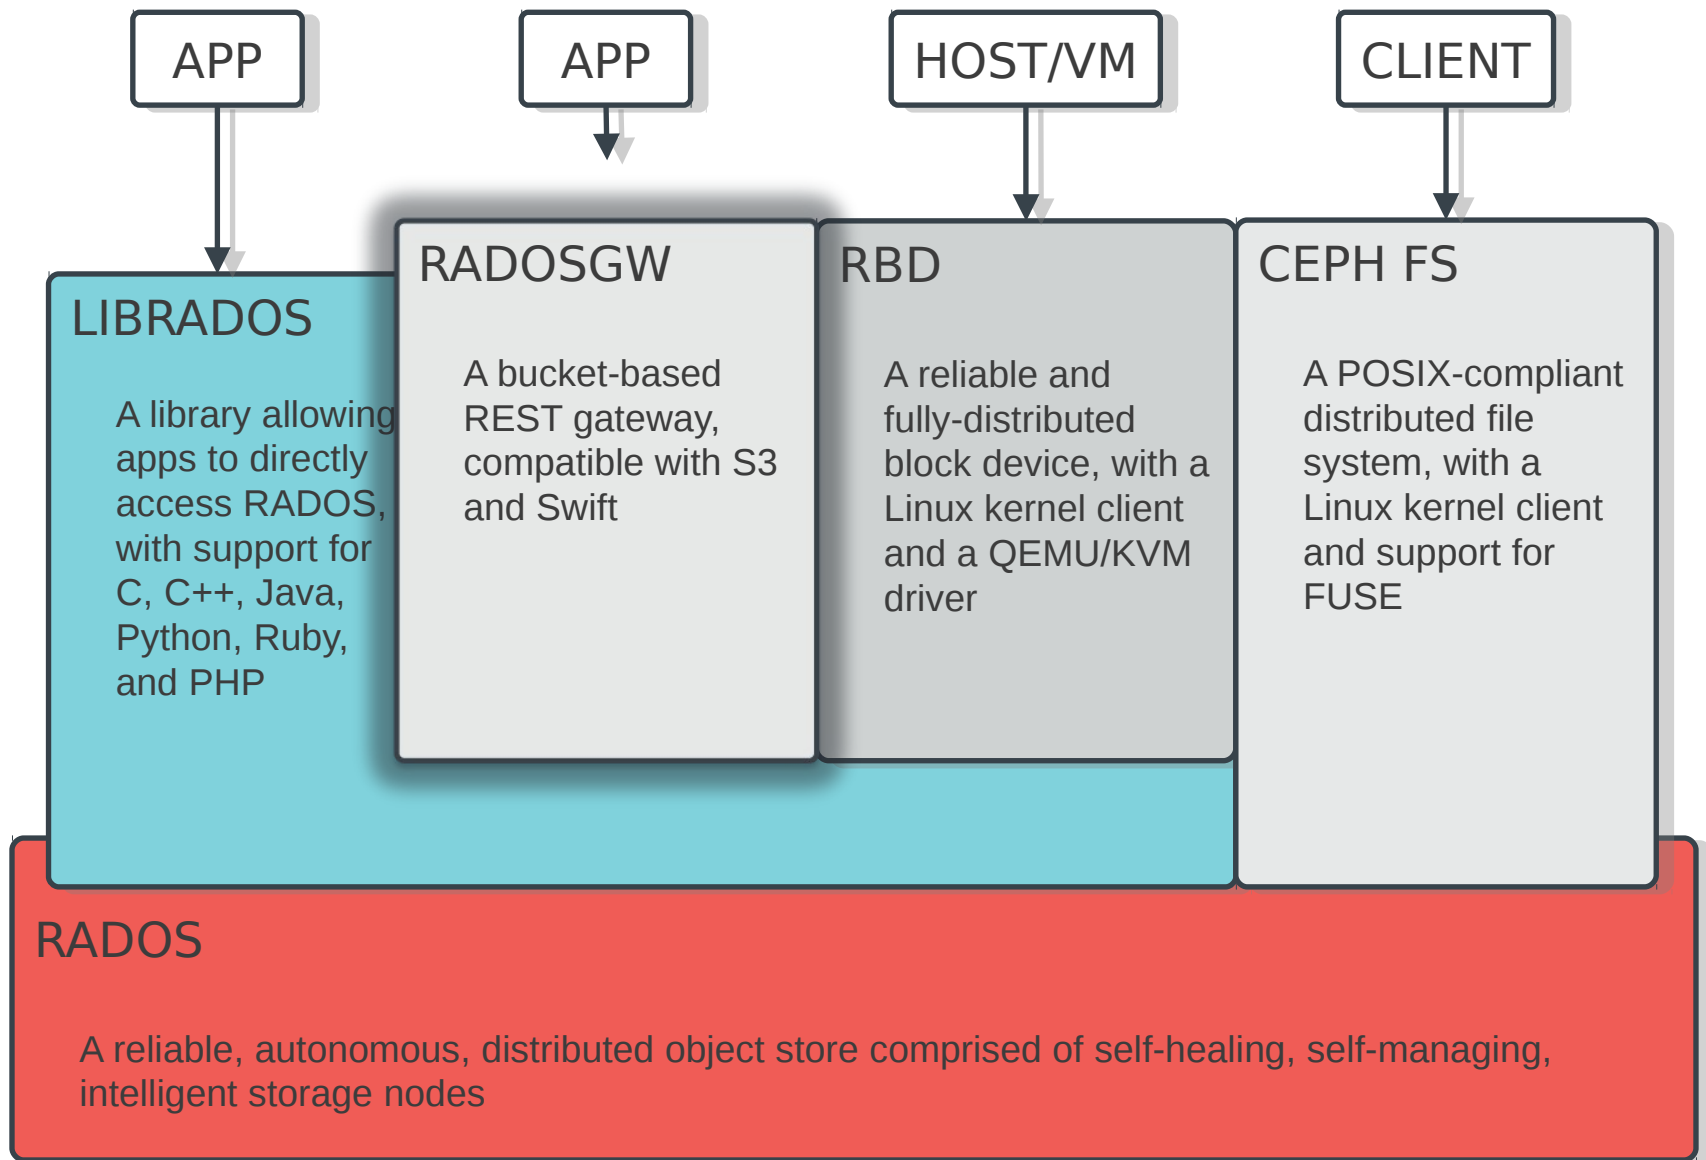


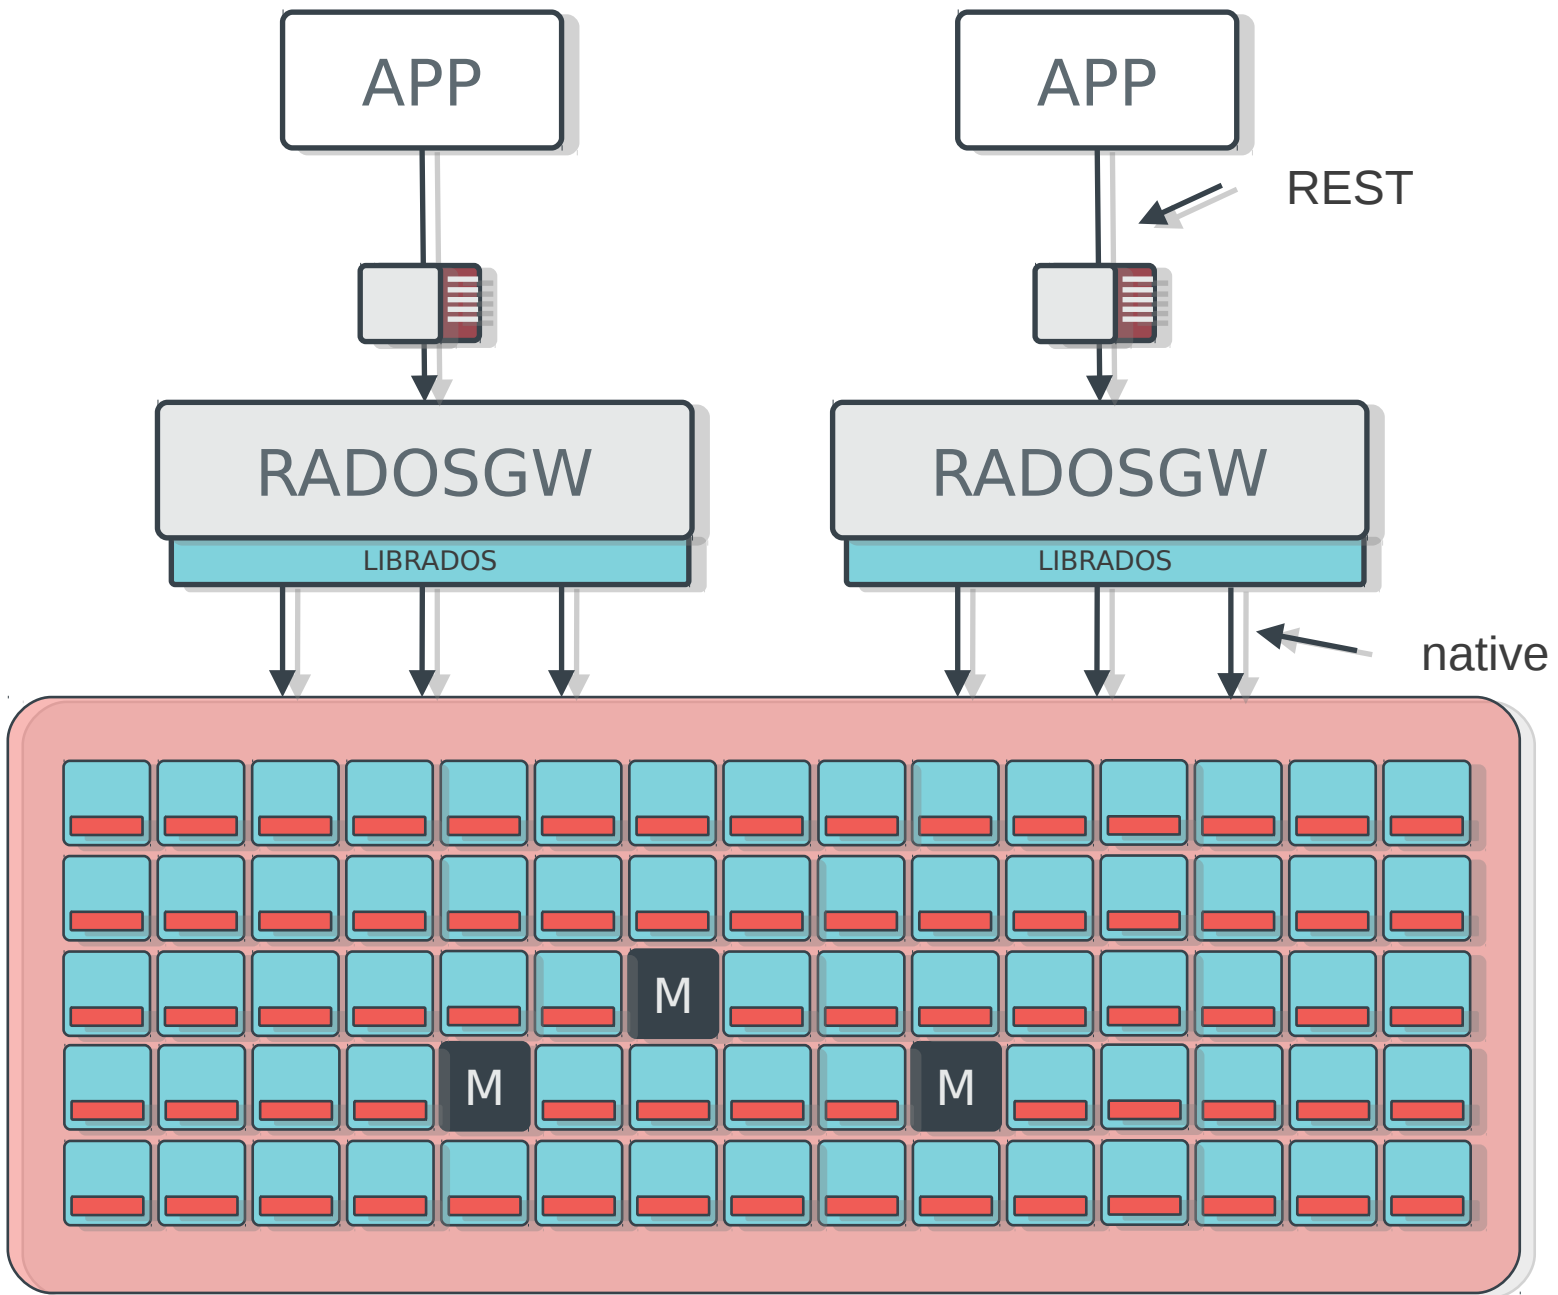
## librados

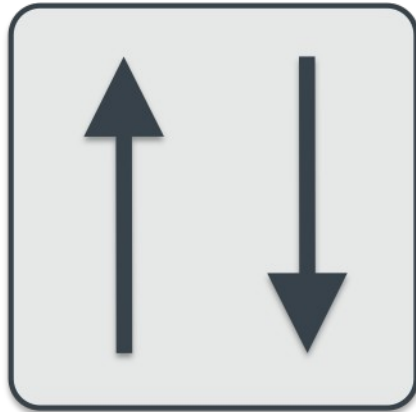
- direct access to RADOS from applications
- C, C++, Python, PHP, Java, Erlang
- direct access to storage nodes
- no HTTP overhead

# rich librados API

- atomic single-object transactions
  - update data, attr together
  - atomic compare-and-swap
- efficient key/value storage inside an object
- object-granularity snapshot infrastructure
- embed code in ceph-osd daemon via plugin API
- inter-client communication via object

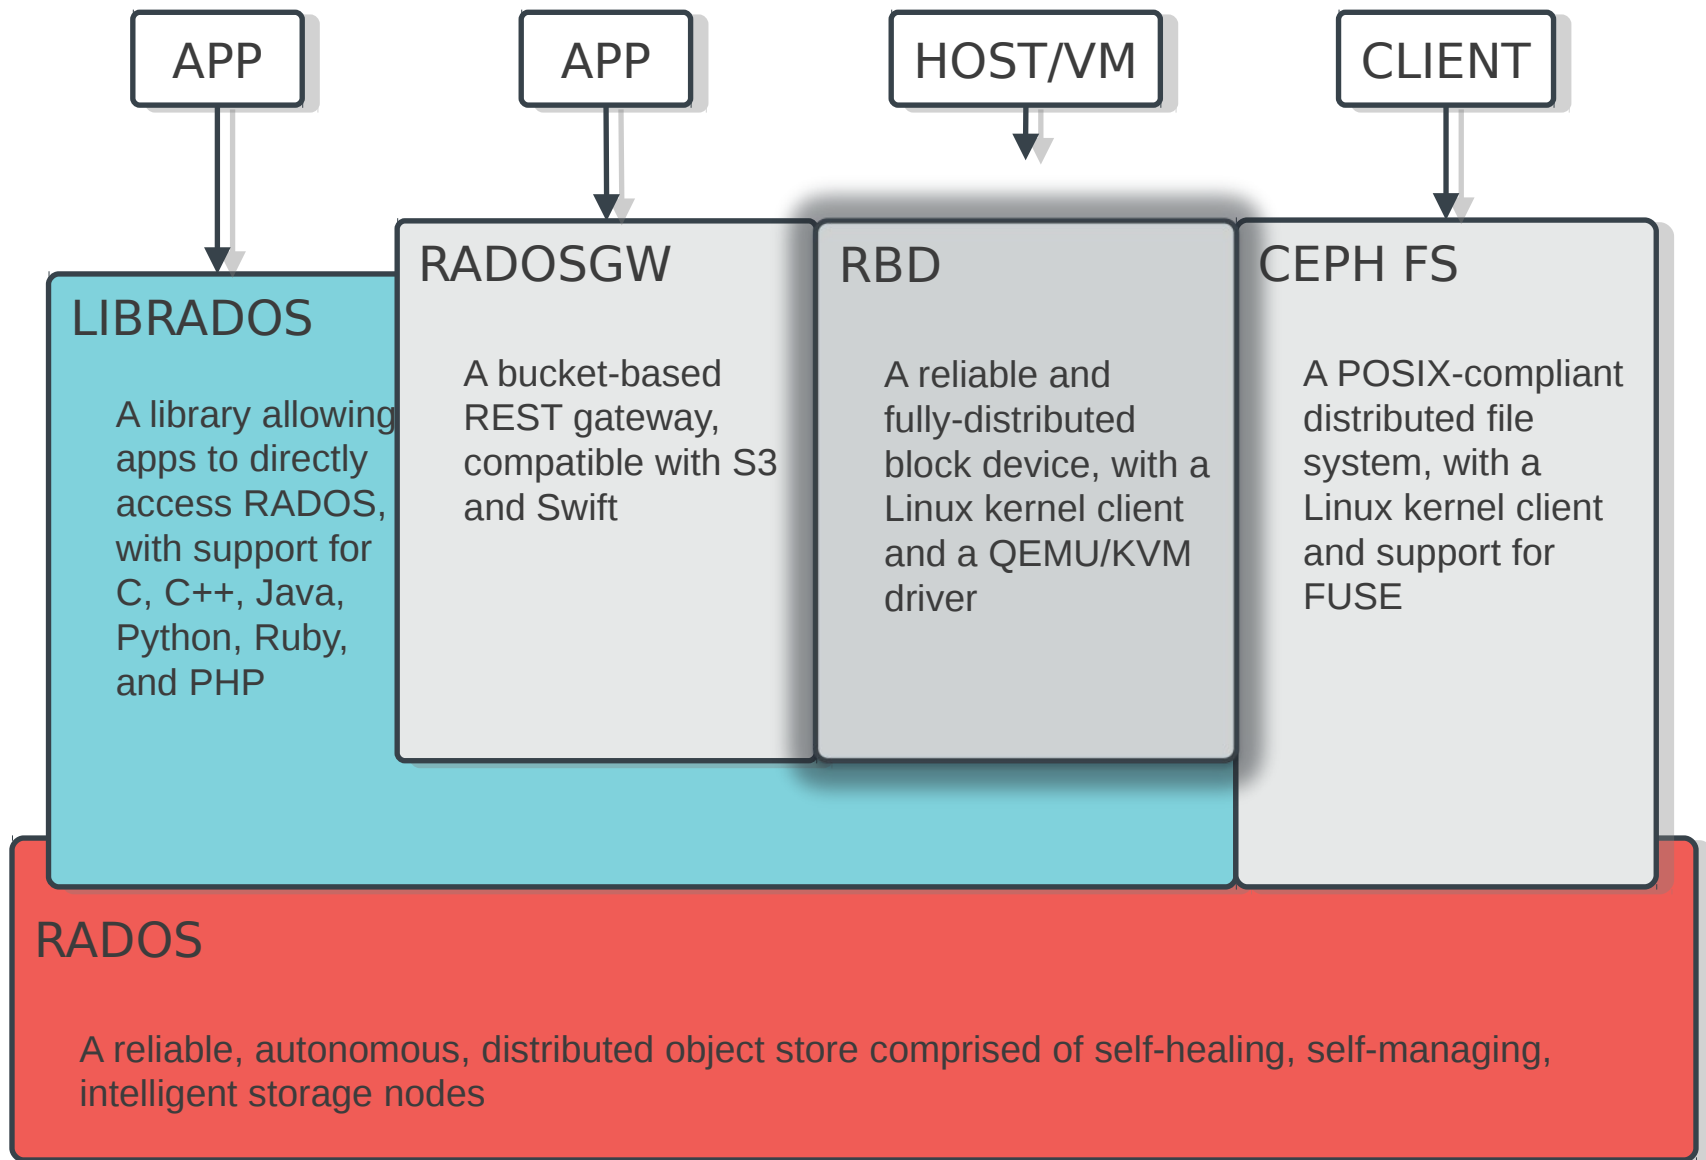




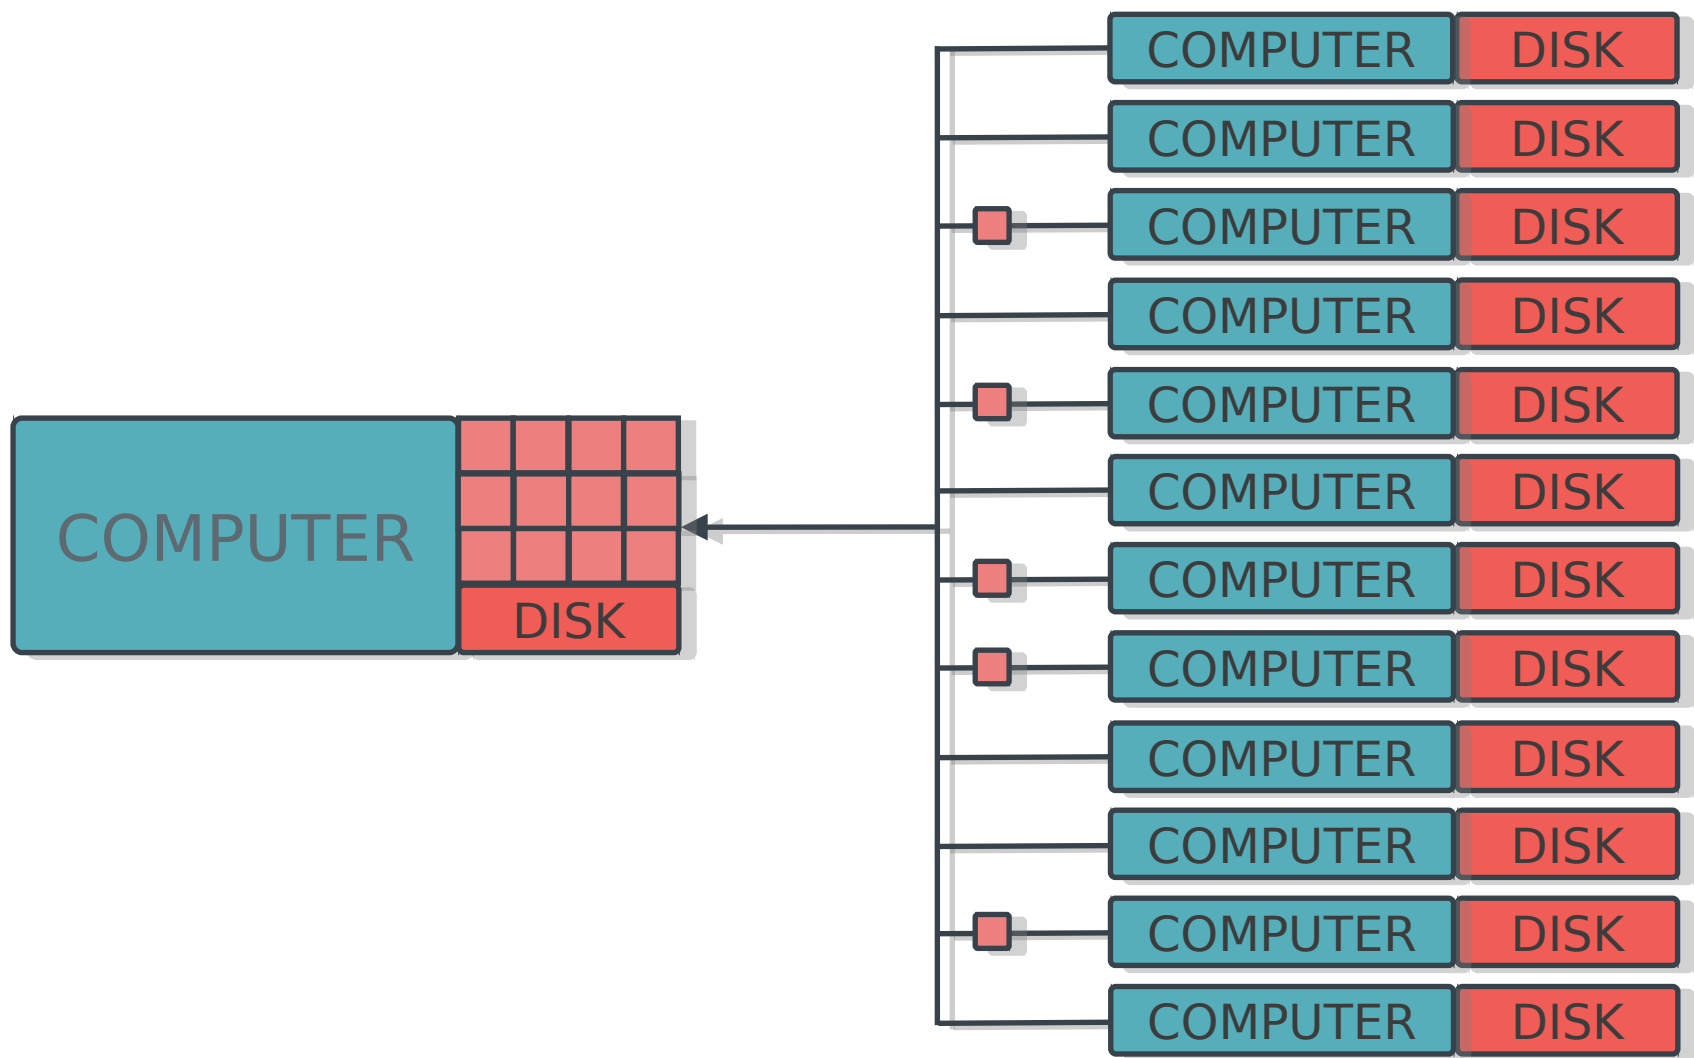


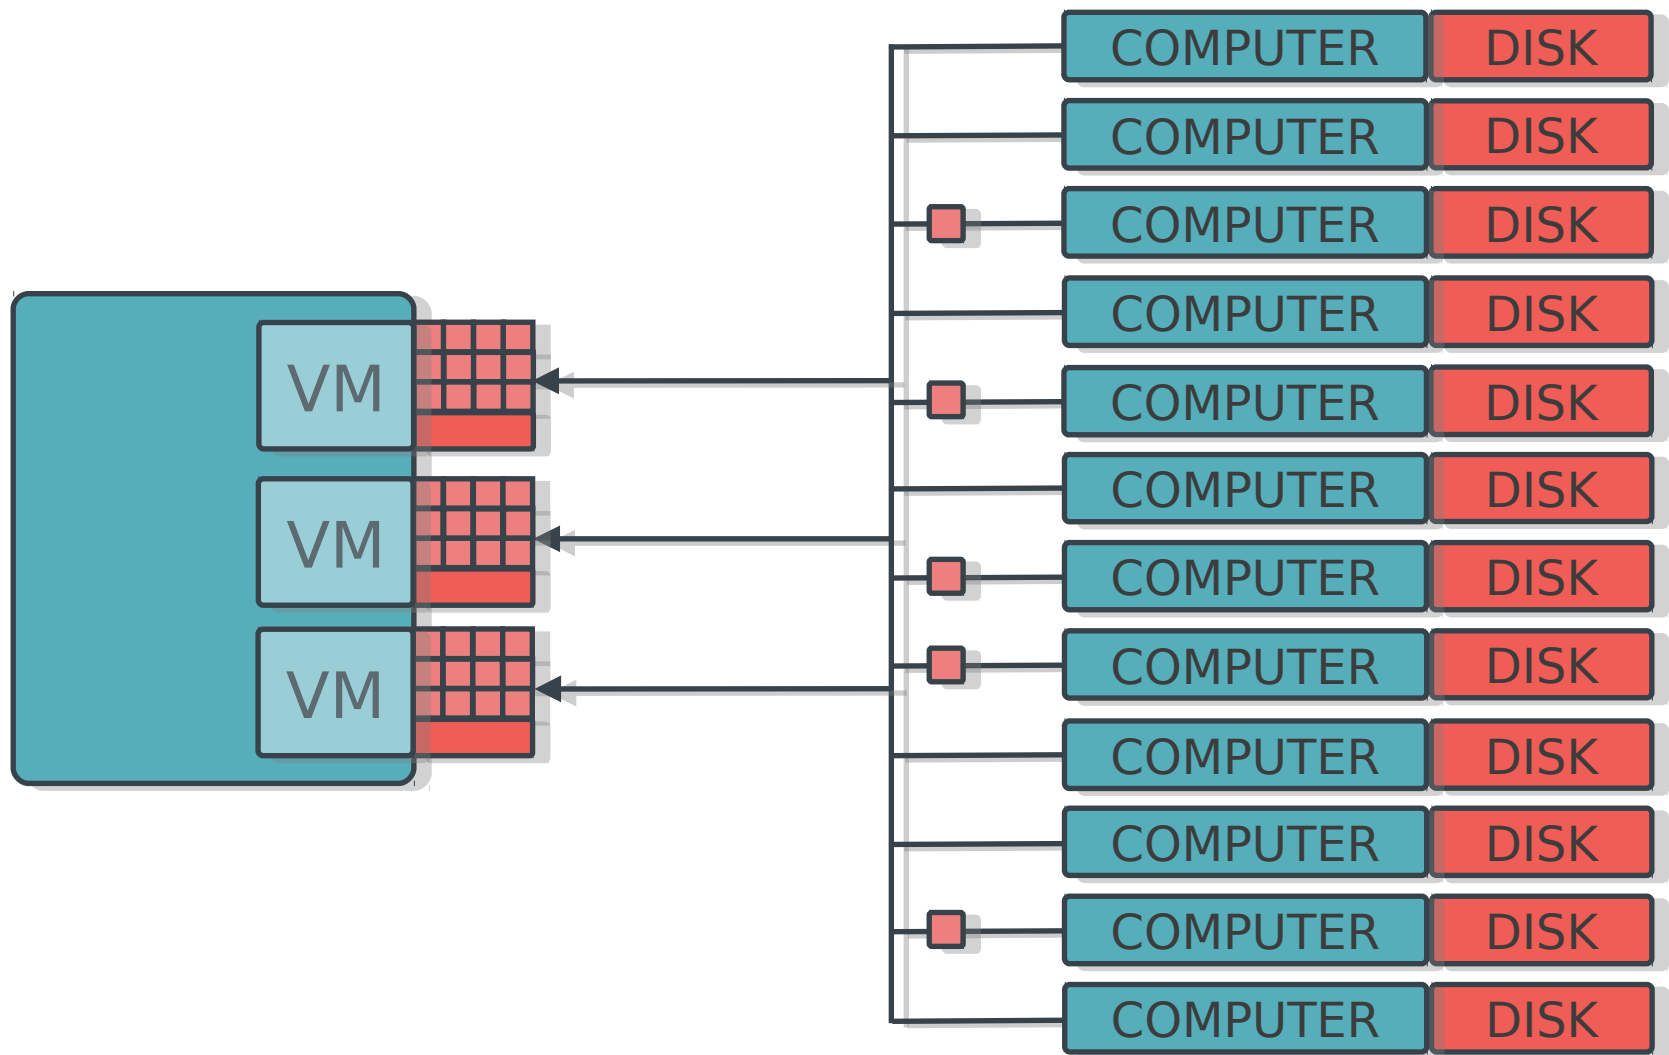
## RADOS Gateway

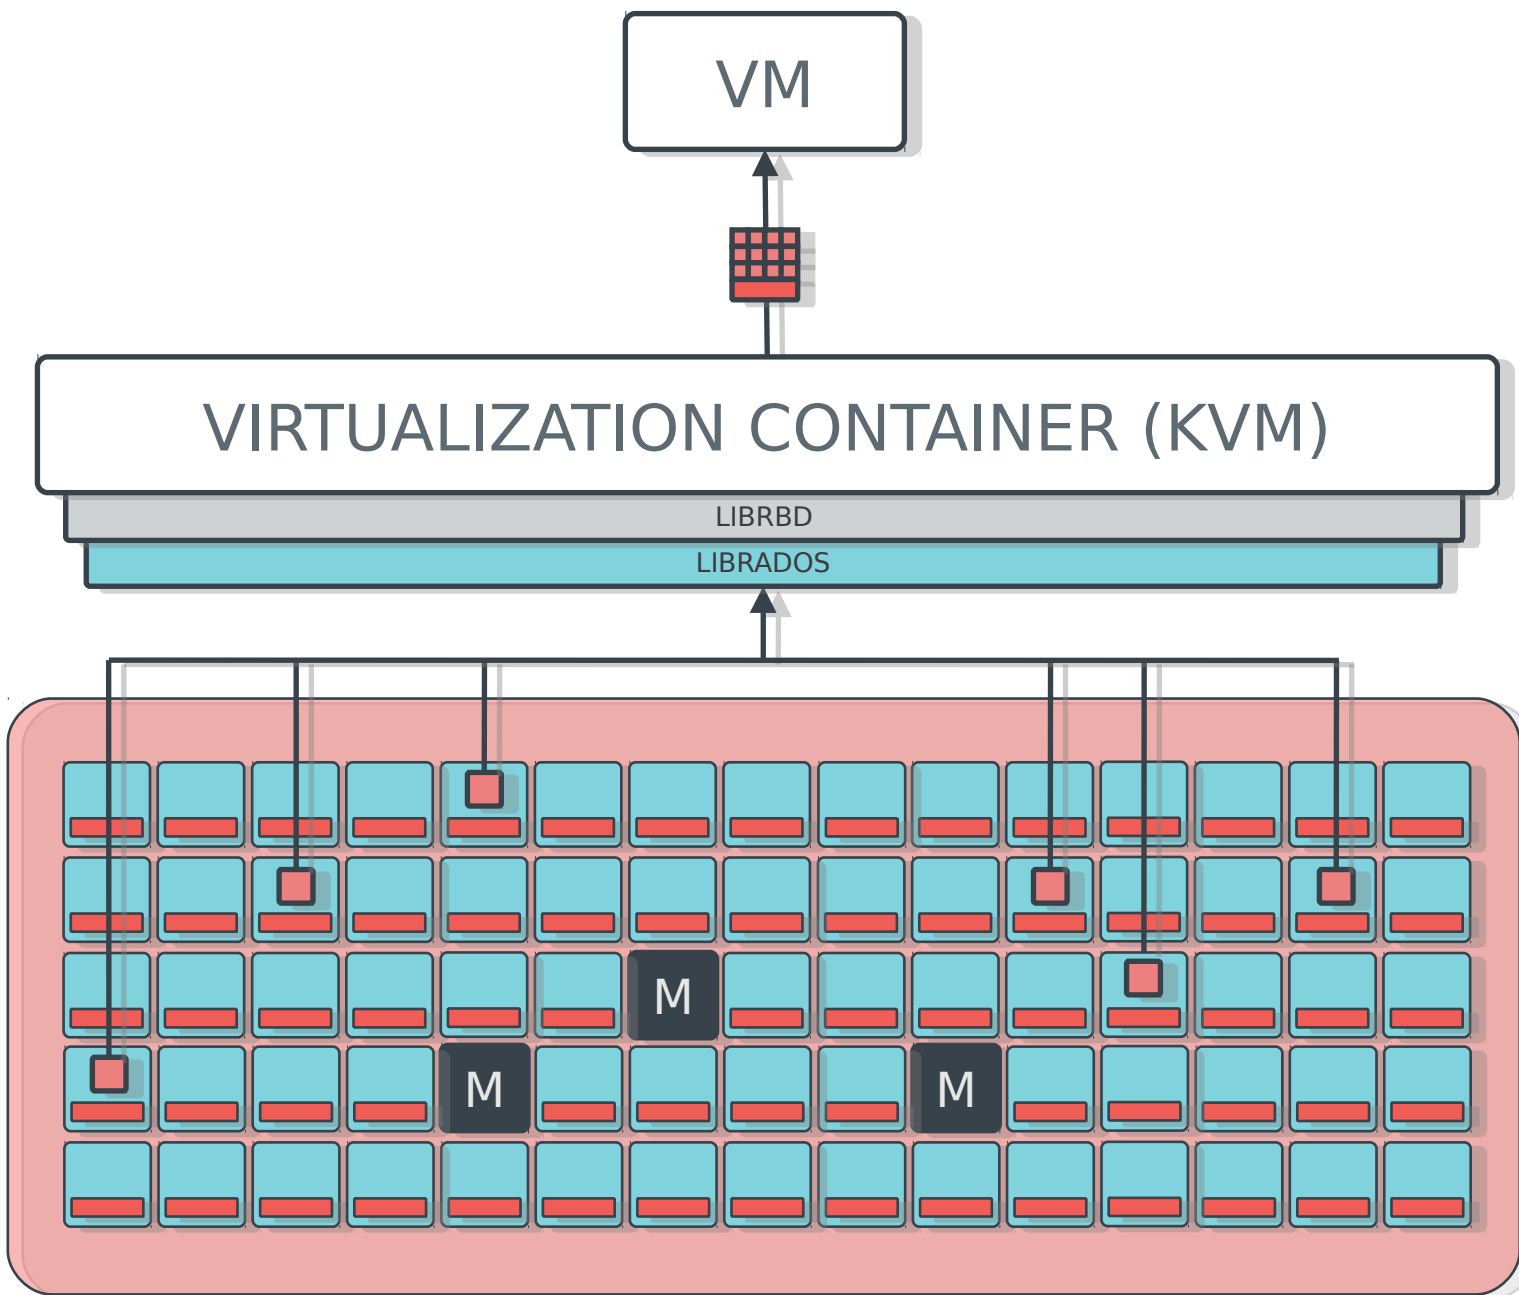
- REST-based object storage proxy
- uses RADOS to store objects
- API supports buckets, accounting
- usage accounting for billing purposes
- compatible with S3, Swift APIs

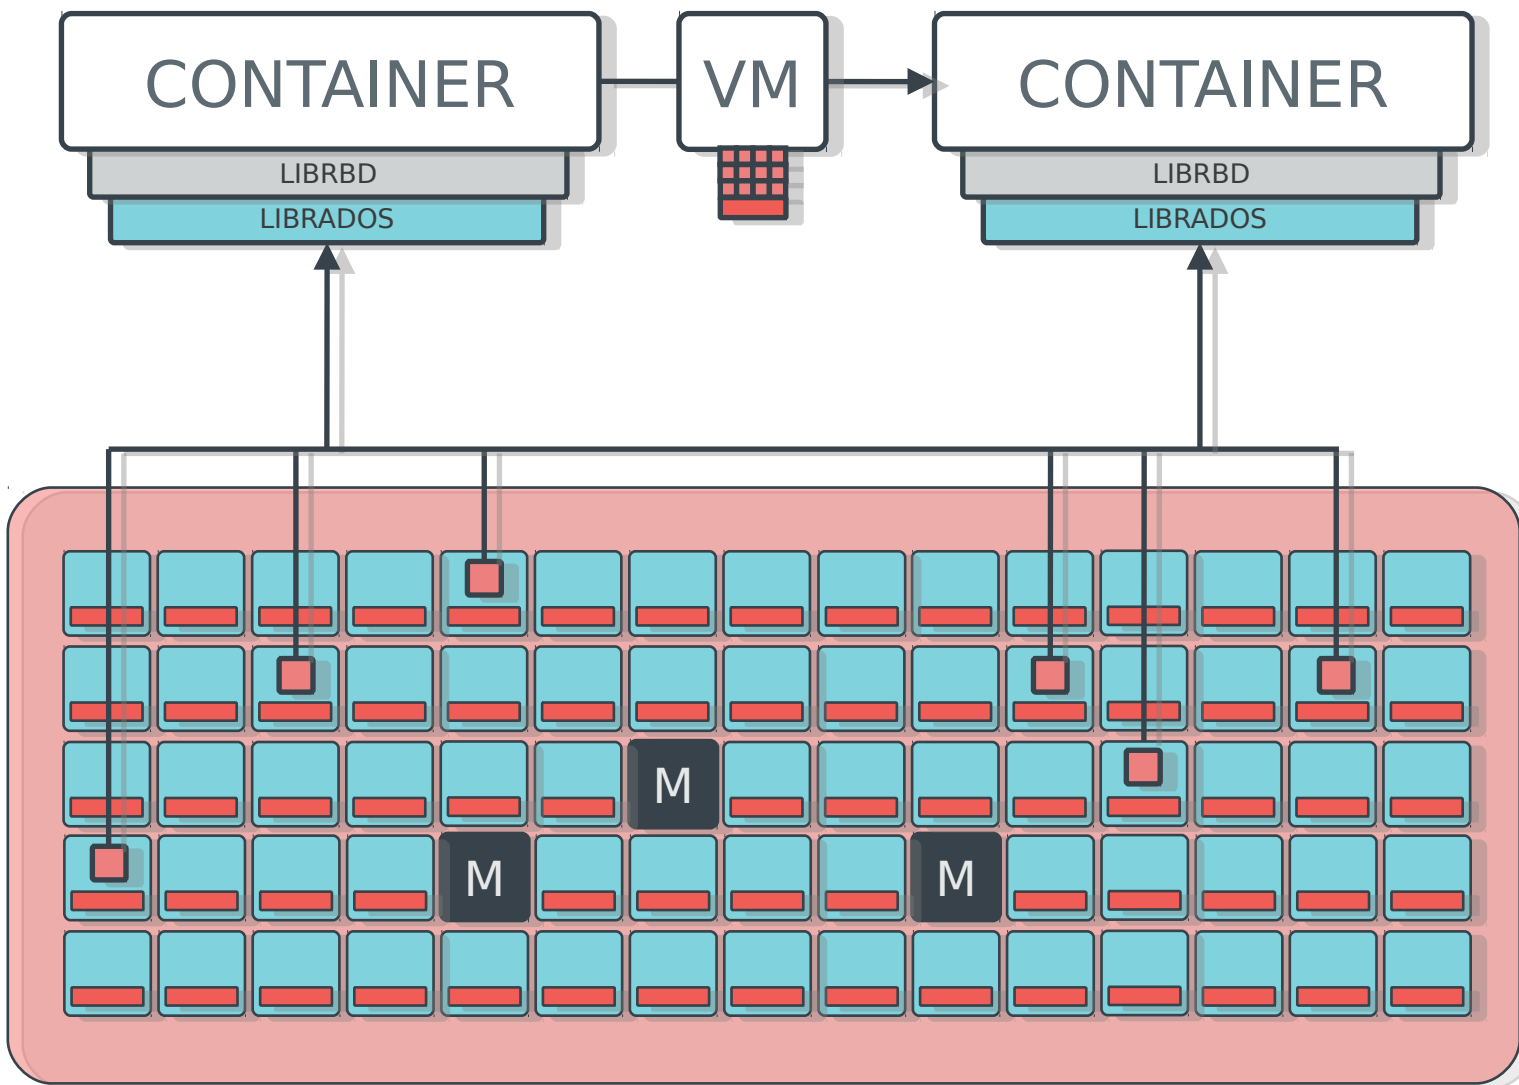


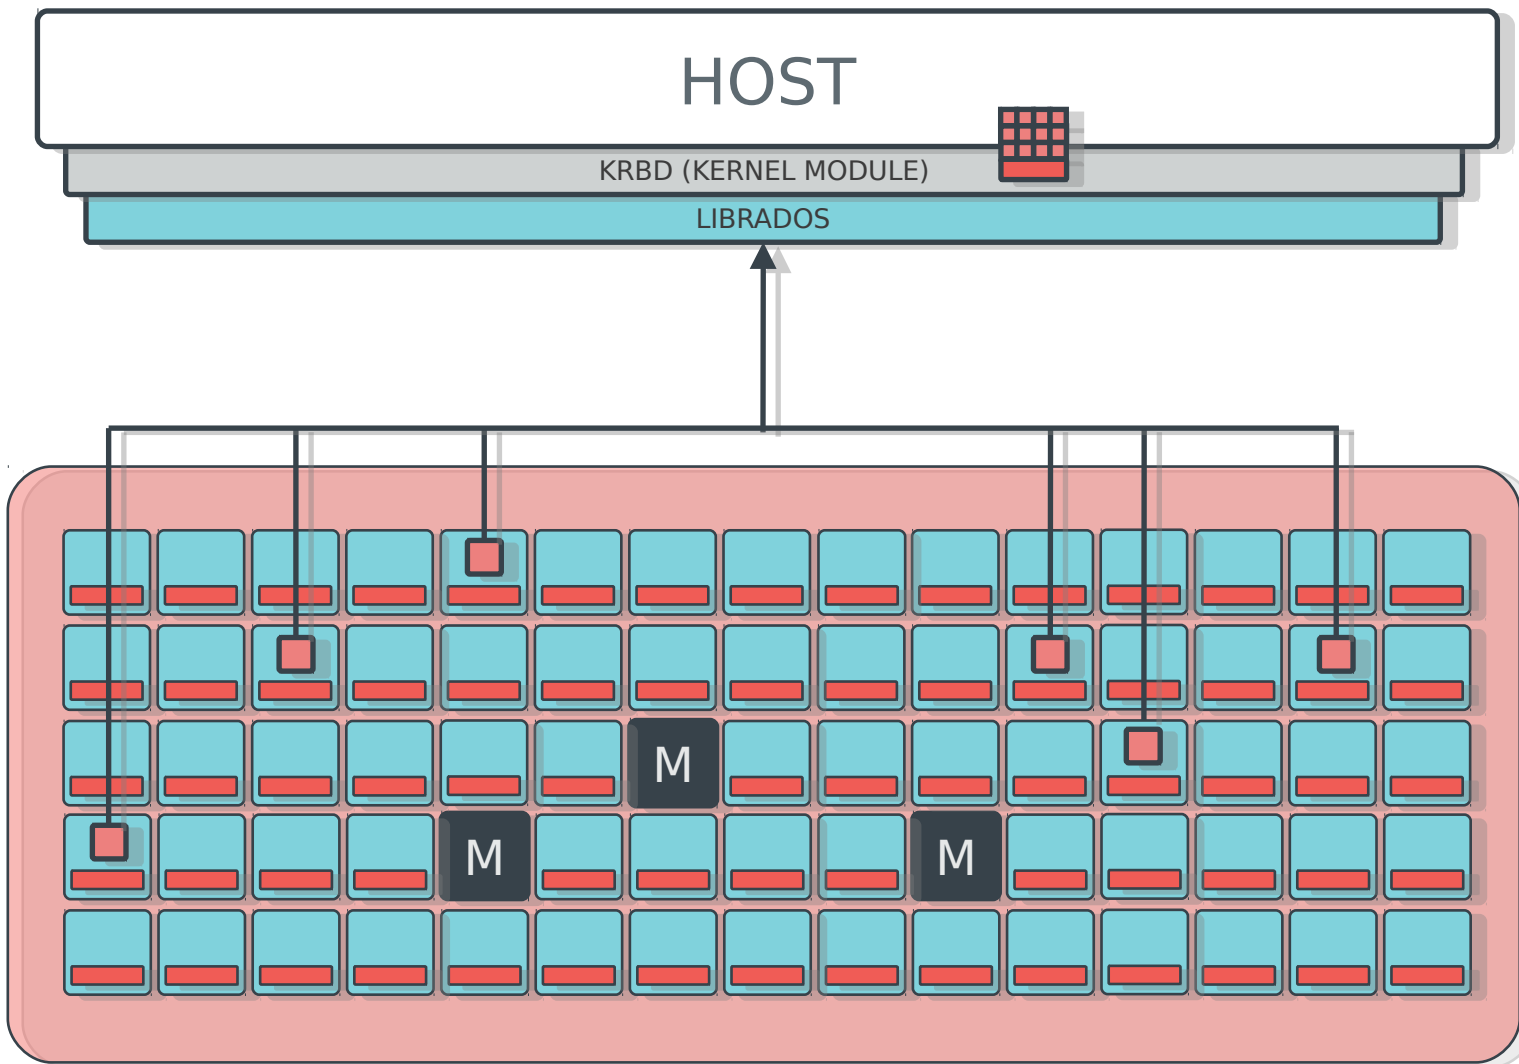


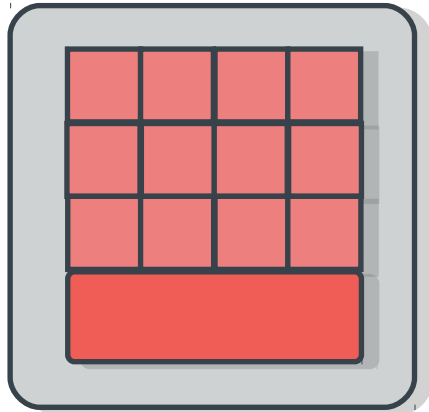






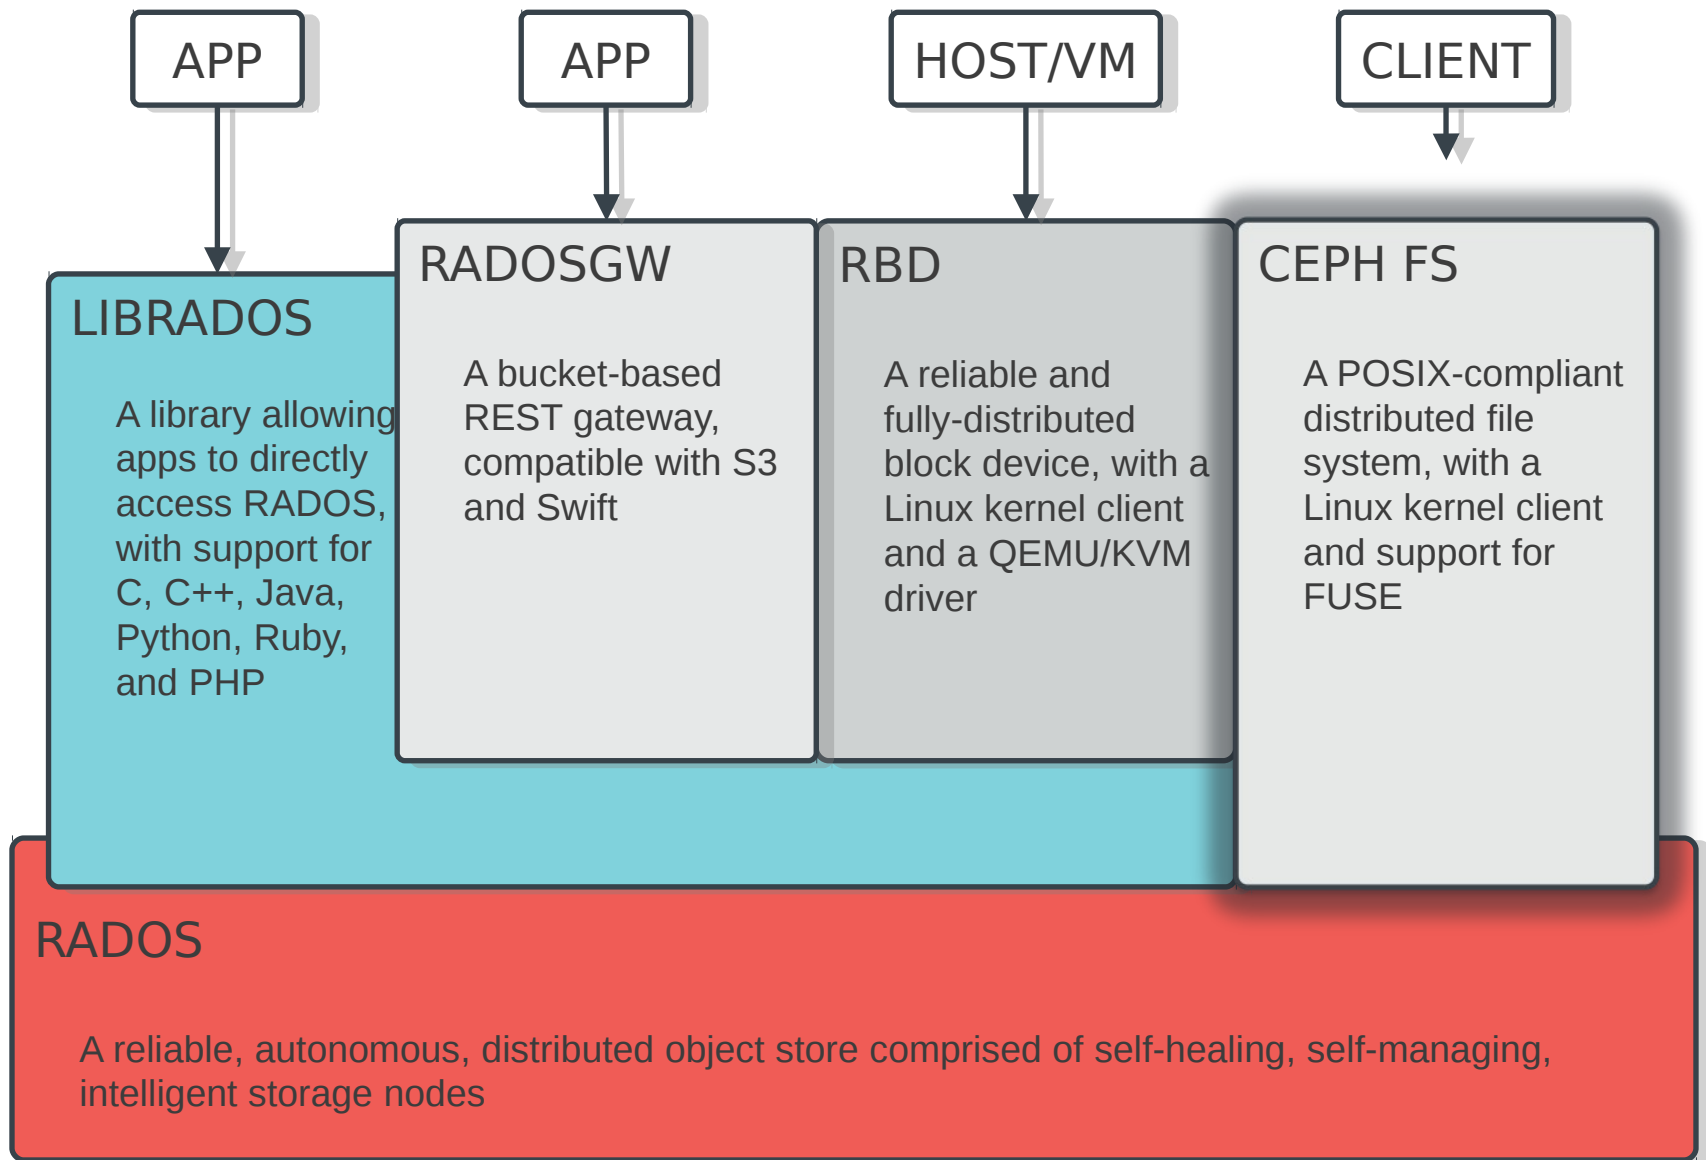


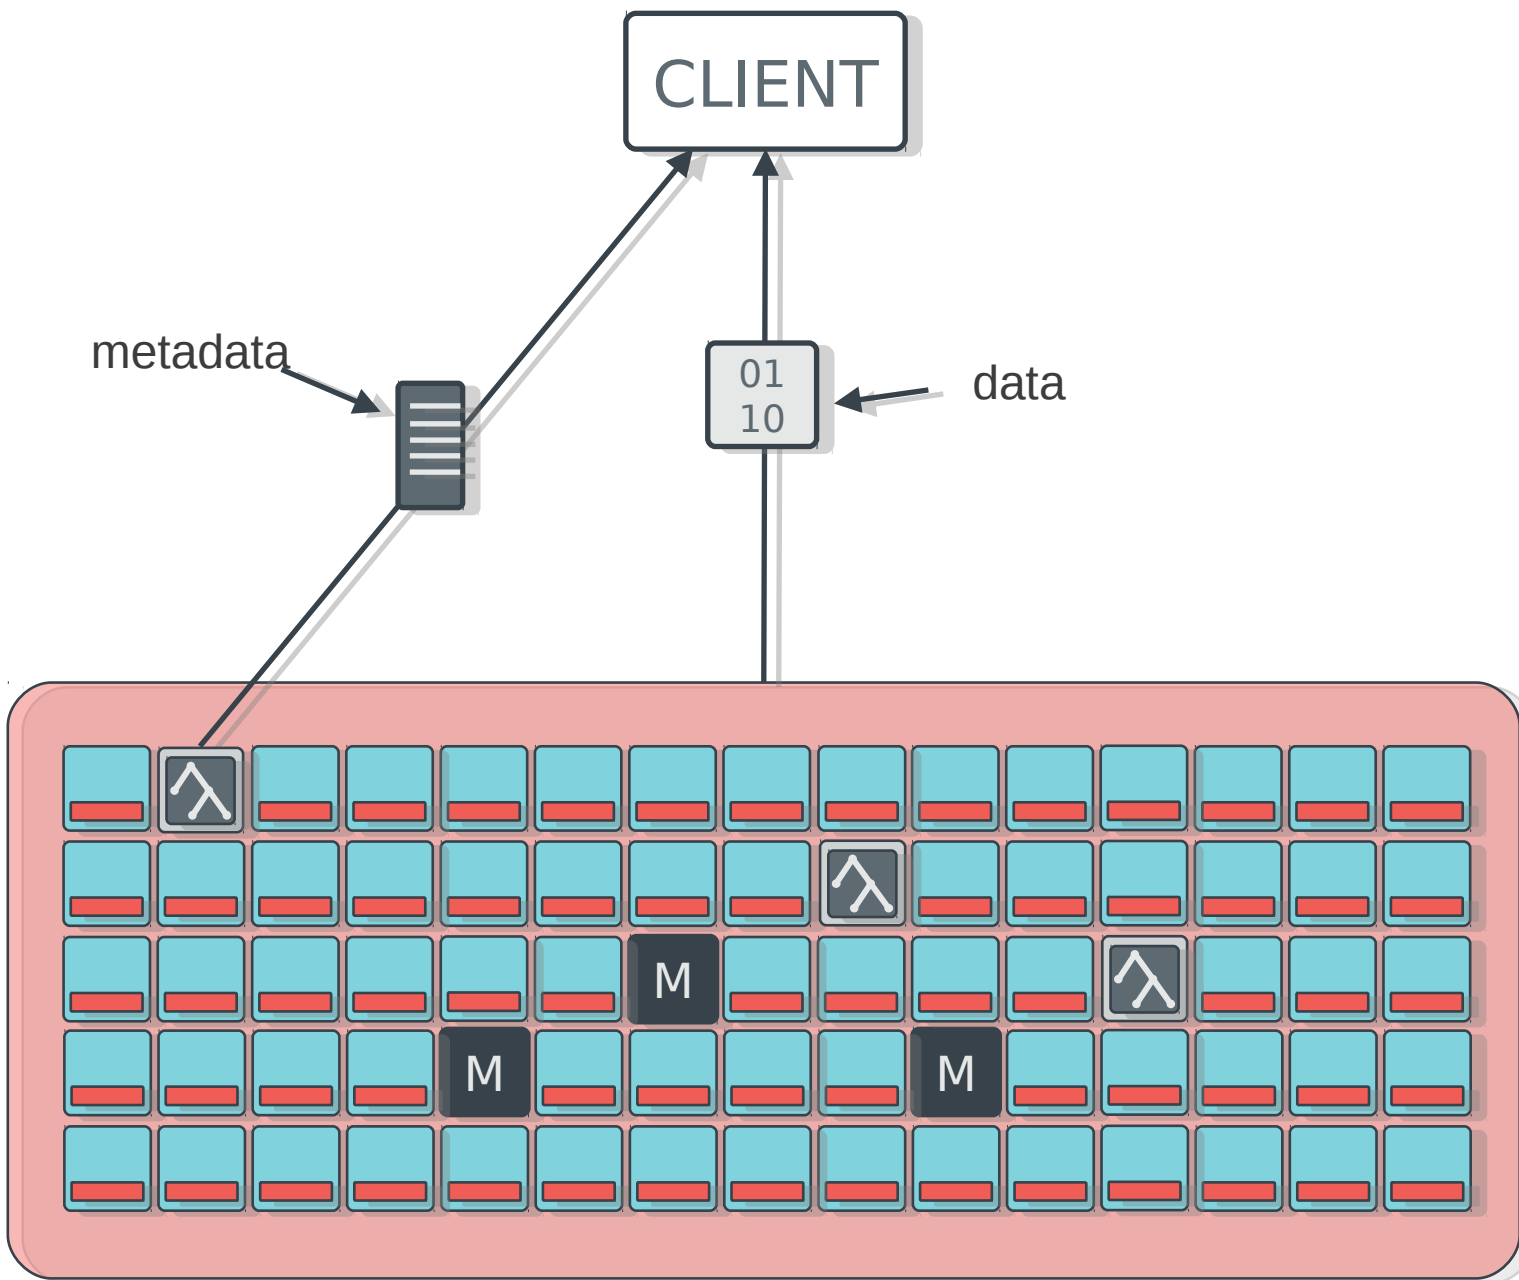




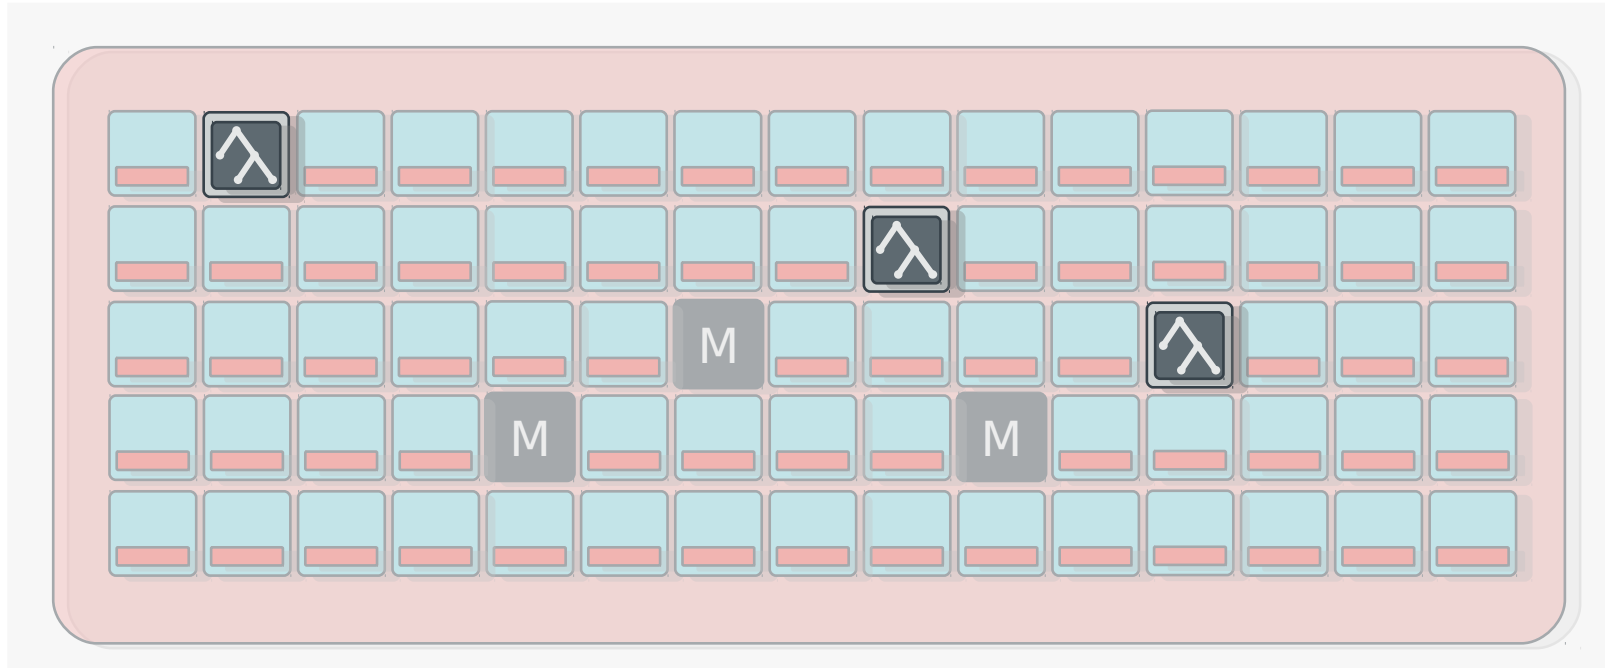
## RADOS Block Device

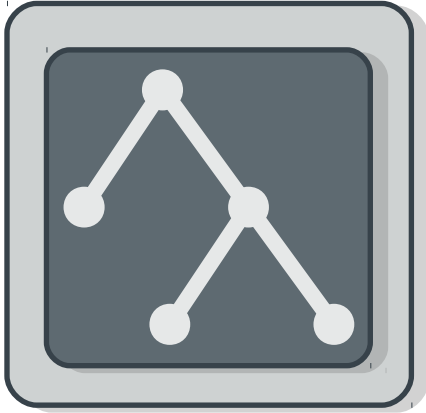
- storage of disk images in RADOS
- decouple VM from host
- images striped across entire cluster (pool)
- snapshots
- copy-on-write clones
- support in
  - Qemu/KVM
  - mainline Linux kernel (2.6.39+)
  - OpenStack, CloudStack





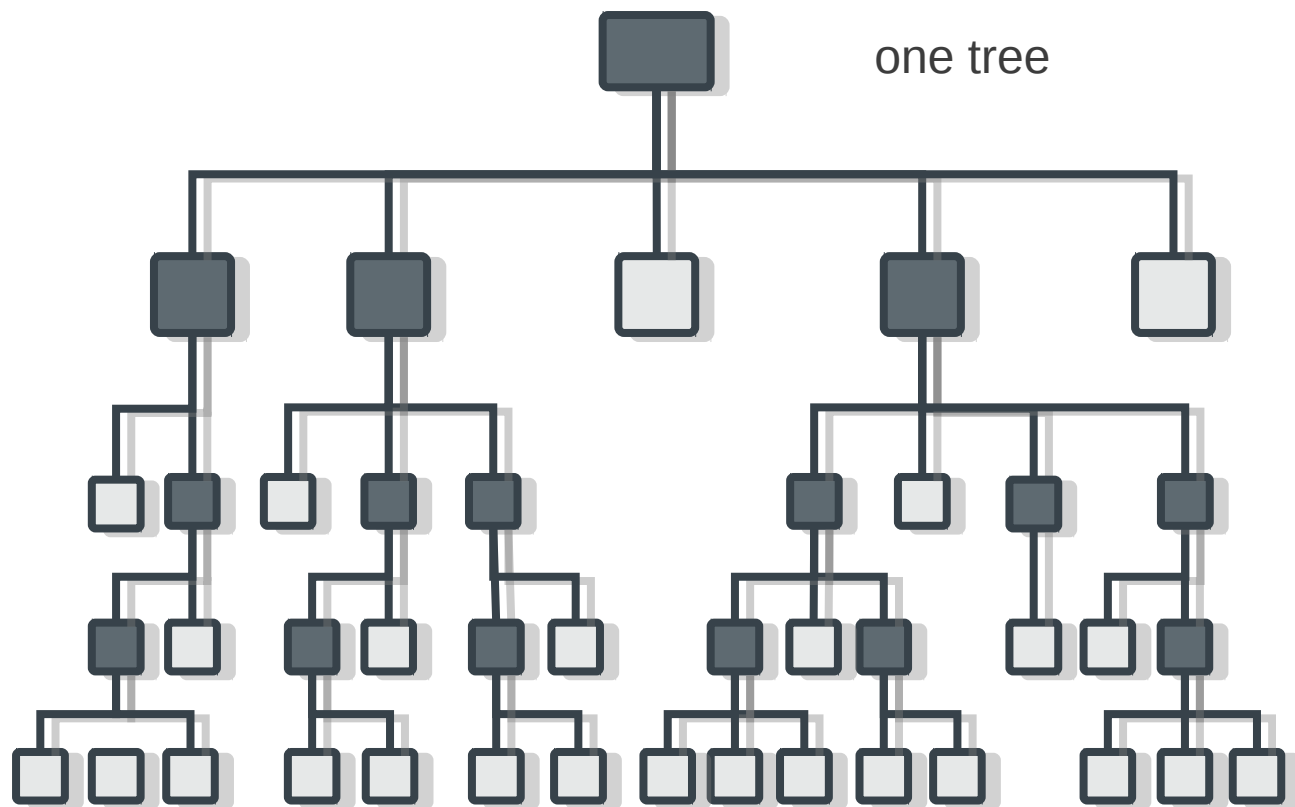




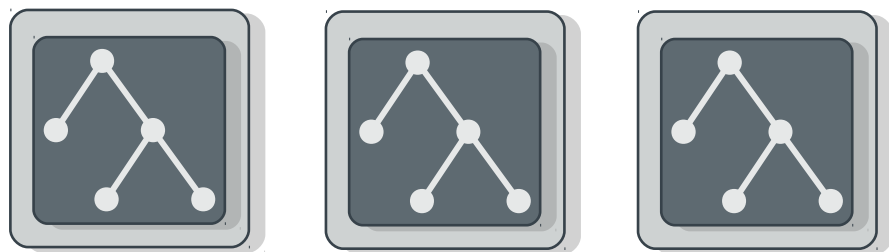


## Metadata Server (MDS)

- manages metadata for POSIX shared file system
  - directory hierarchy
  - file metadata (size, owner, timestamps)
- stores metadata in RADOS
- does not serve file data to clients
- only required for the shared file system

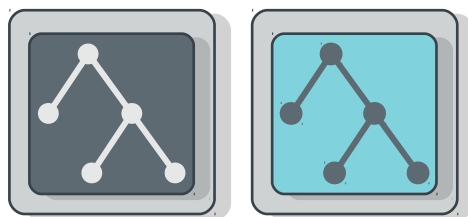
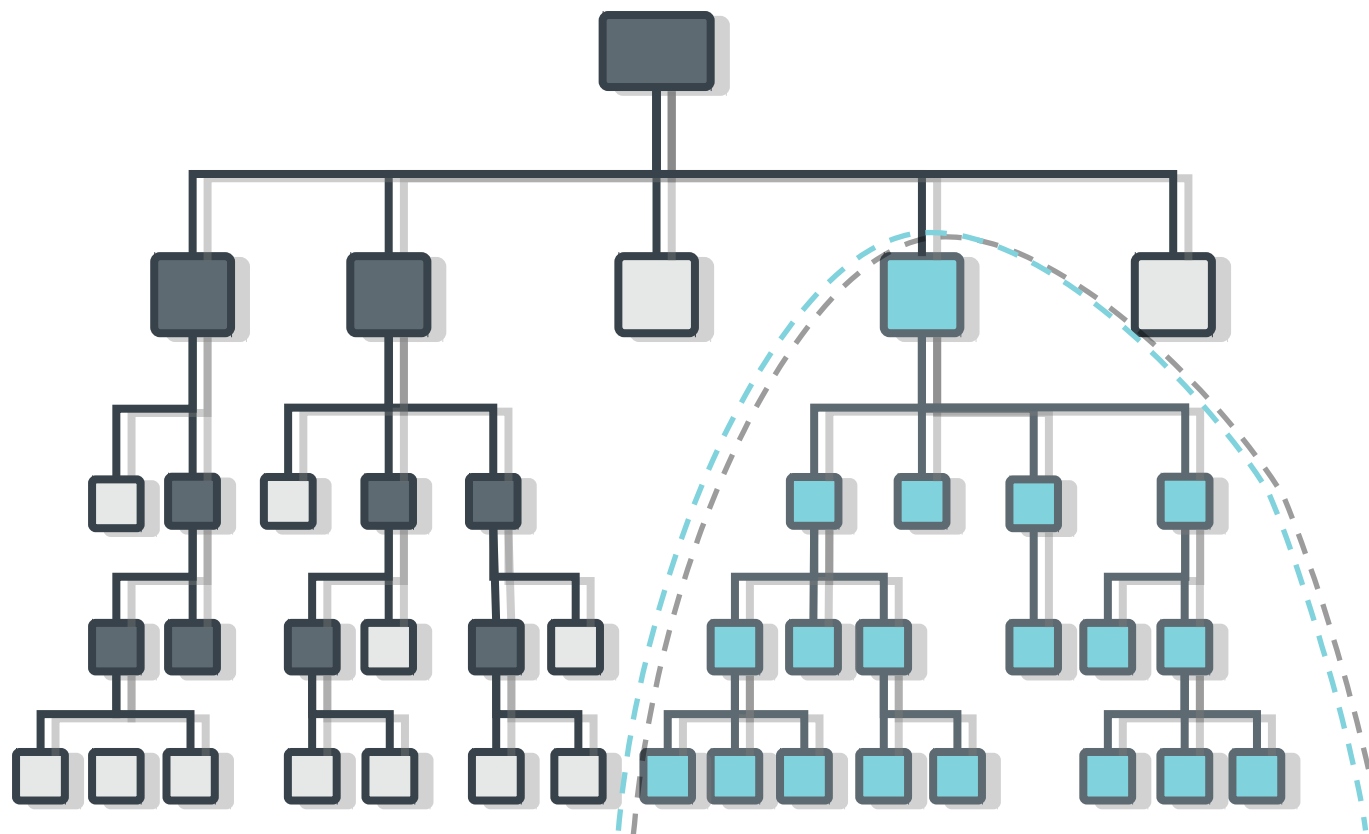


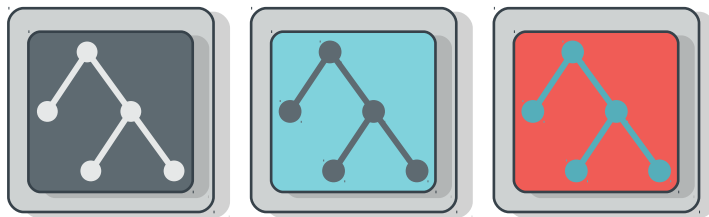
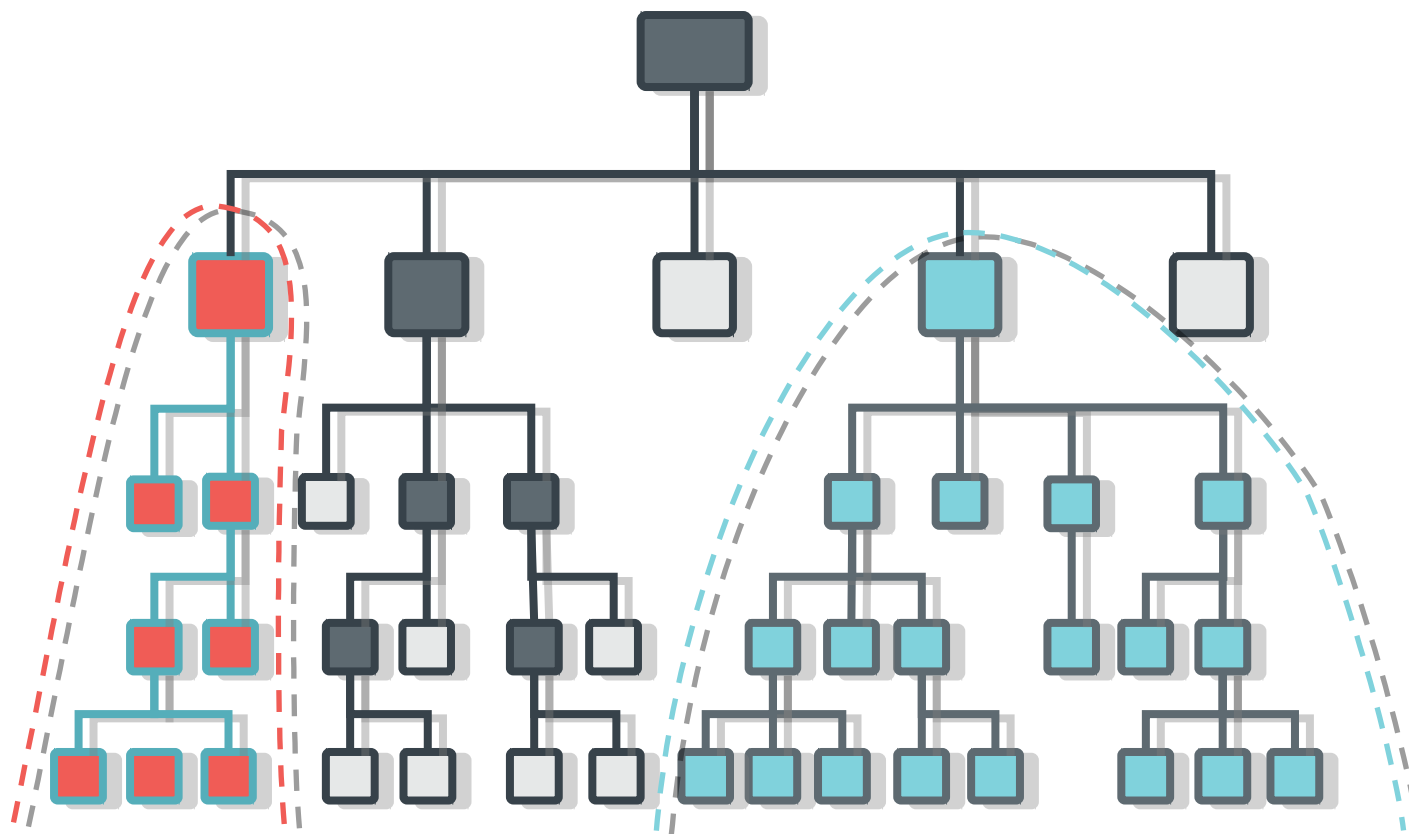
three metadata servers

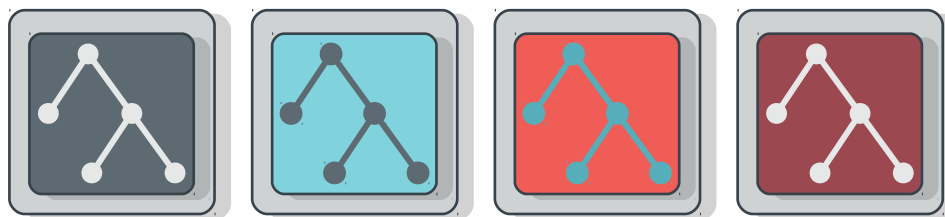


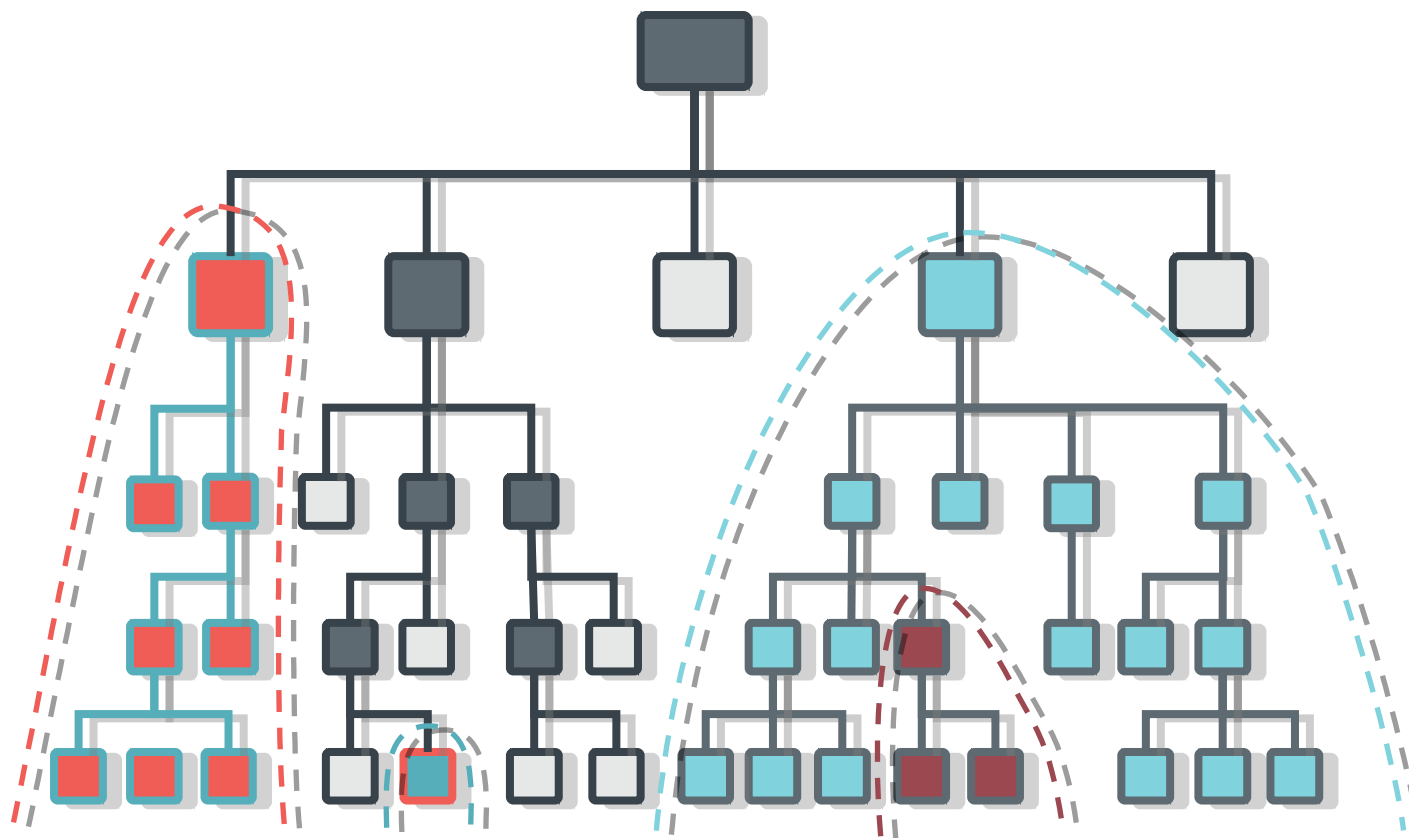
??











DYNAMIC SUBTREE PARTITIONING



# recursive accounting

- ceph-mds tracks recursive directory stats
  - file sizes
  - file and directory counts
  - modification time
- virtual xattrs present full stats
- efficient

```
$ ls -alSh | head
total 0
drwxr-xr-x 1 root      root      9.7T 2011-02-04 15:51 .
drwxr-xr-x 1 root      root      9.7T 2010-12-16 15:06 ..
drwxr-xr-x 1 pomceph   pg4194980 9.6T 2011-02-24 08:25 pomceph
drwxr-xr-x 1 mcg_test1 pg2419992 23G  2011-02-02 08:57 mcg_test1
drwx--x--- 1 luk0      adm       19G  2011-01-21 12:17 luk0
drwx--x--- 1 eest      adm       14G  2011-02-04 16:29 eest
drwxr-xr-x 1 mcg_test2 pg2419992 3.0G 2011-02-02 09:34 mcg_test2
drwx--x--- 1 fuzyceph  adm       1.5G 2011-01-18 10:46 fuzyceph
drwxr-xr-x 1 dallasceph pg275     596M 2011-01-14 10:06 dallasceph
```

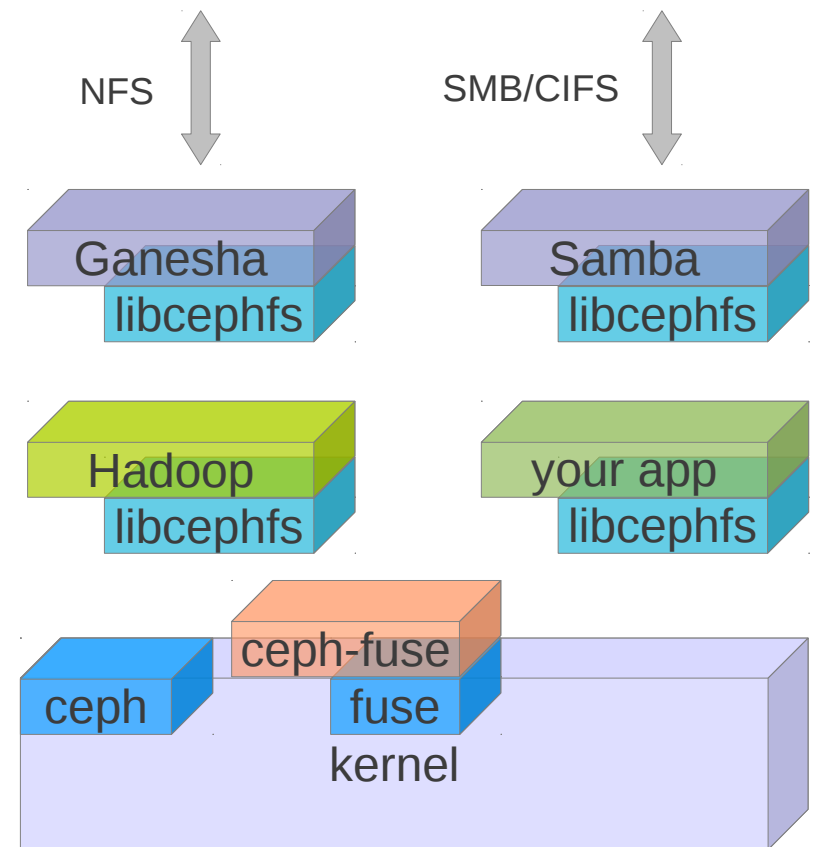
# snapshots

- volume or subvolume snapshots unusable at petabyte scale
  - snapshot arbitrary subdirectories
- simple interface
  - hidden '.snap' directory
  - no special tools

```
$ mkdir foo/.snap/one      # create snapshot
$ ls foo/.snap
one
$ ls foo/bar/.snap
_one_1099511627776        # parent's snap name is mangled
$ rm foo/myfile
$ ls -F foo
bar/
$ ls -F foo/.snap/one
myfile  bar/
$ rmdir foo/.snap/one     # remove snapshot
```

# multiple protocols, implementations

- Linux kernel client
  - `mount -t ceph 1.2.3.4:/ /mnt`
  - export (NFS), Samba (CIFS)
- ceph-fuse
- libcephfs.so
  - your app
  - Samba (CIFS)
  - Ganesha (NFS)
  - Hadoop (map/reduce)



How do you deploy and manage all this stuff?

# installation

- work closely with distros
  - packages in Debian, Fedora, Ubuntu, OpenSUSE
- build releases for all distros, recent releases
  - e.g., point to our deb src and apt-get install
- allow mixed-version clusters
  - upgrade each component individually, each host or rack one-by-one
  - protocol feature bits, safe data type encoding
- facilitate testing, bleeding edge
  - automatic build system generates packages for all branches in git
  - each to test new code, push hot-fixes

# configuration

- minimize local configuration
  - logging, local data paths, tuning options
  - cluster state is managed centrally by monitors
- specify option via
  - config file (global or local)
  - command line
  - adjusted for running daemon
- flexible configuration management options
  - global synced/copied config file
  - generated by management tools
    - Chef, Juju, Crowbar
- be flexible

# provisioning

- embrace **dynamic** nature of the cluster
  - disks, hosts, rack may come online at any time
  - anything may fail at any time
- simple scriptable sequences
  - 'ceph osd create <uuid>' → allocate osd id
  - 'ceph-osd --mkfs -i <id>' → initialize local data dir
  - 'ceph osd crush ...' → add to crush map
- identify minimal amount of central coordination
  - monitor cluster membership/quorum
- provide hooks for external tools to do the rest

# old school HA

- deploy a pair of servers
- heartbeat
- move a floating IP between them
- clients contact same “server”, which floats
- cumbersome to configure, deploy
- the “client/server” model doesn't scale out



# make client, protocol cluster-aware

- clients learn topology of the distributed cluster
  - all OSDs, current ips/ports, data distribution map
- clients talk to the server they want
  - servers scale out, client traffic distributes too
- servers dynamically register with the cluster
  - when a daemon starts, it updates its address in the map
  - other daemons and clients learn map updates via gossip
- servers manage heartbeat, replication
  - no fragile configuration
  - no fail-over pairs

# ceph disk management

- label disks
  - GPT partition type (fixed uuid)
- udev
  - generate event when disk is added, on boot
- 'ceph-disk-activate /dev/sdb'
  - mount the disk in the appropriate location (/var/lib/ceph/osd/NNN)
  - possibly adjust cluster metadata about disk location (host, rack)
- upstart, sysvinit, ...
  - start the daemon
  - daemon “joins” the cluster, brings itself online
- **no manual per-node configuration**

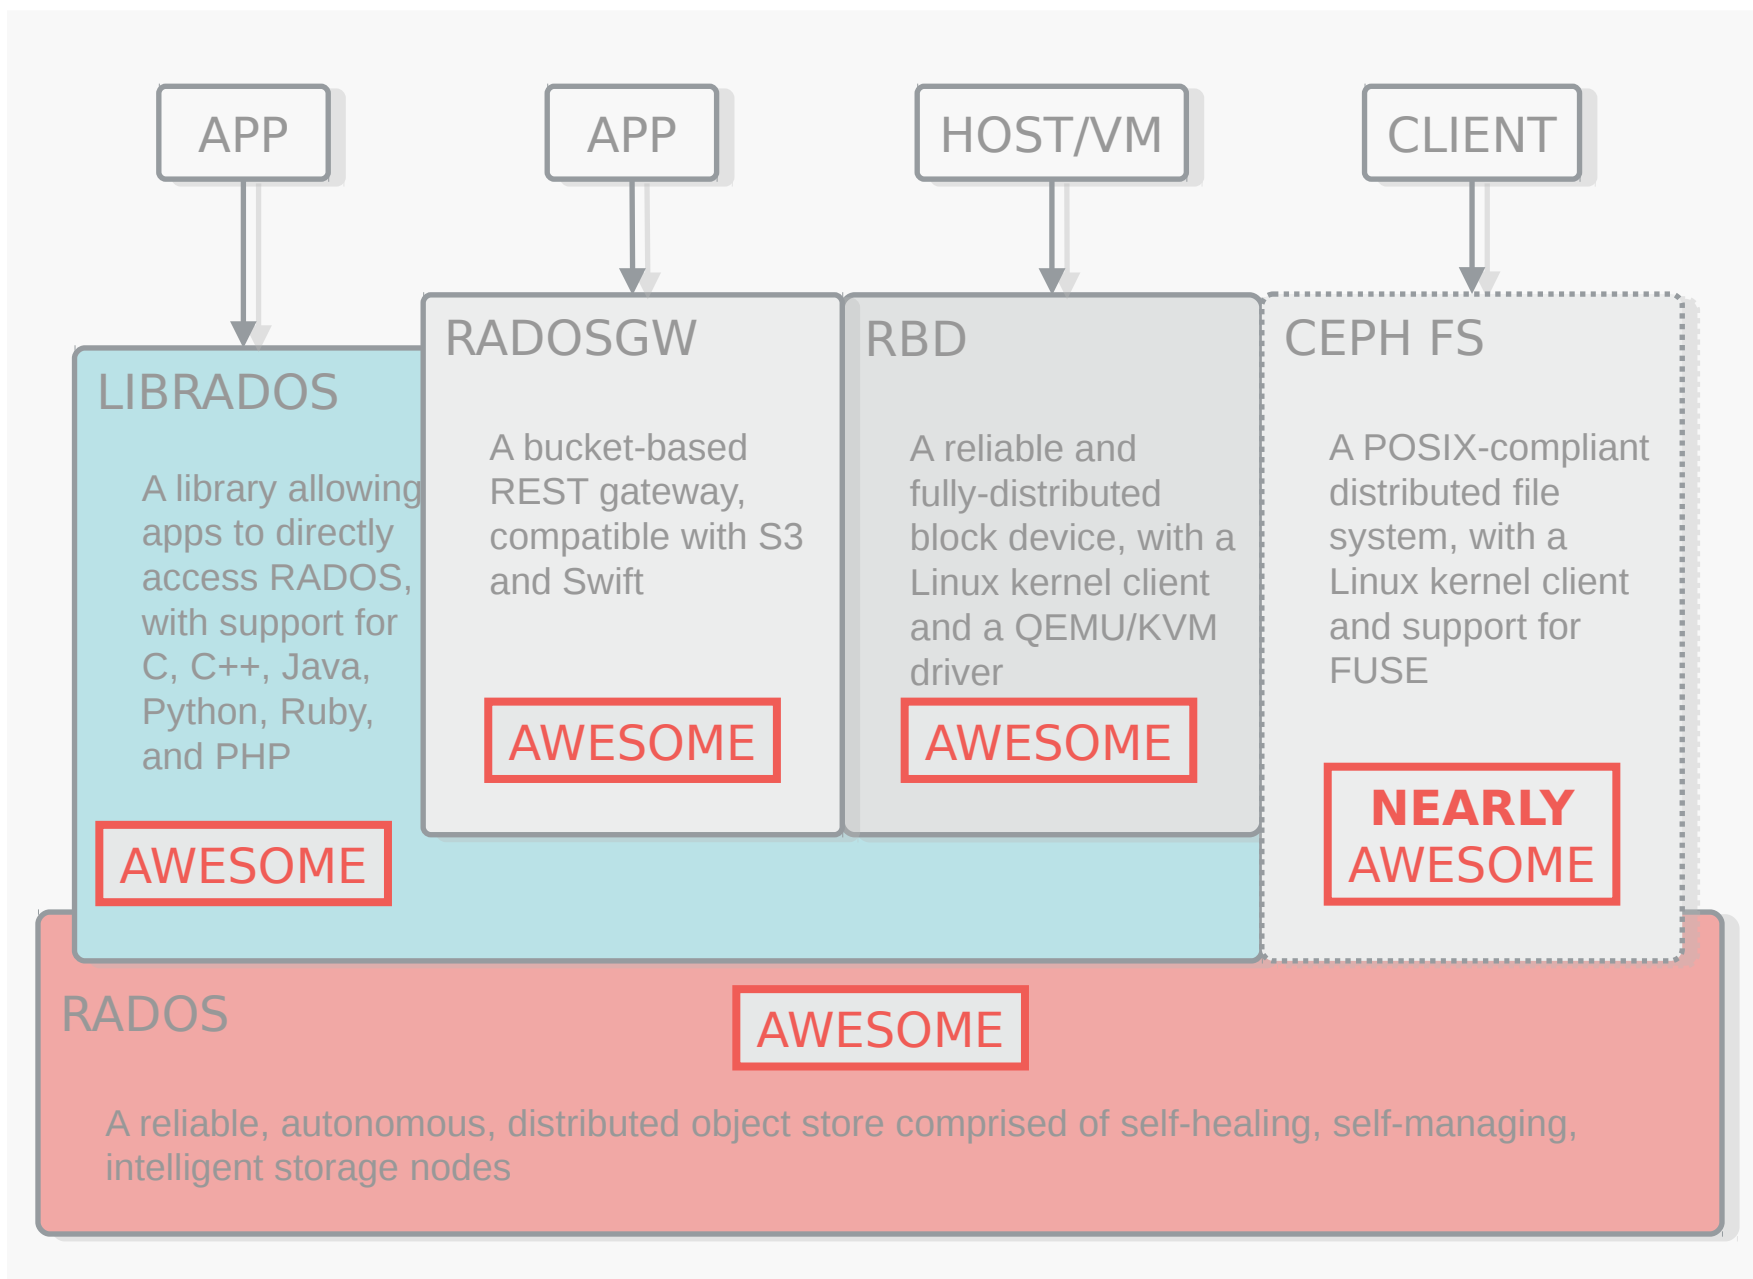
# distributed system health

- ceph-mon collects, aggregates basic system state
  - daemon up/down, current IPs
  - disk utilization
- simple hooks query system health
  - 'ceph health [detail]'
  - trivial plugins for nagios, etc.
- daemons instrumentation
  - query state of running daemon
  - use external tools to aggregate
    - collectd, graphite
    - statsd

# many paths

- do it yourself
  - use ceph interfaces directly
- Chef
  - ceph cookbooks, DreamHost cookbooks
- Crowbar
- Juju
- pre-packaged solutions
  - Piston AirFrame, Dell OpenStack-Powered Cloud Solution, Suse Cloud, etc.

# Project status and roadmap



# current status

- **argonaut** stable release v0.48
  - rados, RBD, radosgw
- **bobtail** stable release v0.56
  - RBD cloning
  - improved performance, scaling, failure behavior
  - radosgw API, performance improvements
  - release in 1-2 weeks

# cuttlefish roadmap

- file system
  - pivot in engineering focus
  - CIFS (Samba), NFS (Ganesha), Hadoop
- RBD
  - Xen integration, iSCSI
- radosgw
  - multi-side federation, disaster recovery
- RADOS
  - geo-replication, disaster recovery
  - ongoing performance improvements



# why we do this

- limited options for scalable open source storage
- proprietary solutions
  - expensive
  - don't scale (well or out)
  - marry hardware and software
- users hungry for alternatives
  - scalability
  - cost
  - features
  - open, interoperable

# licensing

*<yawn>*

- promote adoption
- enable community development
- prevent ceph from becoming proprietary
- allow organic commercialization

# LGPLv2

- “copyleft”
  - free distribution
  - allow derivative works
  - changes you distribute/sell must be shared
- ok to link to proprietary code
  - allow proprietary products to include and build on ceph
  - does not allow proprietary derivatives of ceph

# fragmented copyright

- we do not require copyright assignment from contributors
  - no single person or entity owns all of ceph
  - no single entity can make ceph proprietary
- strong community
  - many players make ceph a safe technology bet
  - project can outlive any single business

# why it's important

- ceph is an ingredient
  - we need to play nice in a larger ecosystem
  - community will be key to success
- truly open source solutions are disruptive
  - frictionless integration with projects, platforms, tools
  - freedom to innovate on protocols
  - leverage community testing, development resources
  - open collaboration is efficient way to build technology

# who we are

- Ceph created at UC Santa Cruz (2004-2007)
- developed by DreamHost (2008-2011)
- supported by **Inktank** (2012)
  - Los Angeles, Sunnyvale, San Francisco, remote
- growing user and developer community
  - Linux distros, users, cloud stacks, SIs, OEMs

<http://ceph.com/>

# thanks

sage weil

[sage@inktank.com](mailto:sage@inktank.com)

@liewegas

<http://github.com/ceph>

<http://ceph.com/>





# why we like btrfs

- pervasive checksumming
- snapshots, **copy-on-write**
- efficient metadata (xattrs)
- inline data for small files
- transparent compression
- integrated volume management
  - software RAID, mirroring, error recovery
  - SSD-aware
- online fsck
- active development community