



## **Making Ceph Fast in the Face of Failure**

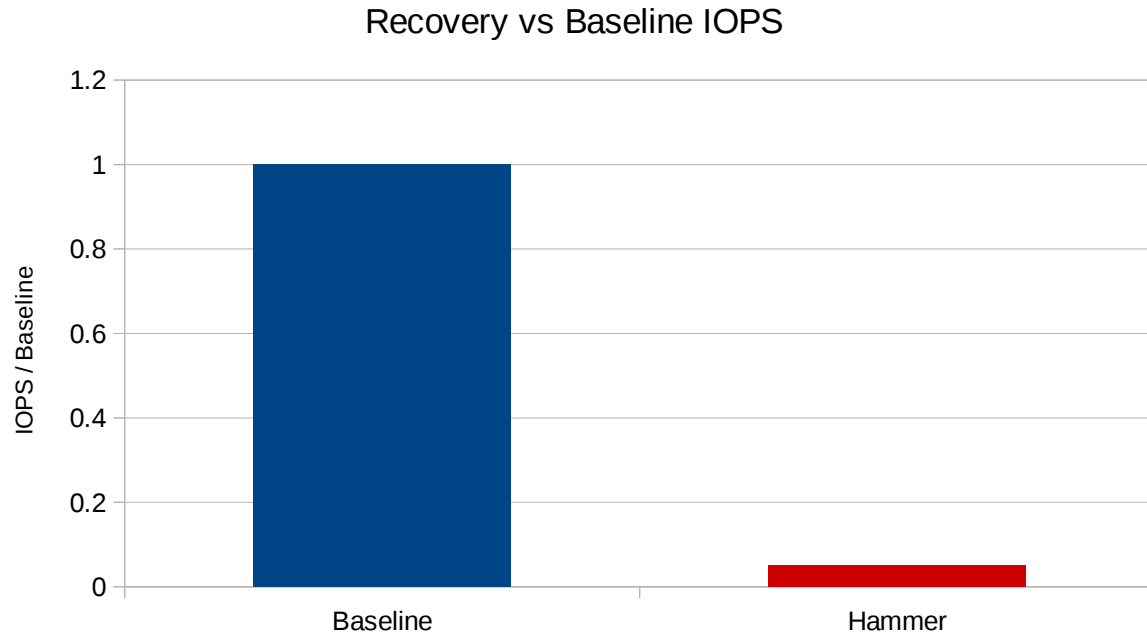
Josh Durgin and Neha Ojha  
2018.08.28

# The Dark Ages



The Fight Between Carnival and Lent by Pieter Bruegel the Elder

# Hammer

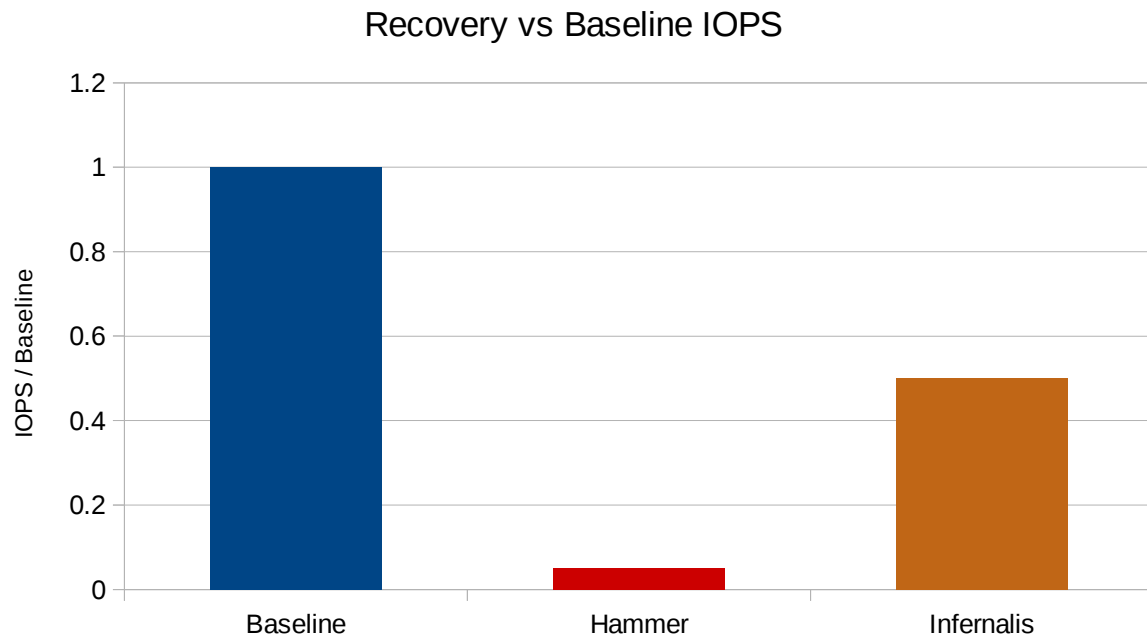


# Hammer

- Favored maximum recovery speed
- Default client impact was huge
- Tuning could help

```
osd max backfills = 10 → 1  
osd recovery max active = 15 → 3  
osd recovery op priority = 10 → 3  
osd recovery max single start = 5 → 1
```

# Infernalis



# Luminous – High-level Priority

Manual way to recover higher-level constructs, e.g. rbd images

pg force-recovery command

```
ceph pg force-recovery {pg-id} [{pg-id #2}] [{pg-id #3} ...]
```

```
ceph pg force-backfill {pg-id} [{pg-id #2}] [{pg-id #3} ...]
```

If you change your mind or prioritize wrong groups, use:

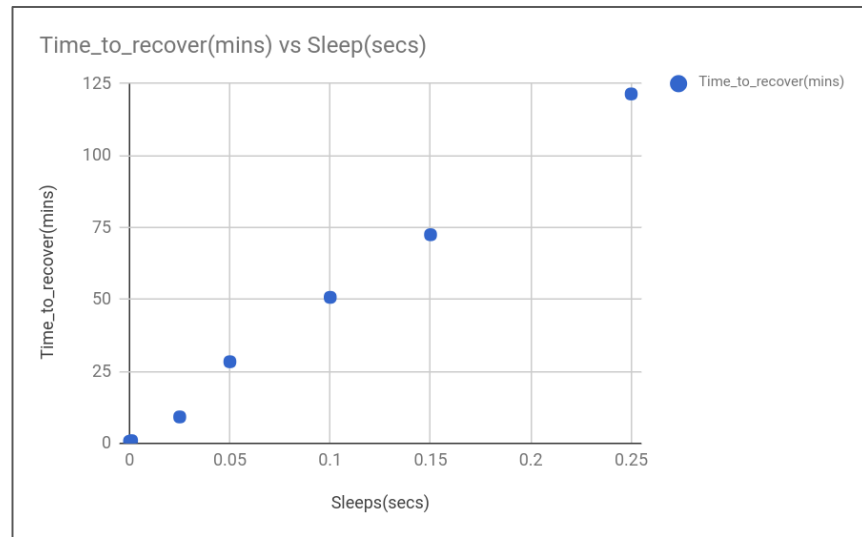

```
ceph pg cancel-force-recovery {pg-id} [{pg-id #2}] [{pg-id #3} ...]
```

```
ceph pg cancel-force-backfill {pg-id} [{pg-id #2}] [{pg-id #3} ...]
```

# Luminous – Throttling Recovery

- `osd_recovery_sleep`
- changing this will shift the balance between recovery and client I/O
- Different configurations based on underlying hardware

# Recovery Sleep - HDDs

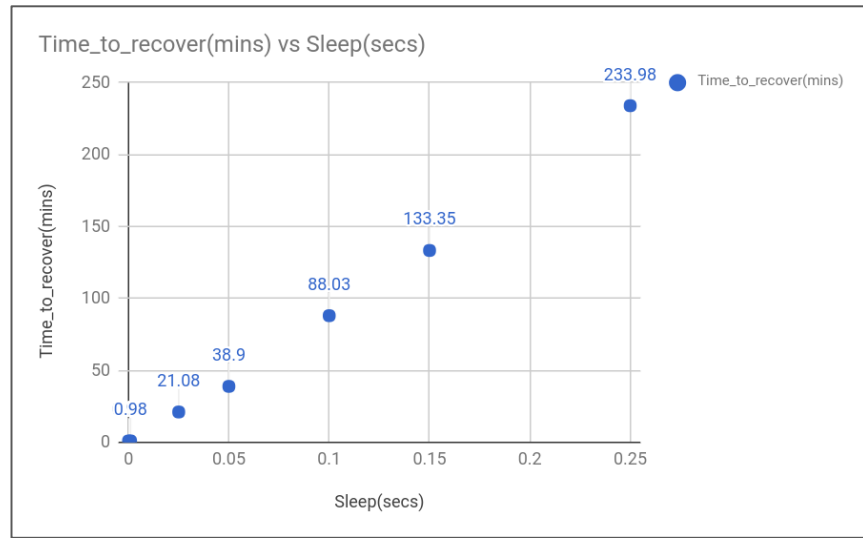
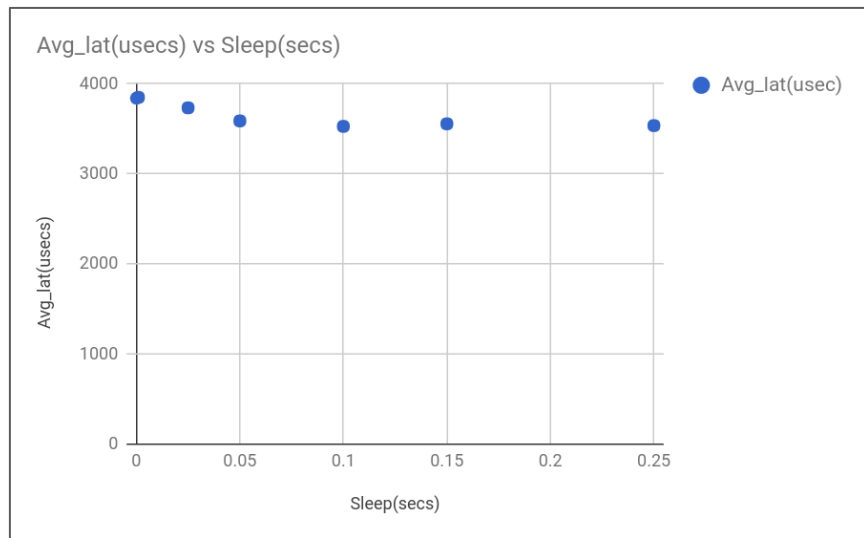


BlueStore on HDDs with Fio 4k random writes

➤ `osd_recovery_sleep_hdd` chosen as 0.1




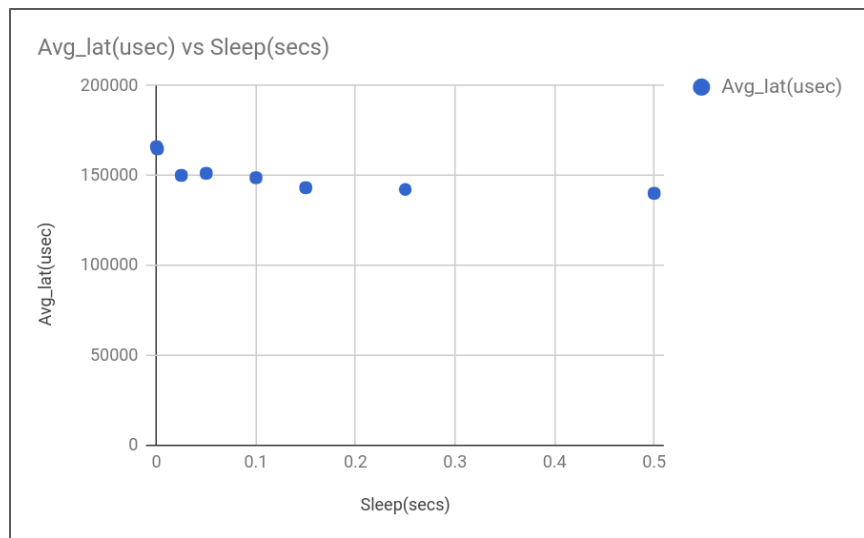
# Recovery Sleep - SSDs



BlueStore on SSDs with Fio 4k random writes

➤ `osd_recovery_sleep_ssd` chosen as 0


# Recovery Sleep - Hybrid



BlueStore on HDDs+SSDs with Fio 4k random writes

➤ `osd_recovery_sleep_hybrid` chosen as 0.025

# Luminous Defaults are much better



# Improving Latency during Recovery

Recovery in Ceph has been a synchronous process - it blocked writes to an object until it was recovered.

Problem: This increases write latencies and affects availability.

## **Solution in Mimic: Asynchronous Recovery**

Do not block writes on objects, which are only missing on non-actingset OSDs.

Perform recovery in the background on an OSD, out of the acting set, similar to backfill, and use the PG log to determine what needs recovery.

# When do we perform asynchronous recovery?

Async recovery targets - OSDs that are not part of the acting set and are chosen based on the following:

- approximate magnitude of the difference in length of logs is used as the cost of recovery, async recovery targets have higher cost to recover
- threshold ***osd\_async\_recovery\_min\_pg\_log\_entries***(default value=100) is used to determine when asynchronous recovery is appropriate
- ***min\_size*** replicas should be available

# Recovery Experiments


RGW Workload generated using Cosbench.

Operations - read, list, write, delete


Cluster prefilled ~ 30%, 40000 objects

Kill one OSD to induce recovery

# Impact on Throughput during Recovery




# Impact on Latency during Recovery






# Throughput Comparison - COSBench

throughput Graph



No Failure

throughput Graph



OSD Failure

# Future Work

- Optimizations
  - Partial object recovery
  - Speed up backfill scanning and transfer
- Adaptive recovery throttling - set the value of `osd_recovery_sleep` based on client load.
- QoS
- Recovery Order

# Summary

- Upgrade! Much better defaults and finer tuning in Luminous and Mimic
- Recovery sleep is all you need if you want to change client i/o vs recovery balance
- Tuning older options (e.g. max active, single start, max backfill) only needed if you want to increase recovery/backfill throughput



# THANK YOU

Josh Durgin : [jdurgin@redhat.com](mailto:jdurgin@redhat.com)

Neha Ojha : [nojha@redhat.com](mailto:nojha@redhat.com)